

Agenda

- **DOCA 2.7**
- **DOCA installation**
- **DPU ARM OS and OVS**
- **DPU Flow Steering and Bond/LAG**
- **Virtio-net Hotplug for Bare-metal**
- **Virtio-net Static for Virtualization**
- **Q&A**

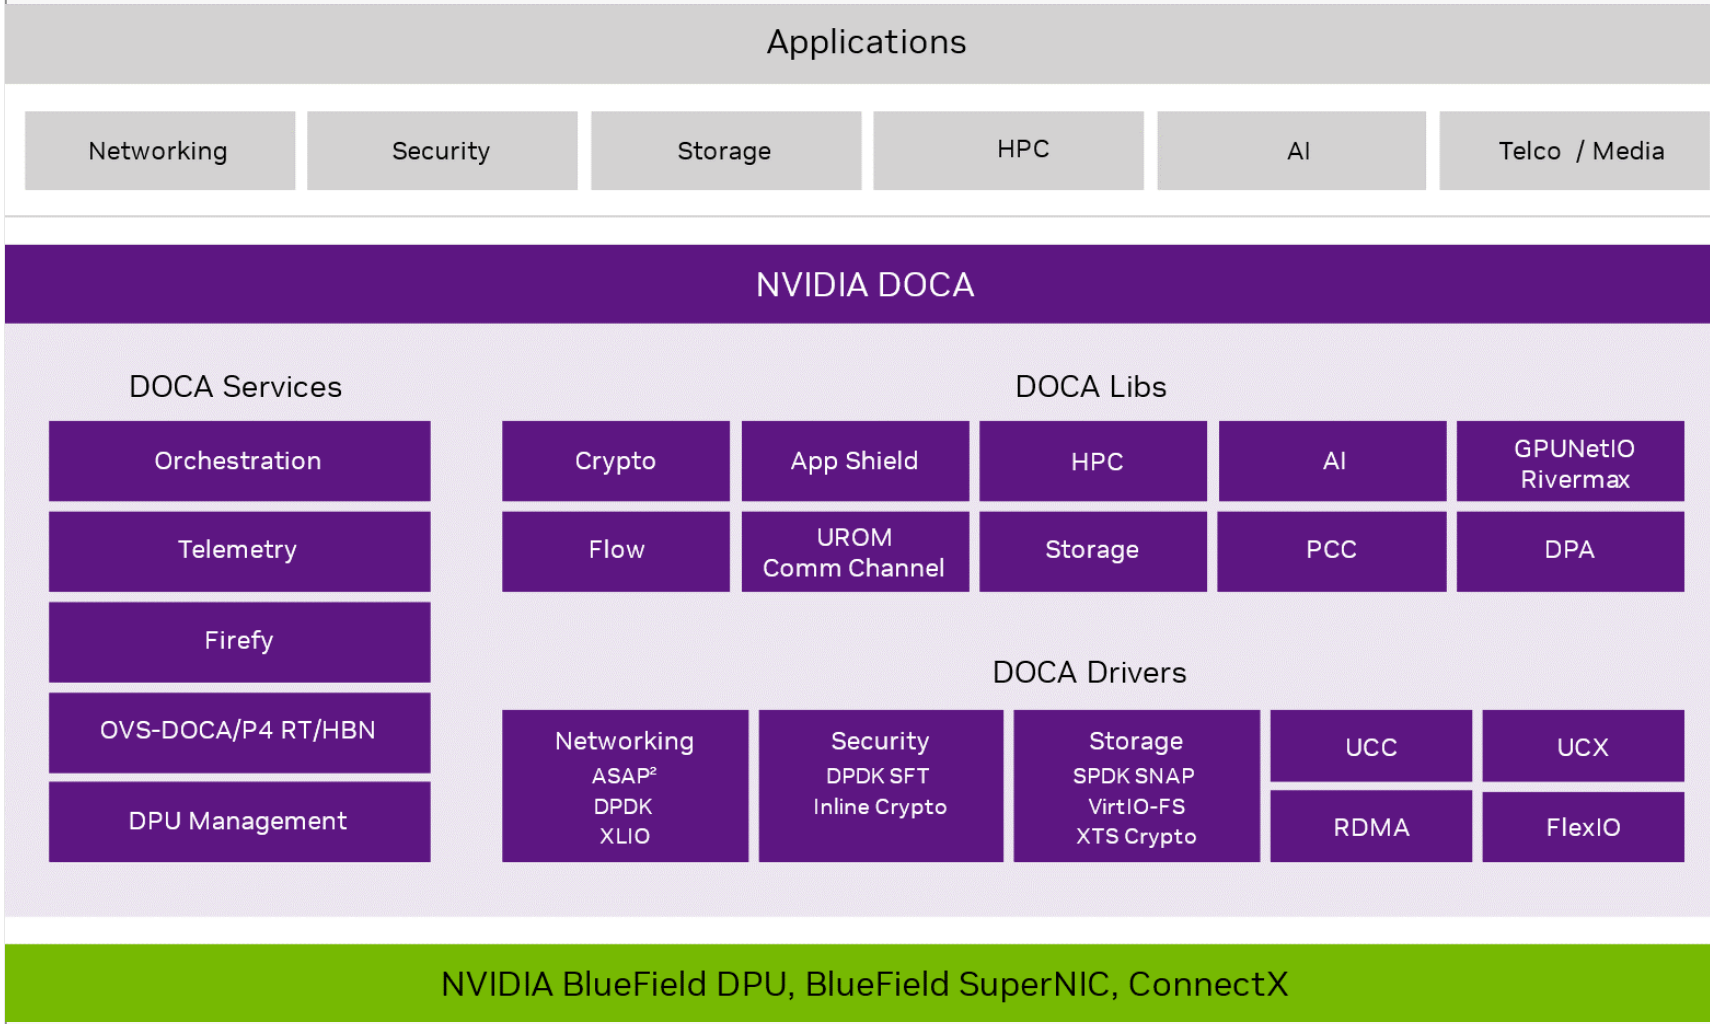
The background features a series of overlapping, diagonal, light green bands that create a sense of depth and movement. The bands are slightly curved and have a soft, glowing appearance. In the top-left corner, there is a solid green vertical bar. The text "DOCA 2.7" is positioned on the left side, partially overlapping the green bar and the background bands.

DOCA 2.7

DOCA 2.7 Major Release Overview

Accelerating AI Cloud East-West and Spectrum-X

- Spectrum-X RA 1.0.1 with BlueField-3 SuperNIC
- DOCA Flow and OVS-DOCA Enhancements
- AI Cloud Traffic Encryption – IPsec GA, PSP support
- DOCA Library Introduction
 - UROM, Device Emulation
- DOCA Services Enhancement and Introduction
 - HBN, Firefly, UROM, DOCA Mgmt. Service (DMS)
- BlueField Platform Software
 - BlueField-3 Upgrade in NIC-Mode and NIC FW Upgrade
 - New BF Bundle Type – BF-FW Bundle
- DOCA Unified Package (MOFED compatible)

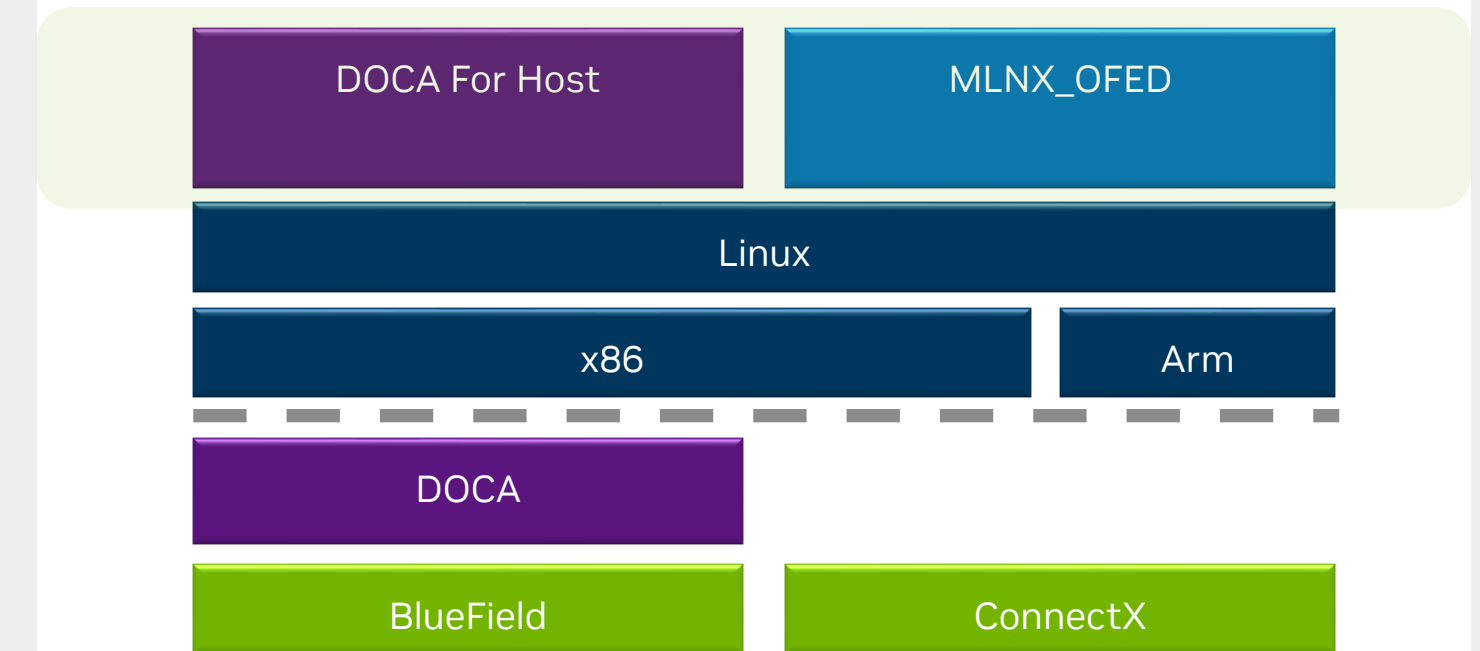


DOCA Unified Package

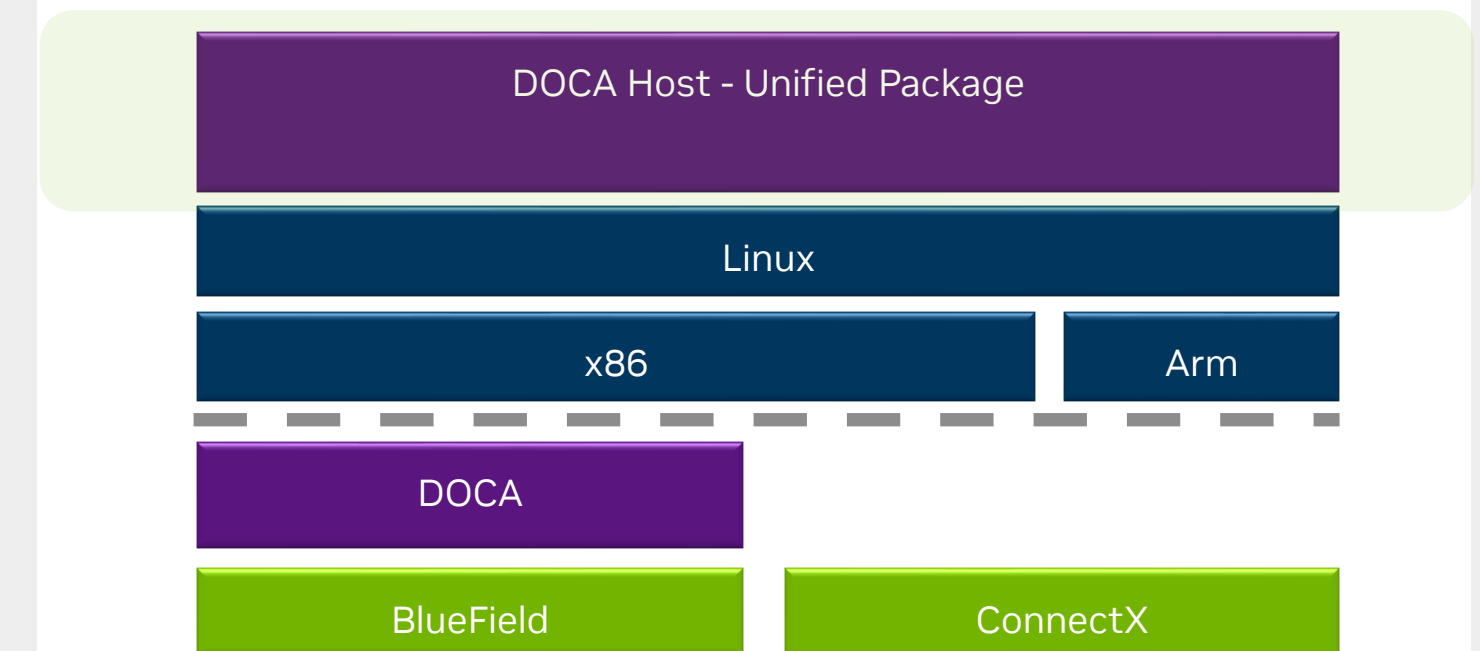
One-Stop-Shop for BlueField and ConnectX Devices

- DOCA-Host
 - Single SW package for the host
 - Replacing MLNX_OFED (last standalone release: Oct 2024)
- Since 2024+
 - DOCA-Host is a one-stop-shop supporting **both BlueField and ConnectX devices**
 - BlueField-3 and ConnectX-8 supported only in DOCA
 - Standard Linux package for host server
 - Host installation profiles: doca-all, doca-networking, doca-ofed
 - **Call for Action:** Promote transition to DOCA Package today
- Use Case and Customers
 - Use DOCA-OFED profile to get the same experience as MLNX-OFED
 - Use the same unified package on a server hosting both BlueField and ConnectX devices

Until DOCA 2.2:



2024+:





New DOCA Download Pages

CUDA-like Download Experience

- DOCA for DPU is what CUDA is for GPU, now also similar download experience
- Remove table from web page, replace with dedicated page

[Download DOCA](#)

- Latest version always in developer.nvidia.com/doca-downloads
- LTS and Previous versions in developer.nvidia.com/doca-archive

DOCA 2.7.0 Downloads

Select

Click on the green buttons that describe your target platform. Only supported platforms will be shown. By downloading and using the software, you agree to fully comply with the terms and conditions of the [DOCA EULA](#).

Host or BlueField	Host-Server	BlueField							
Deployment Package	DOCA-Host	BF-NIC-Firmware	CX-Firmware						
Operating System	Linux	Windows							
Architecture	x86_64	arm64-sbsa							
Profile	doca-ofed	doca-all	doca-networking						
Distribution	Mariner	Ubuntu	Debian	KylinOS	TencentOS	CTyunOS	Fedora	OracleLinux	
	RHEL	UOS	XenServer	SLES	openEuler	AnolisOS	BCLinux	EulerOS	
Version	9.0	8.3	9.3	8.7	8.5	8.4	8.0	9.2	8.1
Installer Type	rpm (local)	rpm (online)							

Download Installer for Linux RHEL 9.3 x86_64

The base installer is available for download below.

> DOCA Online Repository doca-ofed

Installation Instructions:

```
$ echo "[doca]
name=DOCA Online Repo
baseurl=https://doca-repo-prod.nvidia.com/public/repo/doca/2.7.0/rhel9.3/x86_64/
enabled=1
gpgcheck=0" > /etc/yum.repos.d/doca.repo

$ sudo dnf clean all
$ sudo dnf -y install doca-ofed
```

Latest Releases

[DOCA 2.7.0 \(April 2024\)](#), [Versioned Online Documentation](#)

[DOCA 2.5.1 LTS Update \(February 2024\)](#), [Versioned Online Documentation](#)

[DOCA 1.5.3 LTS Update \(December 2023\)](#), [Versioned Online Documentation](#)

Archived Releases

[DOCA 2.6.0 \(January 2024\)](#), [Versioned Online Documentation](#)

[DOCA 2.5.0 LTS \(October 2023\)](#), [Versioned Online Documentation](#)

[DOCA 2.2.1 BF-3 \(October 2023\)](#), [Versioned Online Documentation](#)

[DOCA 2.2.0 BF-2 \(August 2023\)](#), [Versioned Online Documentation](#)

[DOCA 1.5.2 LTS Update \(June 2023\)](#), [Versioned Online Documentation](#)

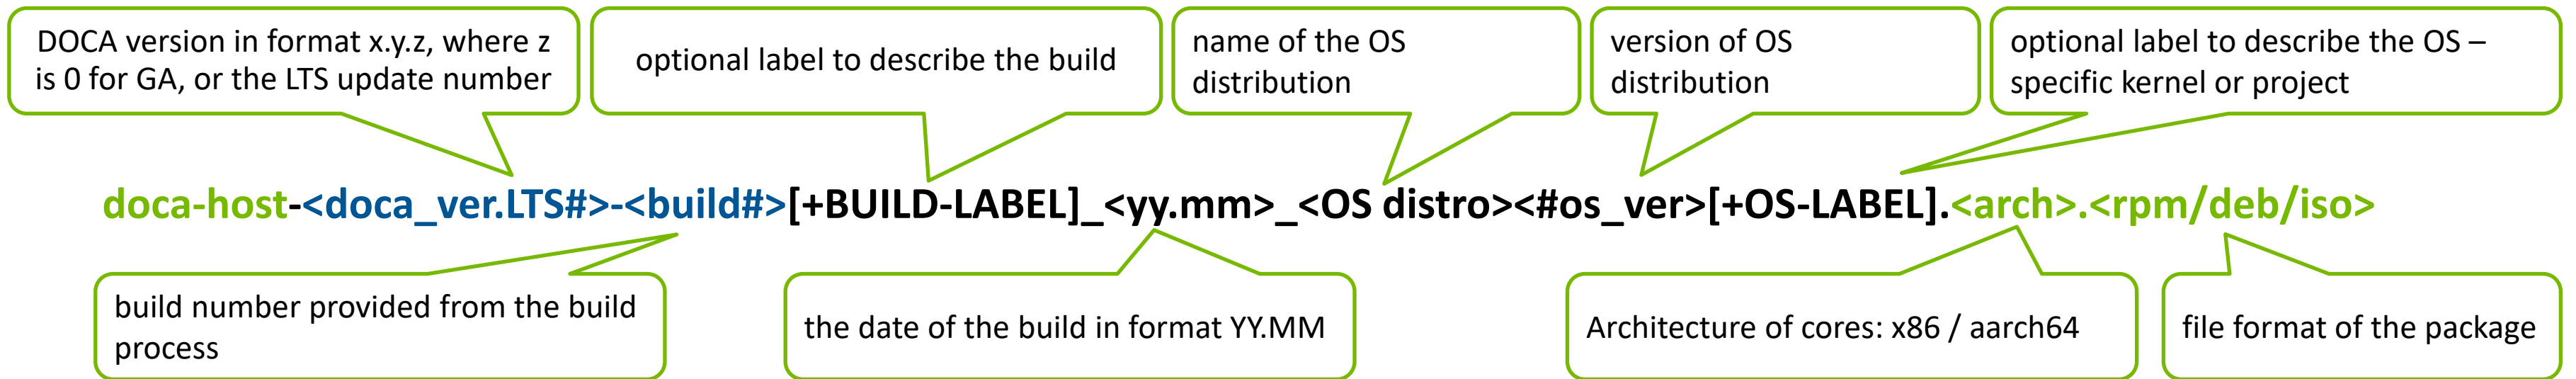
[DOCA 2.0.2 \(May 2023\)](#), [Versioned Online Documentation](#)

[DOCA 1.5.1 LTS Update \(November 2022\)](#), [Versioned Online Documentation](#)

[DOCA 1.5.0 LTS \(October 2022\)](#), [Versioned Online Documentation](#)

DOCA-Host Image Name Format

doca-host



The above is for .rpm and .iso files.

For deb: doca-host_<doca_ver.LTS#>-<build#>[+BUILD-LABEL]-<yy.mm>-<OS distro><#os_ver no decimal point>[+OS-LABEL]_<arch>.deb

(to comply with rpm/deb parsing format)

Examples:

doca-host-2.6.0-1234+L0123_24.01_ubuntu2204+61.x86.rpm

LTS: doca-host-2.5.1-1250+L3_24.02_ubuntu2204+61.x86.rpm

FUR: doca-host-2.6.0-1344+MG54_24.03_ubuntu2204+61.x86.rpm

New BF Bundle Types


What's new in DOCA 2.7

- BF-Bundle – default image
 - The full image with BlueField components, including – ATF, UEFI, nic-fw, bmc-fw, eROT, Ubuntu OS and DOCA runtime
- BF-Firmware Bundle (BF-FWBundle) - ‘firmware only’
 - Smaller image for day2 upgrades – skipping the OS and DOCA upgrade
 - Includes – ATF, UEFI, nic-fw, bmc-fw, eROT
 - DOCA and OS can be separately upgraded using standard Linux tools – apt/yum
- Supported Host and DPU OS Distributions
 - <https://docs.nvidia.com/networking/display/bluefielddpuosv470/supported+platforms+and+interoperability>

NEW

 [DOCA Installation Guide](#)

BF-Bundle BFB Image DOCA 2.7 (BFB)



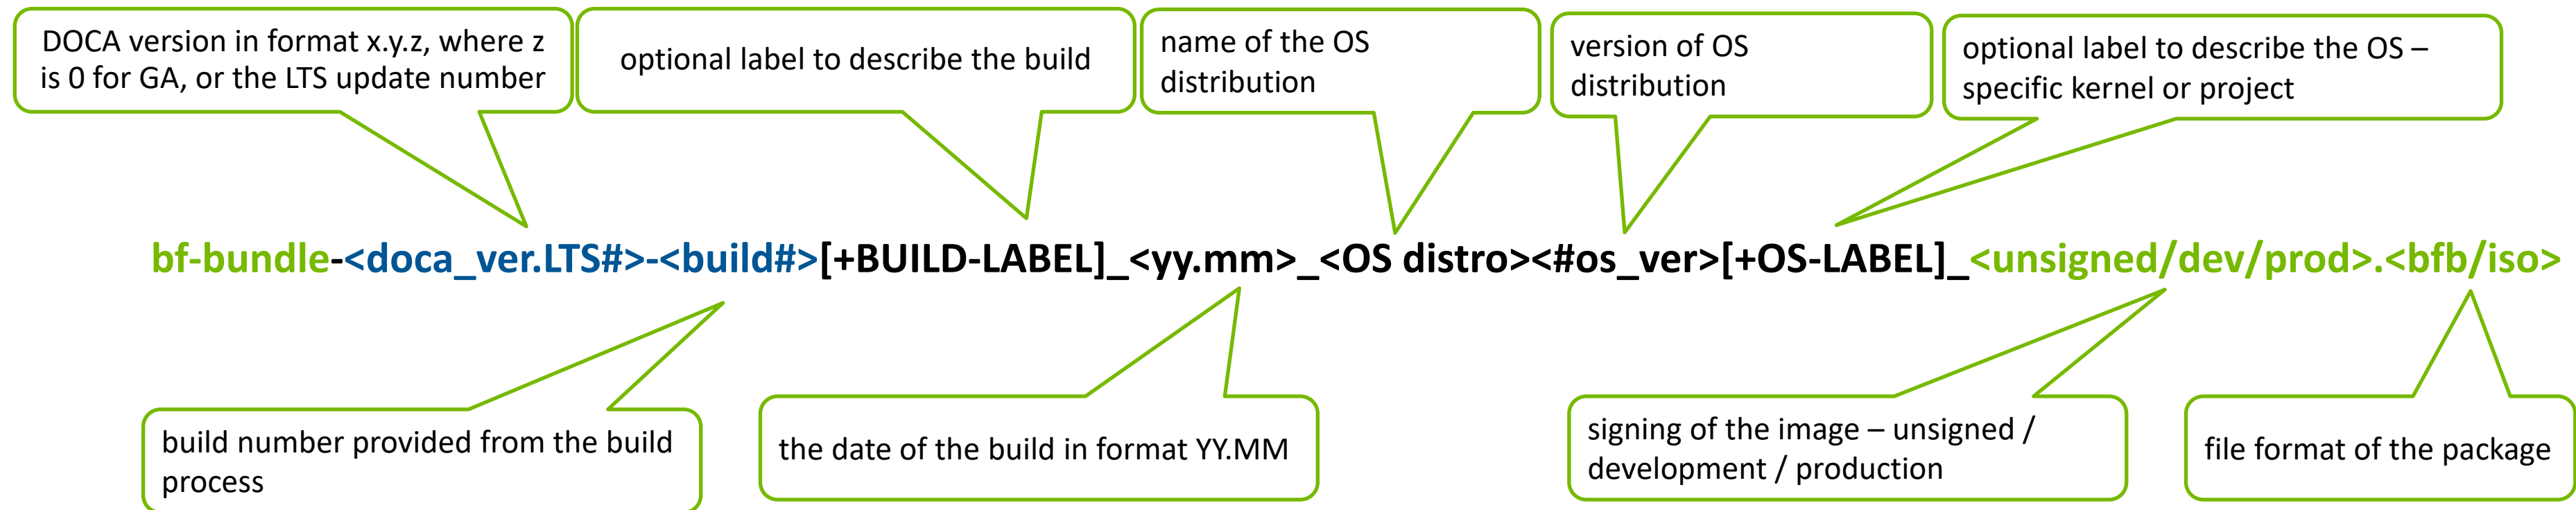
BF-Firmware-Bundle BFB Image DOCA 2.7 (BFB)



BMC/Host/RSHIM updates

BlueField Bundle Image Name Format

bf-bundle



Examples:

bf-bundle-2.6.0-1737_0567-24.01-rhel-8.6_4.19_dev.iso

LTS: bf-bundle-2.5.1-1421_24.01_rhel86_prod.bfb

FUR on LTS: bf-bundle-2.5.1+**1422**_24.02_rhel86_prod.bfb

FUR: bf-bundle-2.5.0-1428+**X123**_24.03_rhel86_dev.iso

DOCA Installation

Installing DOCA on the Host

Host or BlueField

Host-Server

BlueField

Deployment Package

DOCA-Host

BF-NIC-Firmware

CX-Firmware

Operating System

Linux

Windows

Architecture

x86_64

arm64-sbsa

Profile

doca-all

doca-networking

doca-ofed

Distribution

Ubuntu

RHEL

OracleLinux

Debian

ALinux

CTyunOS

Version

22.04


20.04

Installer Type

deb (local)

deb (online)

DOCA functionality is limited by the specific device capabilities.



DOCA-HOST Profiles

DOCA Local Repository doca-all

```
wget https://www.mellanox.com/downloads/DOCA/DOCA_v2.7.0/host/doca-host_2.7.0-204000-24.04-ubuntu2204_amd64.deb
sudo dpkg -i doca-host_2.7.0-204000-24.04-ubuntu2204_amd64.deb
sudo apt-get update
sudo apt-get -y install doca-all mlnx-fw-updater
```

DOCA-OFED profile

- This profile is intended for users who wish to have the exact same user experience and content as MLNX_OFED but with DOCA Package. doca-ofed installs the MLNX_OFED drivers and tools but not other DOCA components.
- Note: The rshim package is included in OFED/this profile as well
- The content of the doca-ofed package is as follows:
 - MLNX_OFED drivers
 - MLNX_OFED tools
- Currently supported operating systems and kernel version per DOCA version and profiles are outlined on our DOCA SDK page (note the column for the doca-ofed profile)

<https://docs.nvidia.com/doca/sdk/profiles/index.html#supported-host-os-per-doca-profile>

- The doca-extra package, located under /opt/mellanox/doca/tools/, contains:
 - doca-info – displays details of all installed dependencies in DOCA
 - doca-kernel-support – if using a kernel not listed in the [Supported Kernel Versions](#) table, this tool downloads the appropriate missing packages for your kernel to support DOCA, unless the packages do not support the existing kernel.

Installing DOCA on Host

Verifying successful installation

- Verify the DOCA packages installed on the host.

```
host# dpkg --get-architecture | grep amd64
ii  doca-apps                1.5.1007-1                amd64  DOCA-Based reference applications
ii  doca-apps-dev            1.5.1007-1                amd64  Development files for DOCA Apps
ii  doca-grpc                1.5.1007-1                amd64  gRPC runtime capabilities for DOCA
ii  doca-grpc-dev            1.5.1007-1                amd64  Development files for DOCA grpc
ii  doca-host-repo-ubuntu1804 1.5.1-0.1.8.1.5.1007.1.5.8.1.1.2.1 amd64  Doca repo bundle package
ii  doca-libs                1.5.1007-1                amd64  Data Center on a Chip Architecture (DOCA)
ii  doca-prime-runtime       1.5.1007-1                amd64  DOCA prime runtime metapackage
ii  doca-prime-sdk           1.5.1007-1                amd64  DOCA prime sdk metapackage
ii  doca-prime-tools         1.5.1007-1                amd64  Runtime DOCA Tools
ii  doca-remote-memory-app   22.07.0                   amd64  Nvidia Bluefield regex benchmarking tool
companion remote memory app.
ii  doca-runtime             1.5.1-0.1.8               amd64  doca-runtime meta-package
ii  doca-samples             1.5.1007-1                amd64  DOCA Samples
ii  doca-sdk                 1.5.1-0.1.8               amd64  doca-sdk meta-package
ii  doca-tools               1.5.1-0.1.8               amd64  doca-tools meta-package
ii  libdoca-libs-dev         1.5.1007-1                amd64  Development files for DOCA Libs
```


Install DOCA on DPU

prepare

- Host side
 - Make sure to install host drivers
 - Check rshim status
- DPU side
 - Connect to your BlueField DPU
 - Check BSP version
 - Check NIC-FW/BMC-FW
- Host side
 - Download BFB Image

```
root@localhost:~# flint -d /dev/mst/mt41692_pciconf0 q
Image type:      FS4
FW Version:      32.40.1000
FW Release Date: 4.2.2024
Product Version: 32.40.1000
Rom Info:        type=UEFI Virtio net version=21.4.13 cpu=AMD64,AARCH64
                  type=UEFI Virtio blk version=22.4.12 cpu=AMD64,AARCH64
                  type=UEFI version=14.33.10 cpu=AMD64,AARCH64
                  type=PXE version=3.7.300 cpu=AMD64
Description:      UID          GuidNumber
Base GUID:       a088c2030075c110    38
Base MAC:        a088c275c110      38
Image VSD:       N/A
Device VSD:      N/A
PSID:            MT_00000000884
Security Attributes: secure-fw
```

```
root@localhost:~# cat /etc/mlnx-release
DOCA_2.6.0_BSP_4.6.0_Ubuntu_22.04-5.24-01.prod
```

```
root@localhost:~# ipmitool mc info
Device ID          : 1
Device Revision    : 1
Firmware Revision  : 23.04
IPMI Version       : 2.0
```

Install DOCA on DPU

Install BFB

```
bfb-install --bfb bf-bundle-2.7.0-33_24.04_ubuntu-22.04_prod.bfb  
--rshim /dev/rshim0
```

- Host side
 - Run the bfb-install script while indicating the BFB image to install and the DPU device.

Pushing bfb

Collecting BlueField booting status. Press Ctrl+C to stop...

INFO[PSC]: PSC BL1 START

INFO[BL2]: UEFI loaded

INFO[UEFI]: eMMC init

INFO[UEFI]: eMMC probed

INFO[MISC]: Ubuntu installation started

INFO[MISC]: Installing OS image

INFO[MISC]: Ubuntu installation completed

WARN[MISC]: Skipping BMC components upgrade.

INFO[MISC]: Updating NIC firmware...

INFO[MISC]: NIC firmware upd

INFO[UEFI]: exit Boot Service

INFO[MISC]: Linux up

INFO[MISC]: DPU is ready

Install DOCA on DPU

Verify the new BFB version

- DPU side
 - After BFB upgrade, the credentials are reset to their default.
 - Default BFB Ubuntu credentials:
 - User: ubuntu
 - Password: ubuntu
 - Check BSP/NIC FW/BMC version

```
ubuntu@localhost:~$ sudo flint -d /dev/mst/mt41692_pciconf0 q
Image type:      FS4
FW Version:      32.41.1000
FW Version(Running): 32.40.1000
FW Release Date: 28.4.2024
Product Version: 32.40.1000
Rom Info:        type=UEFI Virtio net version=21.4.13 cpu=AMD64,AARCH64
                  type=UEFI Virtio blk version=22.4.12 cpu=AMD64,AARCH64
                  type=UEFI version=14.33.10 cpu=AMD64,AARCH64
                  type=PXE version=3.7.300 cpu=AMD64
Description:      UID          GuidNumber
Base GUID:        a088c2030075c110    38
Base MAC:         a088c275c110      38
Image VSD:        N/A
Device VSD:       N/A
PSID:             MT_00000000884
Security Attributes: secure-fw
```

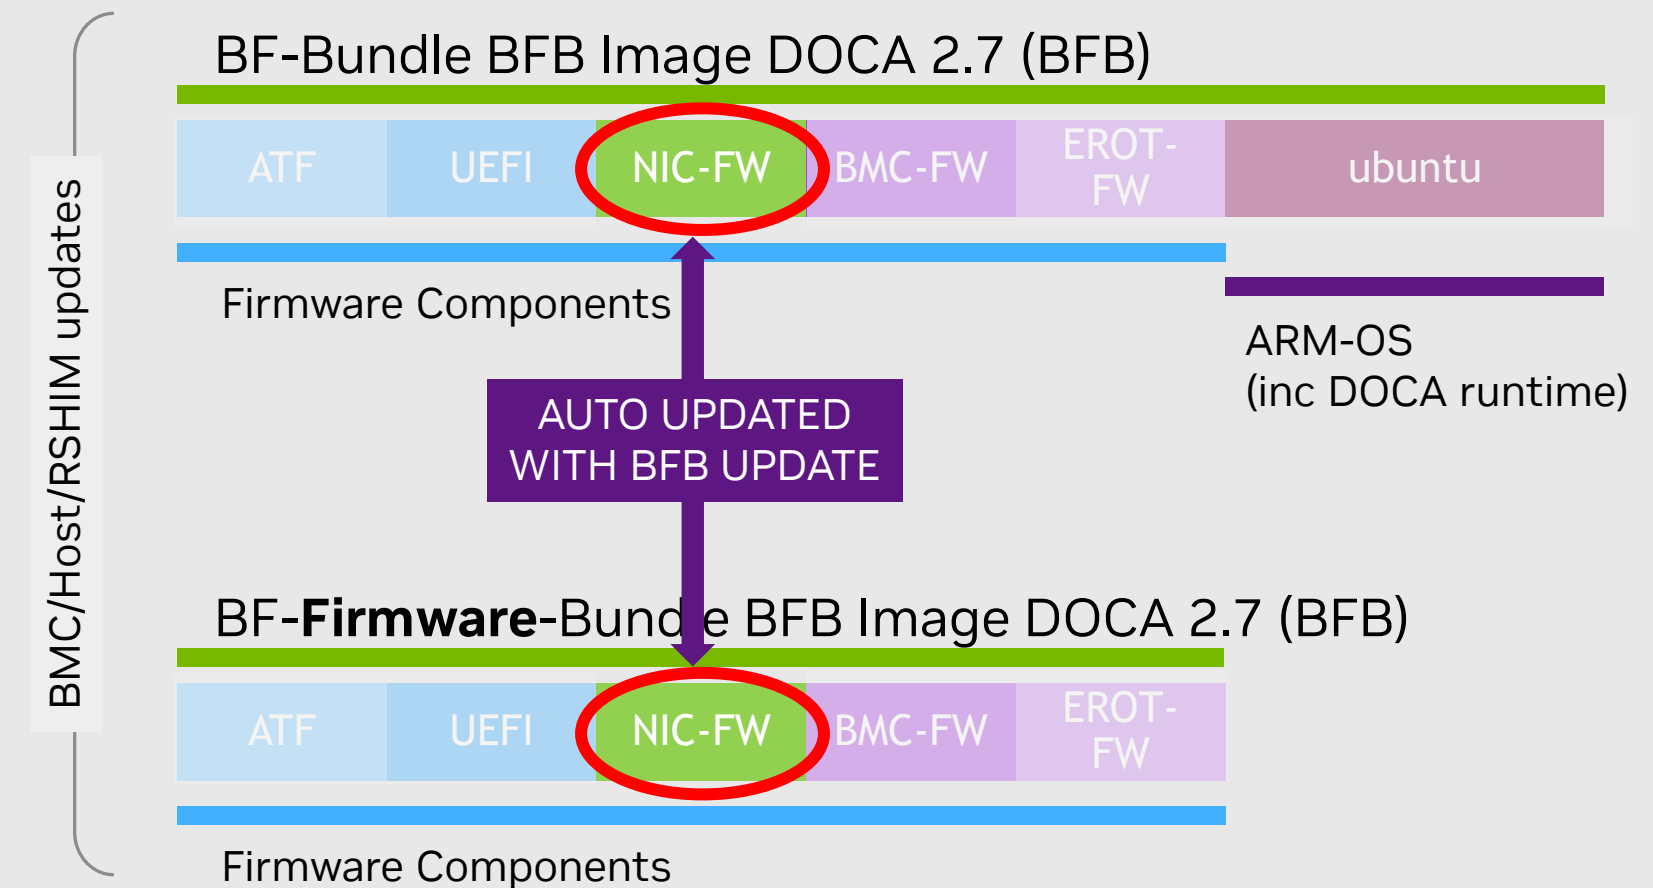
```
ubuntu@localhost:~$ cat /etc/mlnx-release
bf-bundle-2.7.0-33_24.04_ubuntu-22.04_prod
```

```
ubuntu@localhost:~$ sudo ipmitool mc info
Device ID          : 1
Device Revision    : 1
Firmware Revision  : 23.04
```

BFB NIC FW Update by Default

Keeping BlueField Components In Sync

- Keeping all BlueField components in sync to an official NVIDIA BlueField release is crucial for the stability of the product
- Starting April/24 DOCA 2.7 – any upgrade done with BFB image (full or ‘small’) will automatically upgrade the NIC-FW image
- Up to Apr/24 release – customers were required to take manual steps to burn BlueField NIC-FW – this is no longer needed



Customizations During BFB Installation

- The BlueField's UEFI system, boot options and more can be customized through the use of configuration parameters in the bf.cfg file.

Command: `bfb-install --bfb DOCA_2.6.0_BSP_4.6.0_Ubuntu_22.04-5.24-01.prod.bfb --config bf.cfg --rshim rshim0`

Tasks:

- Install target OS if `UPDATE_DPU_OS='yes'`(default)
- Update ATF and UEFI if `UPDATE_ATF_UEFI="yes"`(default)
- Update BMC components: `UPDATE_BMC_FW="yes"`(default)+`UPDATE_CEC_FW="yes"`(default)
- Update NIC firmware if `WITH_NIC_FW_UPDATE="yes"`(default)
- , Ubuntu users can provide a unique password that will be applied at the end of the BFB installation.
- Step1: Create password hash

```
# openssl passwd -1
```

Password:

Verifying - Password:

```
$1$3B0RIrfX$TlHry93NFUJzg3Nya00rE1
```

- Step2: Add the password hash in quotes to the bf.cfg file

```
# vim bf.cfg
```

```
ubuntu_PASSWORD=' $1$3B0RIrfX$TlHry93NFUJzg3Nya00rE1 '
```

Custom BFB

- How to build your own customized bfb
- [GitHub - Mellanox/bfb-build: BFB \(BlueField boot stream and OS installer\) build environment](#)

BFB/DOCA upgrade process - Today

Upgrade repo, Run:

- ☐ \$ wget "doca-dpu-repo"
- ☐ \$ dpkg -i "doca-dpu-repo"
- ☐ \$ sudo apt-get update
- ☐ \$ sudo apt-get upgrade

NOTICE: Today, needed download "doca-dpu-repo" from the artifactory location.

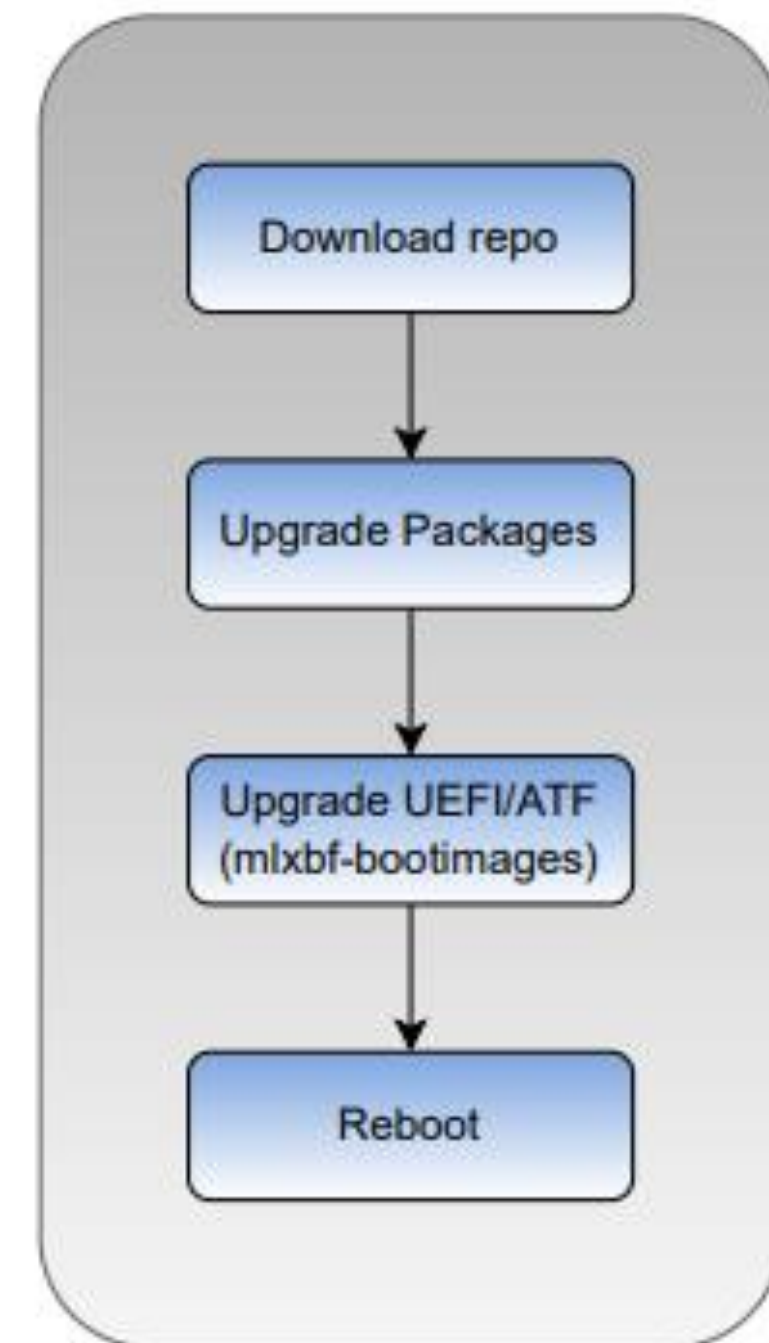
[NVIDIA DOCA Installation Guide for Linux - NVIDIA Docs](#)

Upgrade UEFI/ATF (included in mlxbf-bootimages DEB package) on the boot partition, Run:

- ☐ \$ bfrec --bootctl --policy dual
- ☐ \$ bfrec --capsule /lib/firmware/mellanox/boot/capsule/boot_update2.cap --policy dual
- ☐ \$ reboot

Upgrade firmware, Run

- ☐ dpu# sudo /opt/mellanox/mlnx-fw-updater/mlnx_fw_updater.pl --force-fw-update
- ☐ Check if the driver is supported sync 1
- ☐ dpu# sudo mlxfwreset -d /dev/mst/mt*_pciconf0 --sync 1 -y reset

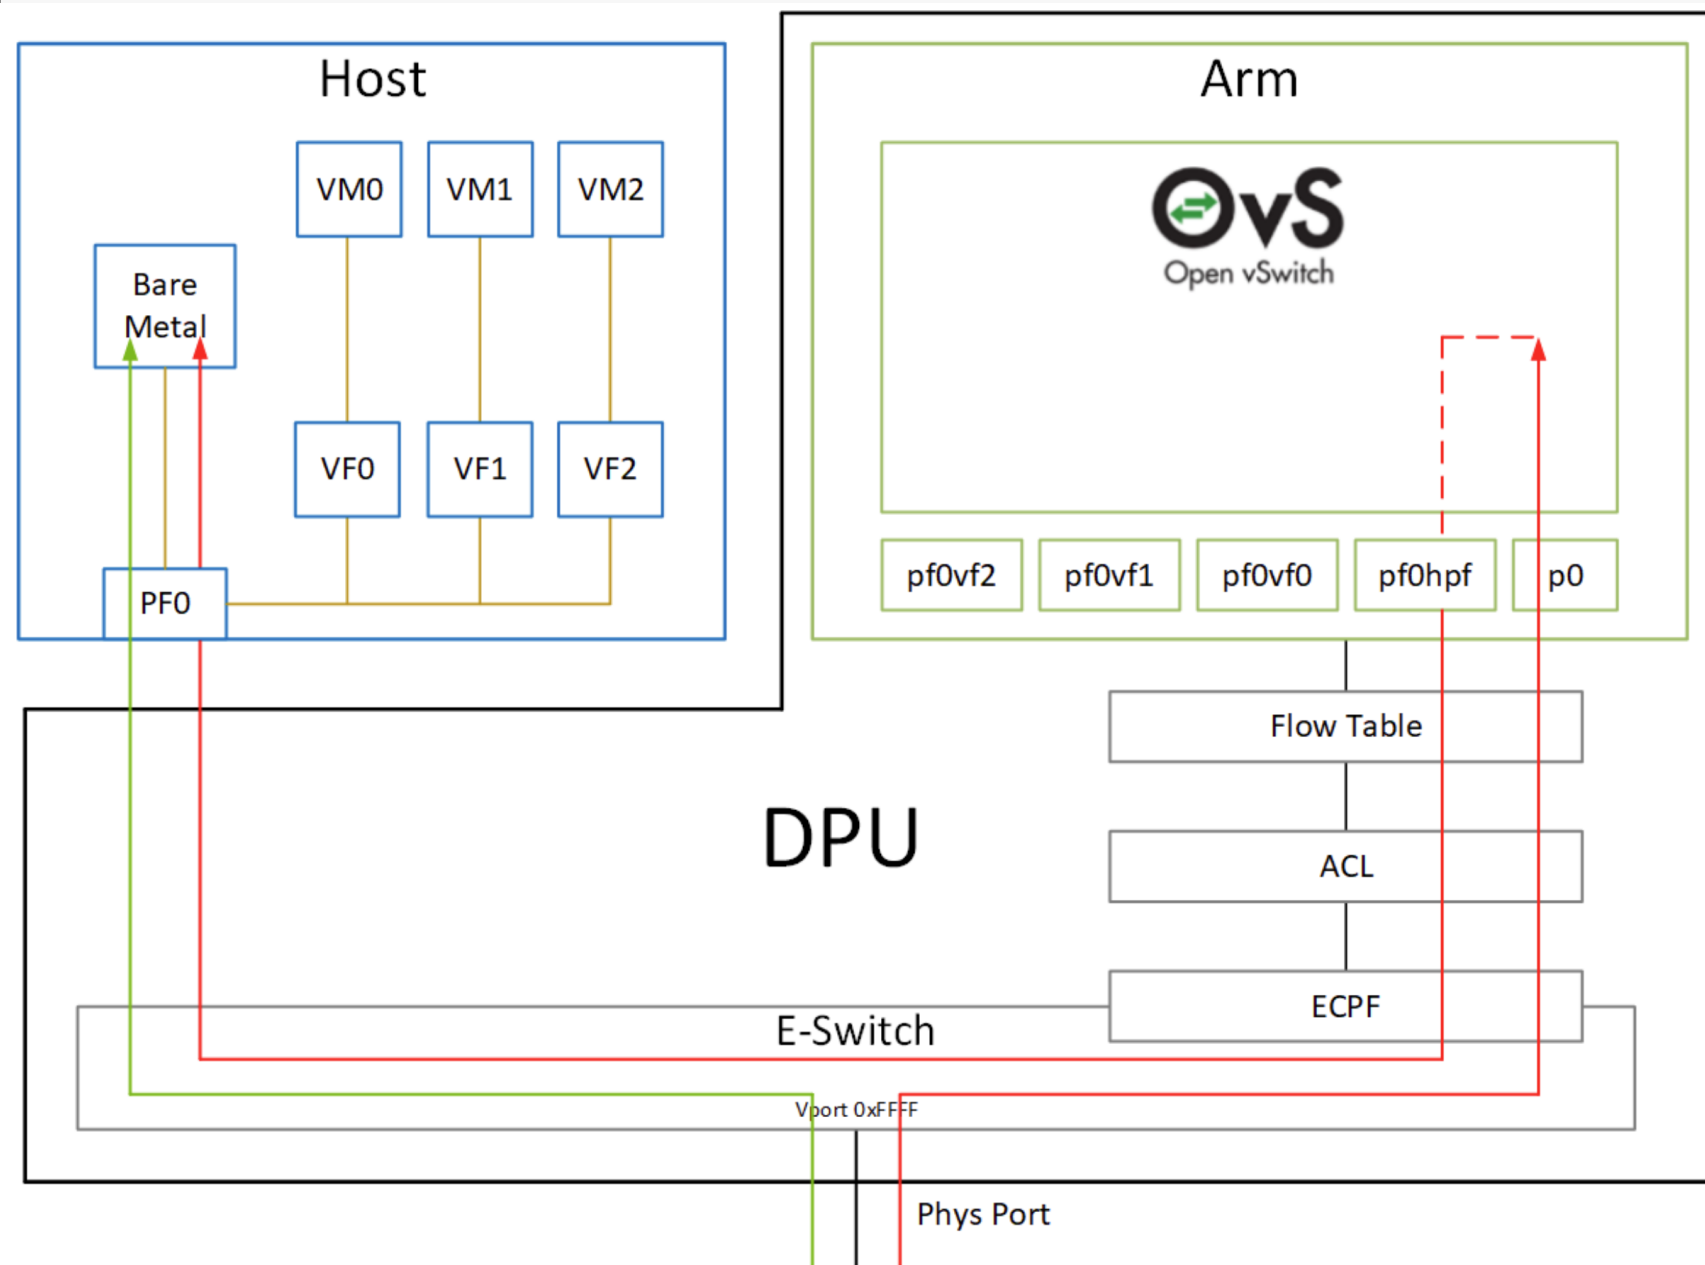


DPU ARM OS and OVS

BlueField DPU Interfaces

- The following interfaces can be seen on the DPU Arm OS :
 - OOB interface: **oob_net0**.
By default, IP address is assigned by DHCP.
 - RShim interface: **tmfifo_net0**.
Pre-configured with a fixed IP address: 192.168.100.2.
 - DPU's PFs: **p0, p1**
 - Host PF representors: **pf0hpf, pf1hpf**
 - Host VF representors (if VFs are created on the host, one per VF): **pf0vf0, pf1vf0...**
 - Default SFs (one per PF): **enp3s0f0s0, enp3s0f1s0**
 - Default SF representors: **en3f0pf0sf0, en3f1pf1sf0**
 - Default OVS bridge ports: **ovsbr1, ovsbr2**
- Applicable only for DPU mode.


Kernel Representors Model



- BlueField DPU uses netdev representors to map each one of the host-side physical and virtual functions.
- These representors serve as:
 - The tunnel to pass traffic for the virtual switch or application running on the Arm cores to the relevant PF or VF on the host side.
 - The channel to configure the embedded switch with rules to the corresponding represented function.
- Those representors are used as the virtual ports being connected to OVS or any other virtual switch running on the Arm cores.

DPU PFs

Representors for DPU network ports



- Two representors are created for each one of the DPU's network ports:
 - For the uplink
 - For the host-side PF
- Naming convention for PF representors:
 - Uplink representor: p<port_number>
 - Host side PF representor: pf<port_number>hpf

DPU PFs Representor Example

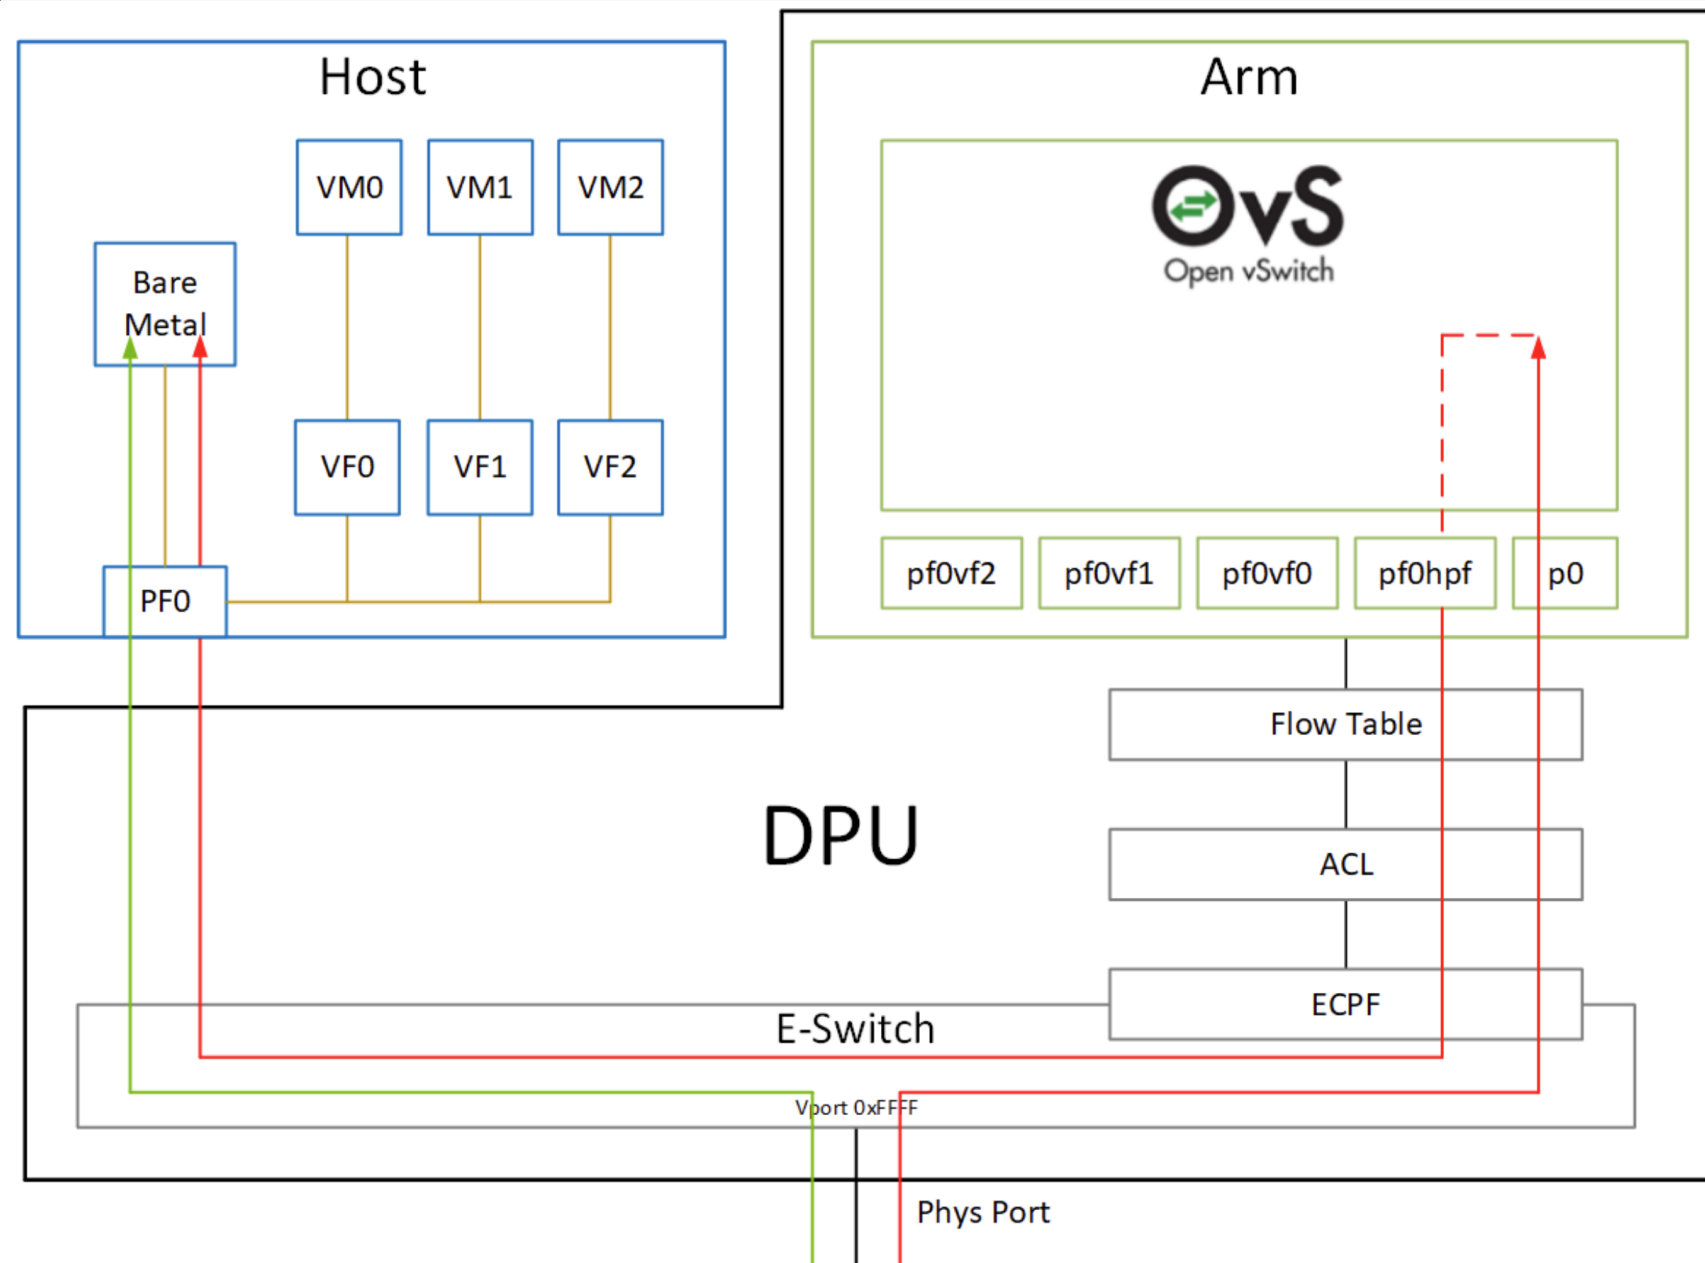
Representors for DPU network ports example

- Run the command “ip a” on the DPU to display a list of network representors:

```
dpu # ip a
4: p0: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc mq master ovs-system state DOWN group default qlen 1000
    link/ether e8:eb:d3:ce:cc:ec brd ff:ff:ff:ff:ff:ff
5: p1: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc mq master ovs-system state DOWN group default qlen 1000
    link/ether e8:eb:d3:ce:cc:ed brd ff:ff:ff:ff:ff:ff
6: pf0hpf: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc mq master ovs-system state UP group default qlen 1000
    link/ether 7e:df:8b:a7:ef:e2 brd ff:ff:ff:ff:ff:ff
    inet6 fe80::7cdf:8bff:fea7:efe2/64 scope link
        valid_lft forever preferred_lft forever
7: pf1hpf: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc mq master ovs-system state UP group default qlen 1000
    link/ether ba:b9:67:c1:e7:e1 brd ff:ff:ff:ff:ff:ff
    inet6 fe80::b8b9:67ff:fec1:e7e1/64 scope link
        valid_lft forever preferred_lft forever
```

- The DPU's first network port is represented by:
 - Uplink representor: p0
 - Host side PF representor: pf0hpf
- The DPU's second network port is represented by:
 - Uplink representor: p1
 - Host side PF representor: pf1hpf

Host VFs Representors



- For each of the VFs created on the host-side, a corresponding representor is created on the DPU Arm-side.
- Naming convention for VF representors:
 - pf<port_number>vf<function_number>

Host VFs Representor Example

- Create VFs on the host:

Number of VFs to create

```
host # echo 2 > /sys/class/net/p1p1/device/sriov_numvfs
```

- Display the list of network representors on the DPU:

```
Dpu # ip a
```

```
16: pf0vf0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc mq state UP group default qlen 1000
    link/ether 22:f0:f2:cc:d4:a3 brd ff:ff:ff:ff:ff:ff
    inet6 fe80::20f0:f2ff:fecc:d4a3/64 scope link
        valid_lft forever preferred_lft forever
17: pf0vf1: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc mq state UP group default qlen 1000
    link/ether 66:48:3d:87:59:c4 brd ff:ff:ff:ff:ff:ff
    inet6 fe80::6448:3dff:fe87:59c4/64 scope link
        valid_lft forever preferred_lft forever
```

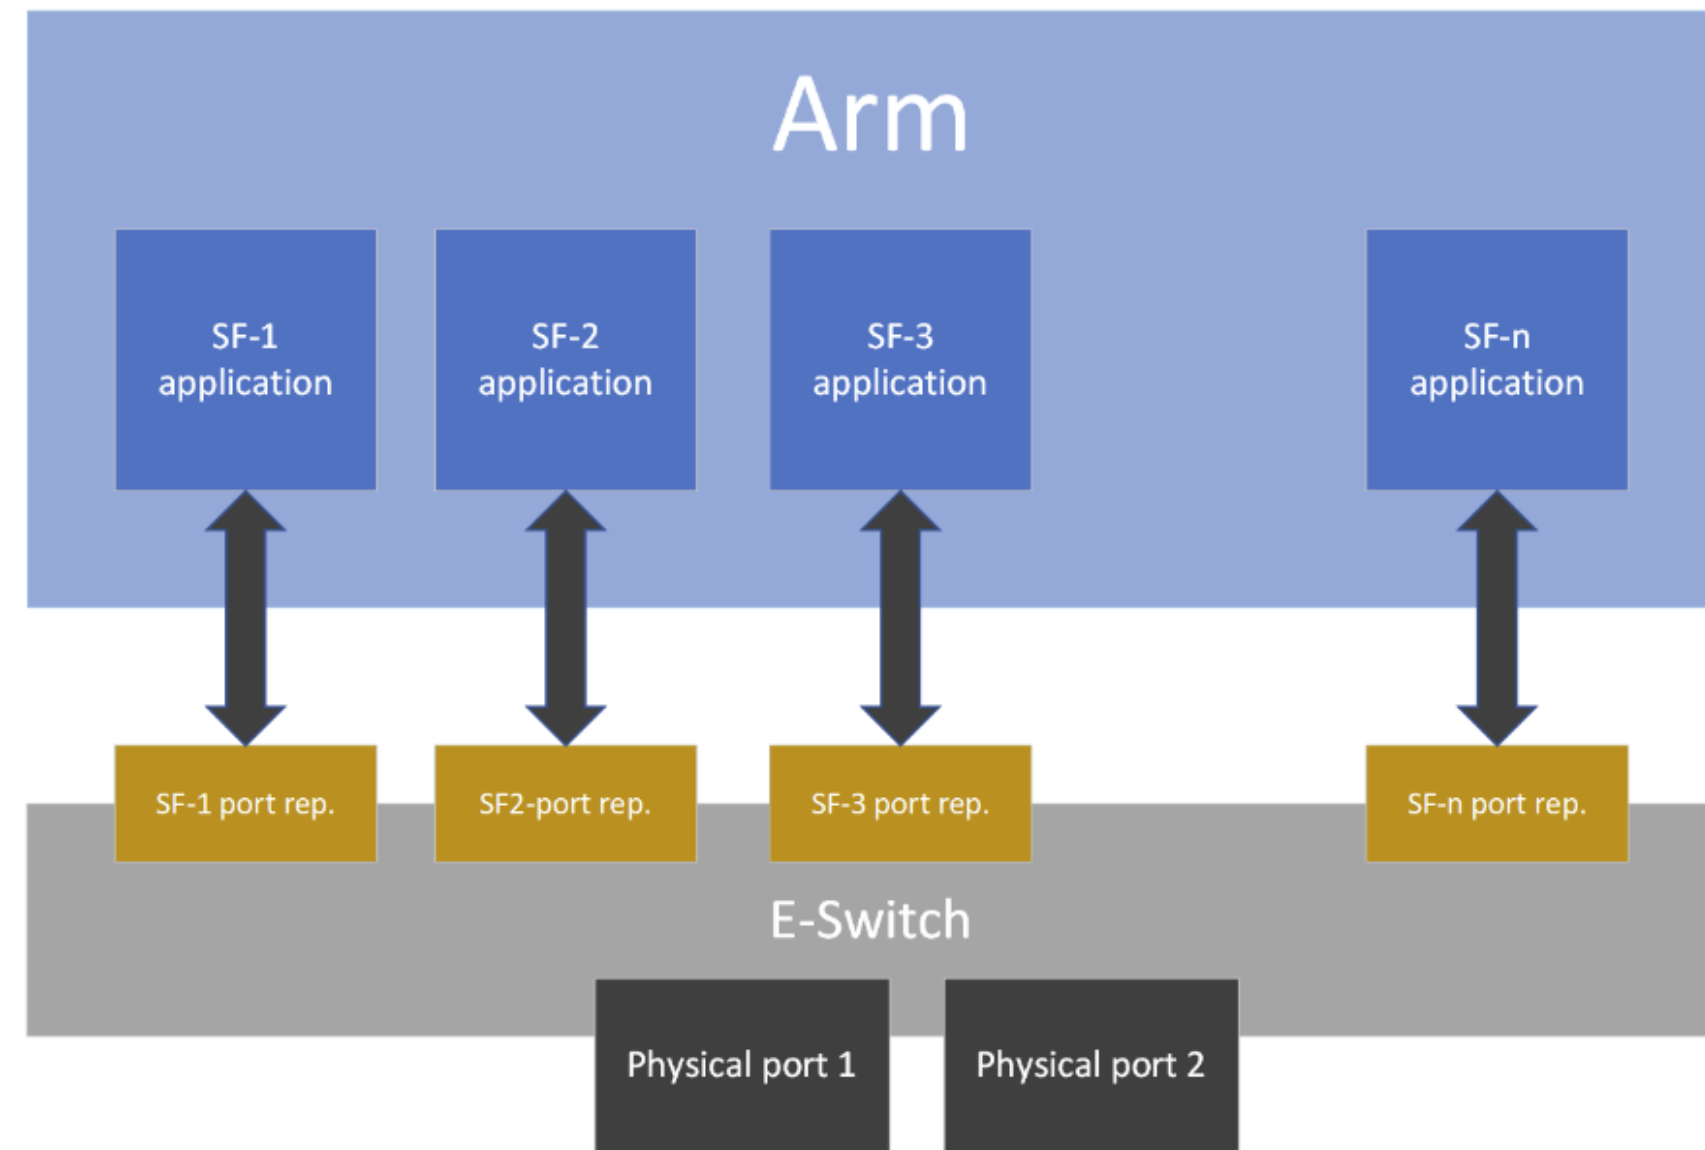
Representor for VF0

Representor for
VF1

- For each VF created on the host, a corresponding representor is created on the DPU.

Scalable Functions (SFs)

- SFs are very similar to VFs which are part of the SR-IOV solution.
- An SF is a lightweight function which has a parent PCIe function on which it is deployed.
- The SF has access to the capabilities and resources of its parent PCIe function, its own function capabilities, and its own resources.
- SFs co-exist with PCIe SR-IOV virtual functions (on the host) but do not require enabling PCIe SR-IOV.
- SFs can be used as:
 - Internal interface in Arm to use between Arm and the external world.
 - RoCE transport between host and Arm and Arm and the external host (for SNAP etc.).



SF Representors

- One SF per DPU PF is created
 - Default SFs: **enp3s0f0s0, enp3s0f1s0**
 - Defaults SF representors: **en3f0pf0sf0, en3f1pf1sf0**
- Additional SFs can be created on the DPU by the user.

SFs Representors Example

- Run the command “ip a” on the DPU to display a list of network representors.

The first DPU's PF:

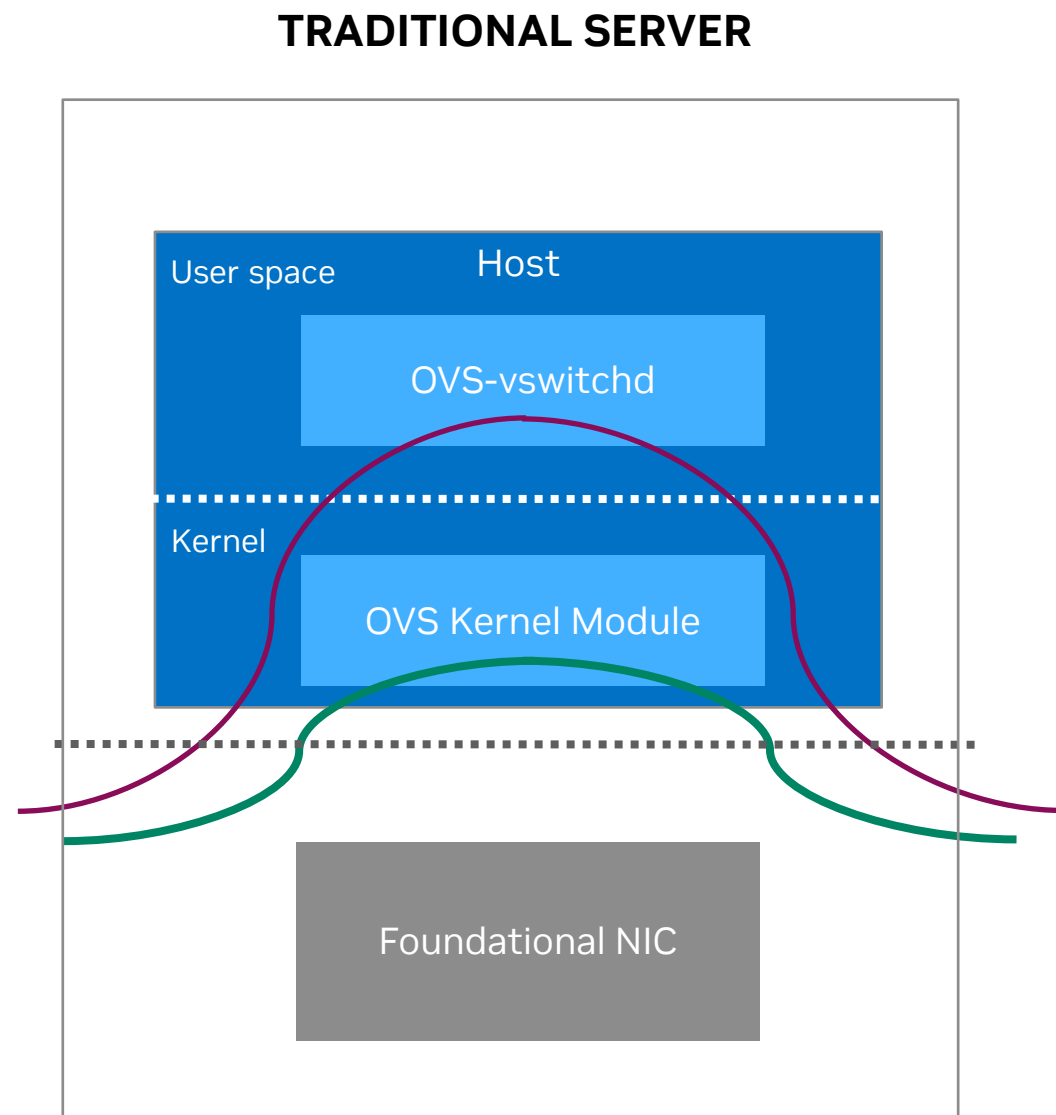
- SF: enp3s0f0s0
- SF representor: en3f0pf0sf0

The second DPU's PF:

- SF: enp3s0f1s0
- SF representor: en3f1pf1sf0

```
dpu# ip a
8: en3f0pf0sf0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc mq master ovs-system state UP
   group default qlen 1000
    link/ether 8a:45:b3:56:6d:75 brd ff:ff:ff:ff:ff:ff
9: enp3s0f0s0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc mq state UP group default qlen 1000
   link/ether 02:04:5f:68:ef:1e brd ff:ff:ff:ff:ff:ff
   inet6 fe80::4:5fff:fe68:ef1e/64 scope link
       valid_lft forever preferred_lft forever
10: en3f1pf1sf0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc mq master ovs-system state UP
   group default qlen 1000
    link/ether 7a:b5:4c:0c:37:55 brd ff:ff:ff:ff:ff:ff
11: enp3s0f1s0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc mq state UP group default qlen
1000
    link/ether 02:7e:f1:e2:c4:a2 brd ff:ff:ff:ff:ff:ff
    inet6 fe80::7e:f1ff:fee2:c4a2/64 scope link
        valid_lft forever preferred_lft forever
```

Open vSwitch (OVS)



Software-only OVS Implementation

High Latency, Low Bandwidth, CPU-Intensive

-  First packet in a flow
-  Fallback FRWD path

- Open vSwitch (OVS) is an open-source software switch that allows VMs and containers to communicate with each other and with the outside world.
- OVS software-only solutions are CPU-intensive, affecting system performance and providing poor network bandwidth and higher latency.

DPU OVS Bridges

- Two OVS bridges are created on the DPU by default after BFB installation.
- Default configuration allows default out-of-the box connection from the host to the DPU and from the host to the outside world.
- Default OVS bridges:
 - **ovsbr1** – allows traffic from the host to the DPU and from the host to the outside world via the first network port.
 - **ovsbr2** – allows traffic from the host to the DPU and from the host to the outside world via the second network port
- Default OVS bridge configuration is customizable.
- Additional bridges can be created.
- Different bridge methods can be used instead of OVS (e.g., DPDK).

DPU OVS Default Configuration

- Run the following command on the DPU to display OVS configuration:

- Bridge **ovsbr1** bridges the representors of the first network port, PF0:

- **p0** – uplink representor
- **pf0hpf** – host-side representor
- **en3f0pf0sf0** – SF representor

- Bridge **ovsbr2** bridges the representors of the second network port, PF1:

- **p1** – uplink representor
- **pf1hpf** – host side representor
- **en3f1pf1sf0** – SF representor

dpu # **ovs-vsctl show**

ea42c6e8-83e9-4449-87f6-
8a1969e1f0ea

Bridge ovsbr1

Port p0

Interface p0

Port pf0hpf

Interface pf0hpf

Port ovsbr1

Interface ovsbr1

type: internal

Port en3f0pf0sf0

Interface en3f0pf0sf0

Bridge ovsbr2

Port p1

Interface p1

Port ovsbr2

Interface ovsbr2

type: internal

Port pf1hpf

Interface pf1hpf

Port en3f1pf1sf0

Interface en3f1pf1sf0

ovs_version: "2.17.2-9e69ff4"

Represor Interfaces

- As noted in previous slides, the DPU OS contains represor interfaces when in ethernet mode for the host OS physical interfaces and virtual function interfaces
- These represor interfaces serve as a tunnel to pass traffic to represented host physical or virtual interface
- It is important to verify OVS (open virtual switch) on the DPU OS side is set correctly to act as the switch/bridge to pass traffic to the host interfaces as needed
- "ovs-vsctl show" can be ran from the DPU OS side to verify this
- If an interface is inside an OVS bridge, the OVS bridge owns the interface
- Assigning an IP to an interface in an OVS bridge will not work. You must either remove that interface from the OVS bridge if possible, or assign an IP to the OVS bridge itself

<https://docs.openvswitch.org/en/latest/faq/issues/>

DPUP Flow Steering and Bond/LAG

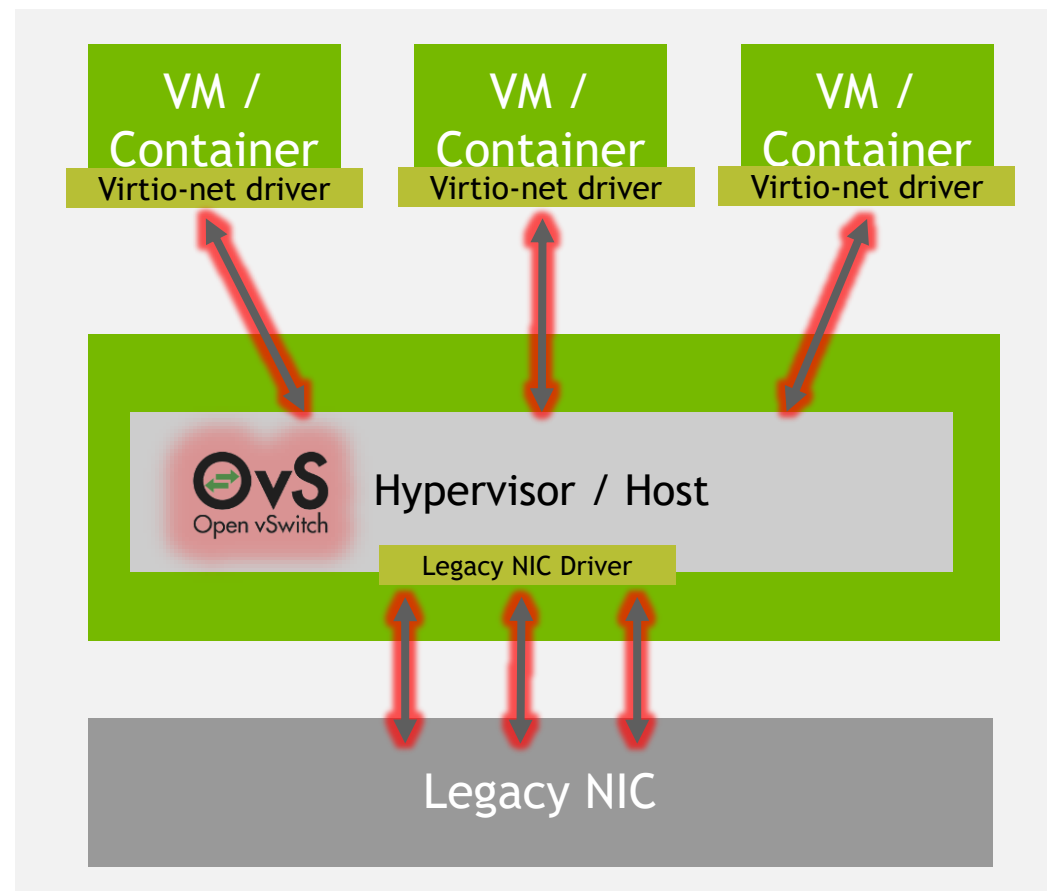
SWS Usage Scenarios

Cloud and GTW

应用场景:

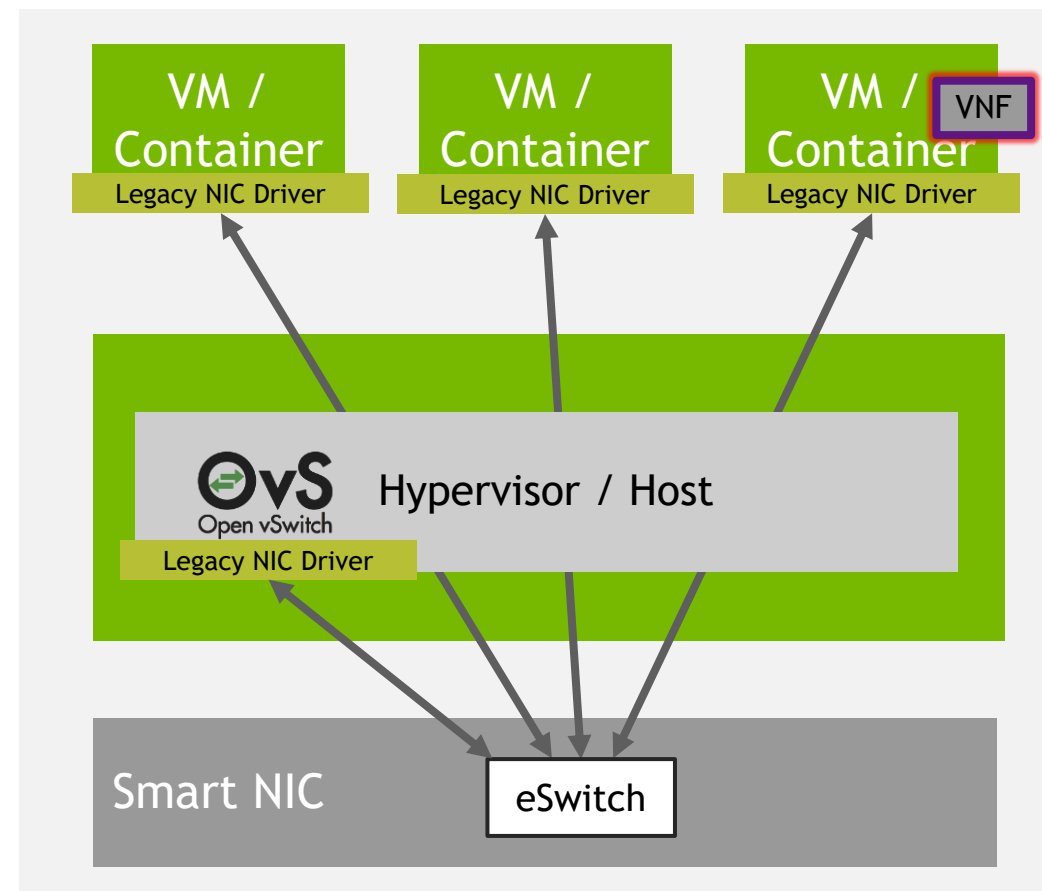
- 云虚拟化和和管理, DPDK, ovs-kernel, ovs-dpdk, 加速流卸载
- 云网关

Para-virtualization



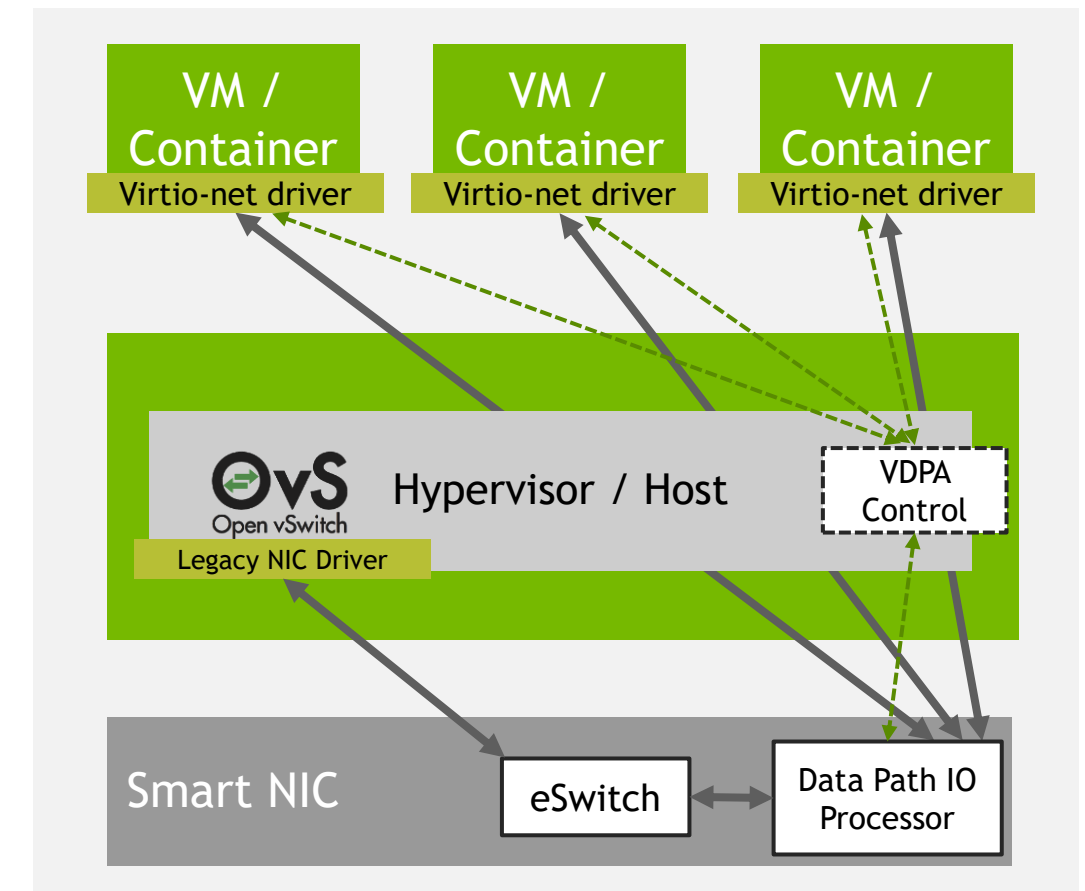
None accelerated

Single Root IO Virtualization (SR-IOV)



Accelerated

Virtual Datapath Acceleration (VDPA)



Accelerated

Flexible steering

Supported Actions (sample)

Forwarding decision

- Drop / allow
- Mirroring & monitor queue

Counter set

- About 8M counters

VLAN push / pop

- Packets coming from the host / network (respectively)
- Up to 2 VLANs

MPLS push / pop

- Packets coming from the host / network (respectively)
- Up to 5 labels

Hairpin

- Forward the packet back to the network

L2/L3 tunnel encapsulation / de-capsulation

- Packets coming from the host / network (respectively)

Header rewrite

- Update a specific header field
- Copy one header field to another
- Add a value to a specific header field

Metadata


- SW \leftrightarrow NIC

Connection Tracking

Encryption - IPsec & TLS


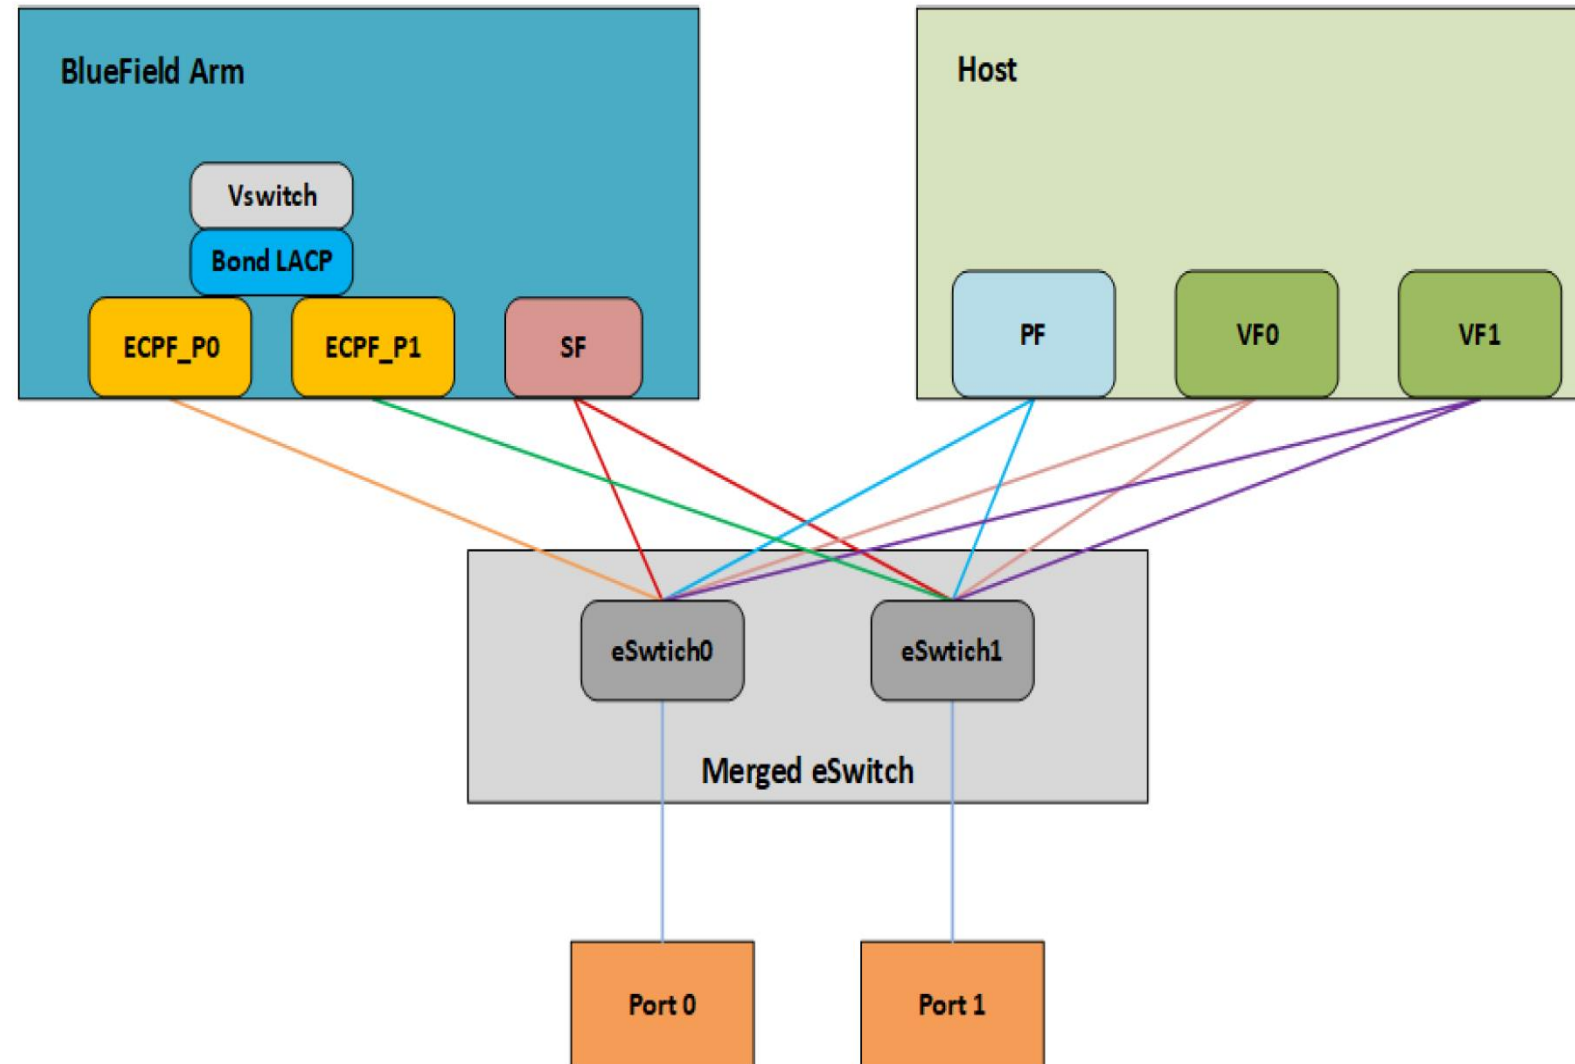
Generic header Push/Pop

Data policing/ Meters at scale (millions)



Link Aggregation Configure

What's LAG



- Combine two ports e-switch into one, no need to do bond on host
- Host devices (PF/VF/SF, virtio-xxx) have LAG properties for bandwidth and fail-over on uplink ports
- Only supported in DPU mode

Link Aggregation Configure

LAG Mode Configure

1. Firmware set

```
mlxconfig -d /dev/mst/mt41686_pciconf0 s  
LAG_RESOURCE_ALLOCATION=1  
  
mlxconfig -d /dev/mst/mt41686_pciconf0 s  
HIDE_PORT2_PF=True NUM_OF_PF=1
```

2. Hash mode

Add in /etc/mellanox/mlnx-bf.conf for LAG configuration

```
LAG_HASH_MODE="yes"
```

3. Delete the existing OVS

```
ovs-vsctl del-port ovsbr1 en3f0pf0sf0  
ovs-vsctl del-port ovsbr1 p0  
ovs-vsctl del-port ovsbr1 pf0hpf  
ovs-vsctl del-port ovsbr2 p1  
ovs-vsctl del-port ovsbr2 en3f1pf1sf0  
ovs-vsctl del-port ovsbr2 pf1hpf
```

4. LACP bonding configure

```
root@demo1:~# cat /etc/netplan/70-mlnx.yaml  
network:  
  renderer: networkd  
  ethernets:  
    p0:  
      dhcp4: no  
    p1:  
      dhcp4: no  
  bonds:  
    bond0:  
      interfaces: [p0,p1]  
      dhcp4: no  
      parameters:  
        mode: 802.3ad  
        transmit-hash-policy: layer3+4  
        mii-monitor-interval: 100  
  version: 2
```

5. Power cycle

```
ipmitool power cycle
```


Link Aggregation Configure

verify and debug

5. Check bonding device is OK

```
[dpu]
root@localhost:~# ibdev2netdev
mlx5_0 port 1 ==> enp3s0f0s0 (Up)
mlx5_1 port 1 ==> enp3s0f1s0 (Up)
mlx5_bond_0 port 1 ==> en3f0pf0sf0 (Up)
```

6. Add into OVS

```
[dpu]
ovs-vsctl add-br bf-lag
ovs-vsctl add-port bf-lag bond0
ovs-vsctl add-port bf-lag pf0hpf
ovs-vsctl add-port bf-lag en3f1pf1sf0
...
```

**** should not add p0,p1**

7. Ovs show

```
root@localhost:~# ovs-vsctl show
8a6decf9-5f0f-4096-aa50-dcf0f822ee5f
    Bridge bf-lag
        Port pf0hpf
            Interface pf0hpf
        Port en3f0pf0sf0
            Interface en3f0pf0sf0
        Port en3f1pf1sf0
            Interface en3f1pf1sf0
        Port bond0
            Interface bond0
        Port bf-lag
            Interface bf-lag
                type: internal
    ovs_version: "2.17.8-3feee121f"
```

8. Run rdma perftest on Host PF

```
[Host]
ib_write_bw -d mlx5_0 -b -p 50001 -q 4

[Peer]
ib_write_bw -d mlx5_0 1.1.1.1 -p 50001 -q 4 -b --report_gbits --run_infinately
```

```
[Dpu]
mlx_perf -i p0 -c 1 | grep x_bytes_phy && mlx_perf -i p1 -c 1 | grep
x_bytes_phy
    rx_bytes_phy: 25,829,049,110 Bps    = 206,632.39 Mbps
    rx_bytes_phy: 25,913,204,164 Bps    = 207,305.63 Mbps
```

9. Run iperf on Host PF

```
[Host]
numactl -N 0 -m 0 iperf -s -p 50001

[Peer]
numactl -N 0 -m 0 iperf -c 1.1.1.1 -t 1000 -i 2 -p 50001 -P 16
```

```
[Dpu]
mlx_perf -i ens2f0np0 -c 1 | grep rx_bytes_phy && mlx_perf -i ens5f1np1
-c 1 | grep rx_bytes_phy
    rx_bytes_phy: 22,170,418,304 Bps    = 177,363.34 Mbps
    rx_bytes_phy: 22,623,055,644 Bps    = 180,984.44 Mbps
```

Other ways to config Bond/LAG

■ Bond config for CentOS

/etc/sysconfig/network-scripts/ifcfg-bond0

```
DEVICE=bond0
NAME=bond0
TYPE=bond
BONDING_MASTER=yes
IPADDR=192.168.99.2
PREFIX=24
ONBOOT=yes
BOOTPROTO=none
BONDING_OPTS="miimon=100 mode=802.3ad lacp_rate=1
xmit_hash_policy=layer3+4"
```

/etc/sysconfig/network-scripts/ifcfg-ens1f0

```
DEVICE=p0
NAME=p0
TYPE=Ethernet
BOOTPROTO=none
ONBOOT=yes
MASTER=bond0
SLAVE=yes
```

/etc/sysconfig/network-scripts/ifcfg-ens1f1

```
DEVICE=p1
NAME= p1
TYPE=Ethernet
BOOTPROTO=none
ONBOOT=yes
MASTER=bond0
SLAVE=yes
```

■ Manual scripts

```
Modprobe -r bonding
```

```
modprobe bonding miimon=100 mode=4
echo +bond0 >/sys/class/net/bonding_masters
echo "layer3+4" >
/sys/class/net/bond0/bonding/xmit_hash_policy
ifconfig p0 up
ifconfig p1 up
ifconfig bond0 up
ifconfig p0 mtu 9508
ifconfig p1 mtu 9508
ifconfig bond0 mtu 9508
ifenslave bond0 p0 p1
```

```
cat /proc/net/bonding/bond0
dmesg | grep "lag map"
```

```
root@localhost:~# dmesg | grep lag
[ 21.967200] mlx5_core 0000:03:00.0: lag map active ports: 1, 2
[ 32.087494] device bf-lag entered promiscuous mode
```

```
root@localhost:~# ibdev2netdev
mlx5_0 port 1 ==> enp3s0f0s0 (Up)
mlx5_1 port 1 ==> enp3s0f1s0 (Up)
mlx5_bond_0 port 1 ==> en3f0pf0sf0 (Up)
mlx5_bond_0 port 10 ==> en3f0pf0sf0 (Up)
mlx5_bond_0 port 100 ==> en3f0pf0sf0 (Up)
mlx5_bond_0 port 101 ==> en3f0pf0sf0 (Up)
```

Bond/LAG OVS config example

■ ovs-kernel

```
ovs-vsctl set Open_vSwitch . Other_config:hw-offload=true
ovs-vsctl add-br ovsbr
ovs-vsctl add-port ovsbr bond1
ovs-vsctl add-port ovsbr pf0hpf
ovs-vsctl add-port ovsbr en3f0pf0sf2000
ovs-vsctl add-port ovsbr en3f0pf0sf2001
ovs-vsctl add-port ovsbr en3f0pf0sf2002
```

■ ovs-dpdk

```
ovs-vsctl clear o . Other_config
ovs-vsctl add-br ovsbr -- set bridge ovsbr datapath_type=netdev
ovs-vsctl --no-wait set o . other_config:hw-offload=true other_config:max-idle=30000
ovs-vsctl --no-wait set o . other_config:dpdk-extra="-w 0000:03:00.0,representor=pf0sf[2000-2030,65535],dv_xmeta_en=1,sys_mem_en=1 mtu_request=9000"
ovs-vsctl --no-wait set o . other_config:dpdk-init=true
```

```
ovs-vsctl add-port ovsbr bond0 -- set Interface bond0 type=dpdk options:dpdk-  
devargs="0000:03:00.0,dv_xmeta_en=1,sys_mem_en=1" mtu_request=9000
```

```
ovs-vsctl add-port ovsbr pf0hpf -- set Interface pf0hpf type=dpdk options:dpdk-  
devargs="0000:03:00.0,representor=sf65535,dv_xmeta_en=1,sys_mem_en=1" mtu_request=9000
```

```
ovs-vsctl add-port ovsbr pf0sf2000 -- set Interface pf0sf2000 type=dpdk options:dpdk-  
devargs="0000:03:00.0,representor=pf0sf2000,dv_xmeta_en=1,sys_mem_en=1" mtu_request=9000
```

BF LAG Take aways

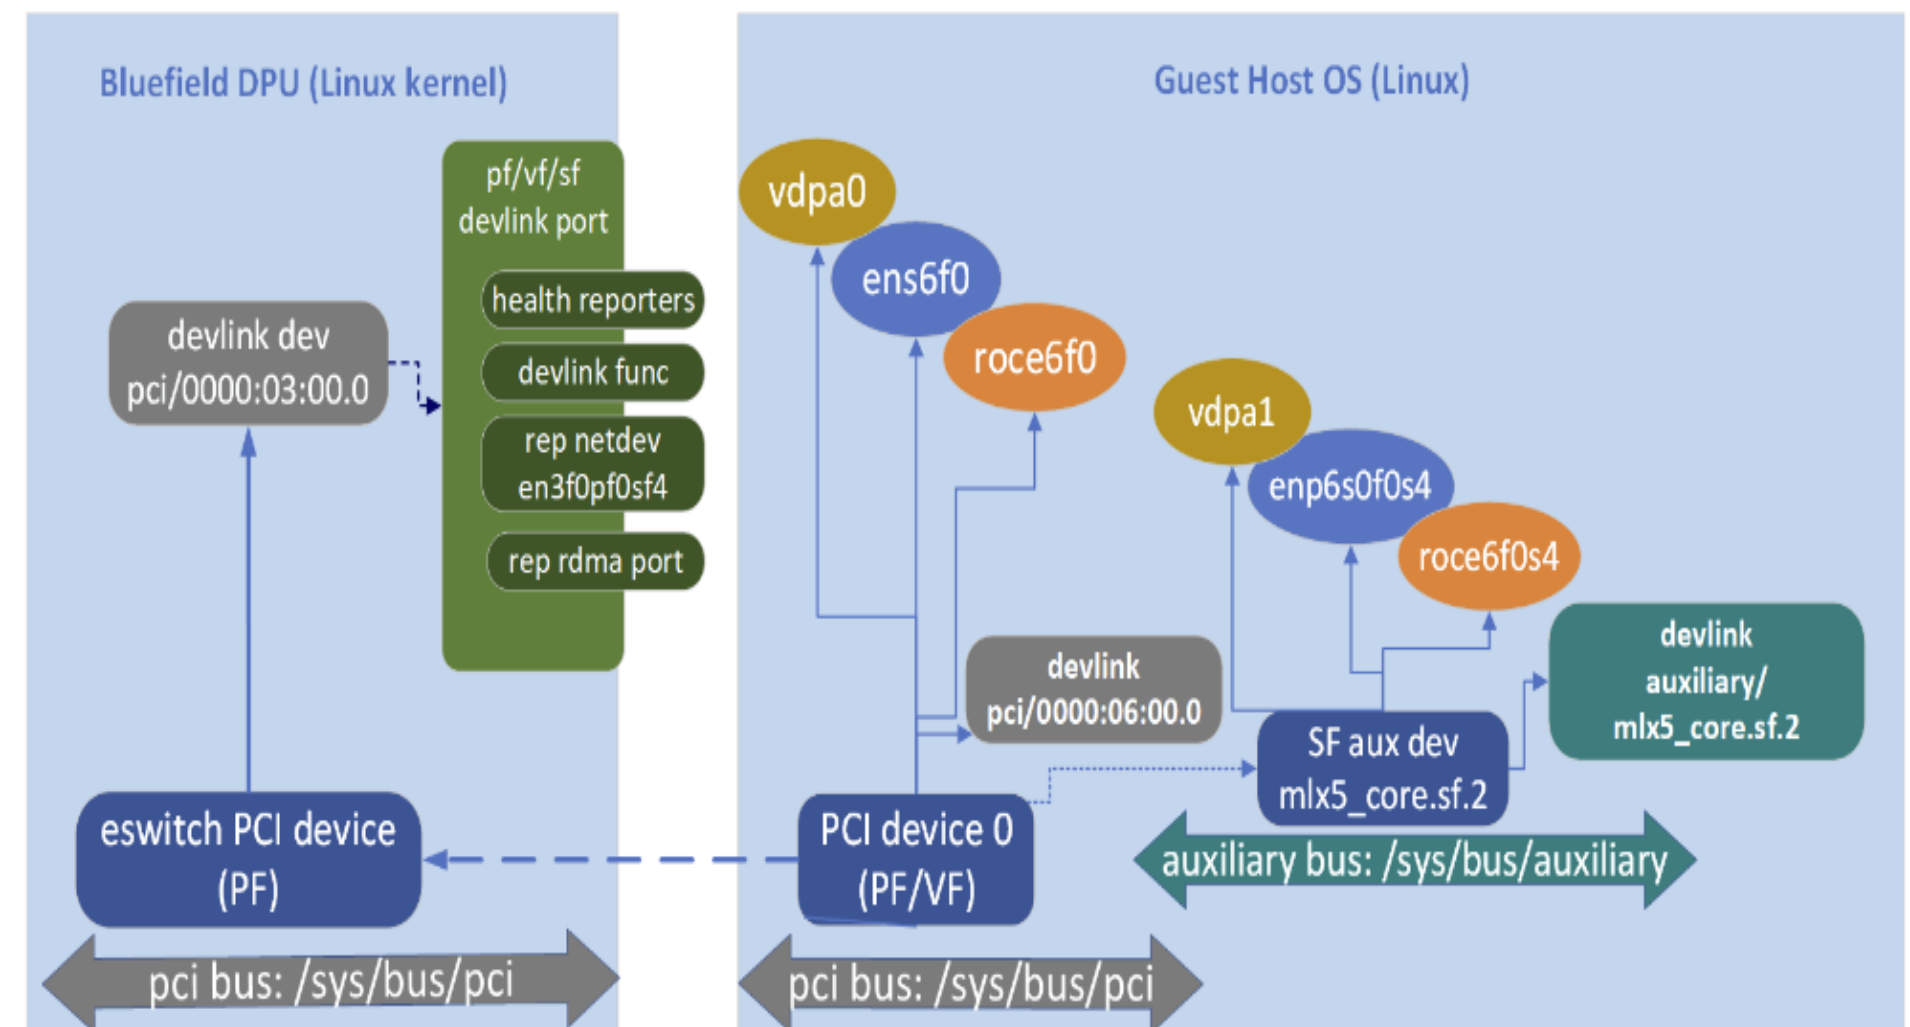
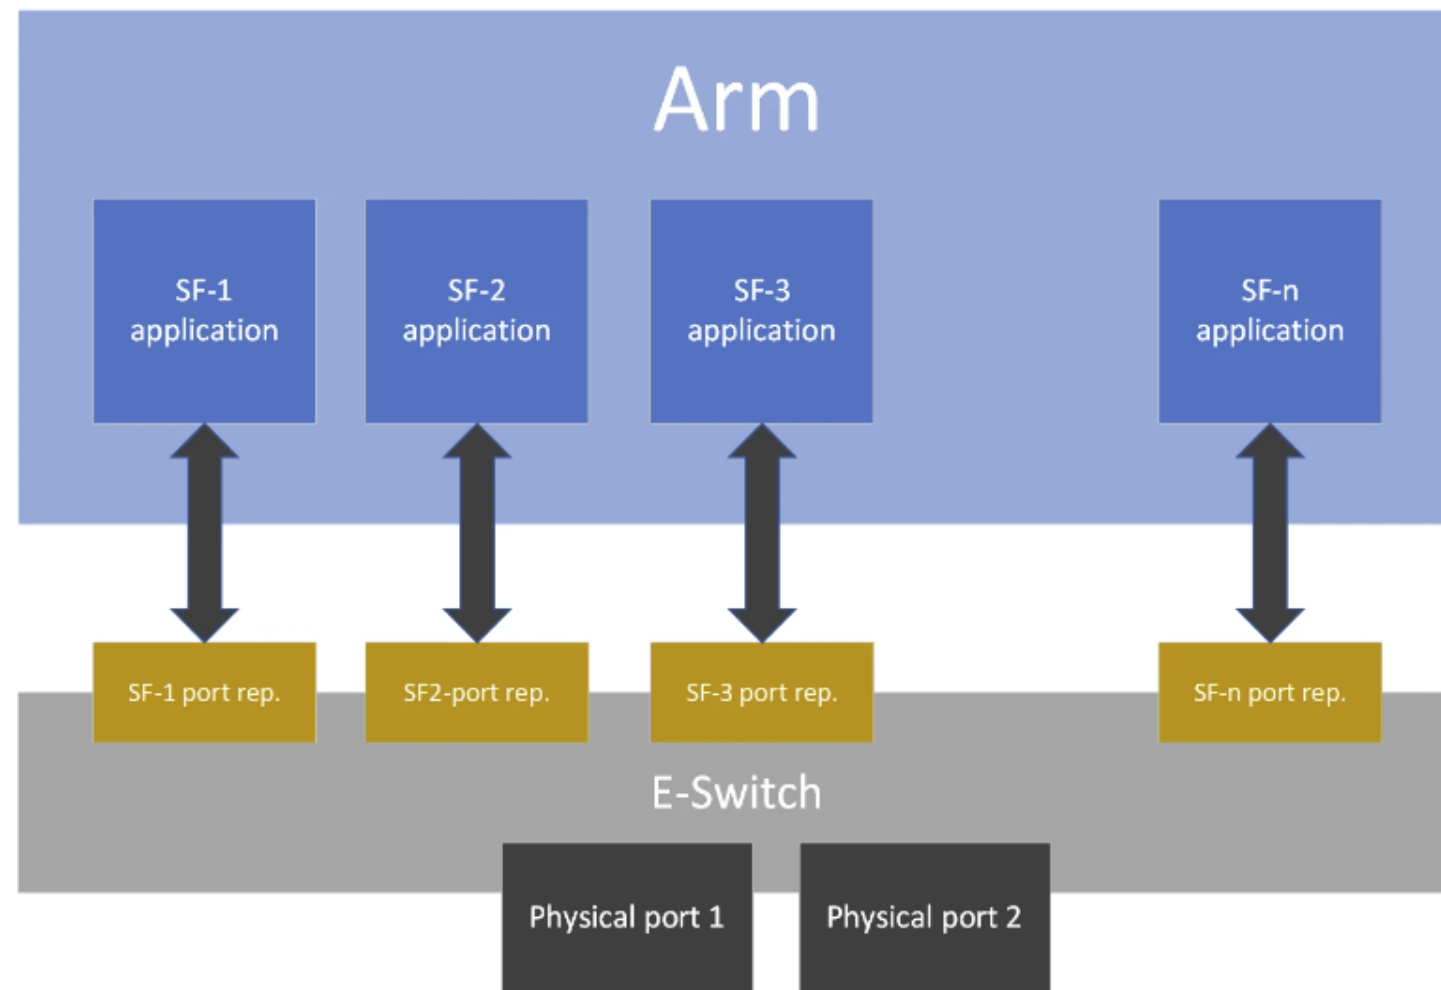
1. Bond/LAG 是否成功看mlx5_bond_0是否生成, bond运行需要在没有创建SRIOV或者OVS stop之前
2. 创建LAG之后, host上的PF, VF, DPU上的SF都具备了LAG属性, 也就是能从两个uplink口收发数据
3. LAG的tx缺省是hash模式 (5元组) , 可以选配queue_affinity模式
4. LAG的OVS中不需要向bridge中加入p0/p1, 只需要加入bond master
5. DPU LAG 抓包
 - `tcpdump -i mlx5_bond_0 # need upgrade libpcap >= 1.10`
 - `ovs-tcpdump -i bond0`



Scalable Functions (SFs)

DPU Scalable Functions (SF)

- Scalable functions (SF) 功能类似于VF，提供的数量更多elastic networking interface (ENI)
- 从PF上衍生的轻量级的功能设备，有独立的硬件资源例如队列
- 不需要SRIOV支持，能和VF共存，并且能在DPU的ARM OS上或者HOST上创建
- 有eth device和对应的RDMA设备，DPU内对应representer用于OVS的flows



DPU Scalable Functions (SF)

• Create SF for DPU internal

FW config

```
[DPU] mlxconfig -d 0000:03:00.0 s PF_BAR2_ENABLE=0 PER_PF_NUM_SF=1 PF_TOTAL_SF=236 PF_SF_BAR_SIZE=10
```

Create the SF

```
[DPU] /opt/mellanox/iproute2/sbin/mlxdevm port add pci/0000:03:00.0 flavour pcisf pfnum 0 sfnun 5
```

Configure the SF

```
[DPU] /opt/mellanox/iproute2/sbin/mlxdevm port function set pci/0000:03:00.0/163874 hw_addr 02:25:f2:8d:a2:5c trust on state active
```

Deploy the SF

```
[DPU] echo mlx5_core.sf.3 > /sys/bus/auxiliary/drivers/mlx5_core.sf_cfg/unbind
```

```
[DPU] echo mlx5_core.sf.3 > /sys/bus/auxiliary/drivers/mlx5_core.sf/bind
```

Delete SF

```
[DPU] /opt/mellanox/iproute2/sbin/mlxdevm port function set pci/0000:03:00.0/163874 state inactive
```

```
[DPU] /opt/mellanox/iproute2/sbin/mlxdevm port del pci/0000:03:00.0/163874
```

```
root@localhost:~# mlx-sf -a show
SF Index: pci/0000:03:00.0/163872
Parent PCI dev: 0000:03:00.0
Representor netdev: en3f0pf0sf0
Function HWADDR: 02:27:c2:6d:31:63
Function trust: off
Function roce: true
Function eswitch: NA
Auxiliary device: mlx5_core.sf.1
netdev: enp3s0f0s0
RDMA dev: mlx5_1
```

```
root@amdgen5sz:~# ibdev2netdev
mlx5_0 port 1 ==> ens3f0np0 (Up)
mlx5_1 port 1 ==> ens3f1np1 (Up)
mlx5_2 port 1 ==> ens2f0np0 (Up)
mlx5_3 port 1 ==> enp2s0f0s88 (Down)
```

• Create SF for DPU host external

FW config (DPU and Host)

```
[DPU] mlxconfig -d 0000:03:00.0 -y s NUM_PF_MSIX_VALID=0
```

```
[DPU] mlxconfig -d 0000:03:00.0 -y s PF_NUM_PF_MSIX_VALID=1
```

```
[DPU] mlxconfig -d 0000:03:00.0 -y s PF_TOTAL_SF=32 PF_NUM_PF_MSIX=128 PF_BAR2_ENABLE=0 PER_PF_NUM_SF=1 PF_SF_BAR_SIZE=10
```

```
[Host] mlxconfig -d 0000:02:00.0 -y s PF_TOTAL_SF=32 PF_NUM_PF_MSIX=128 PF_BAR2_ENABLE=0 PER_PF_NUM_SF=1 PF_SF_BAR_SIZE=10
```

Create and active SF

```
[DPU] /opt/mellanox/iproute2/sbin/mlxdevm port add pci/0000:03:00.0 flavour pcisf pfnum 0 sfnun 88 controller 1
```

```
[DPU] /opt/mellanox/iproute2/sbin/mlxdevm port function set pci/0000:03:00.0/163840 hw_addr 02:25:f2:8d:aa:4c state active
```

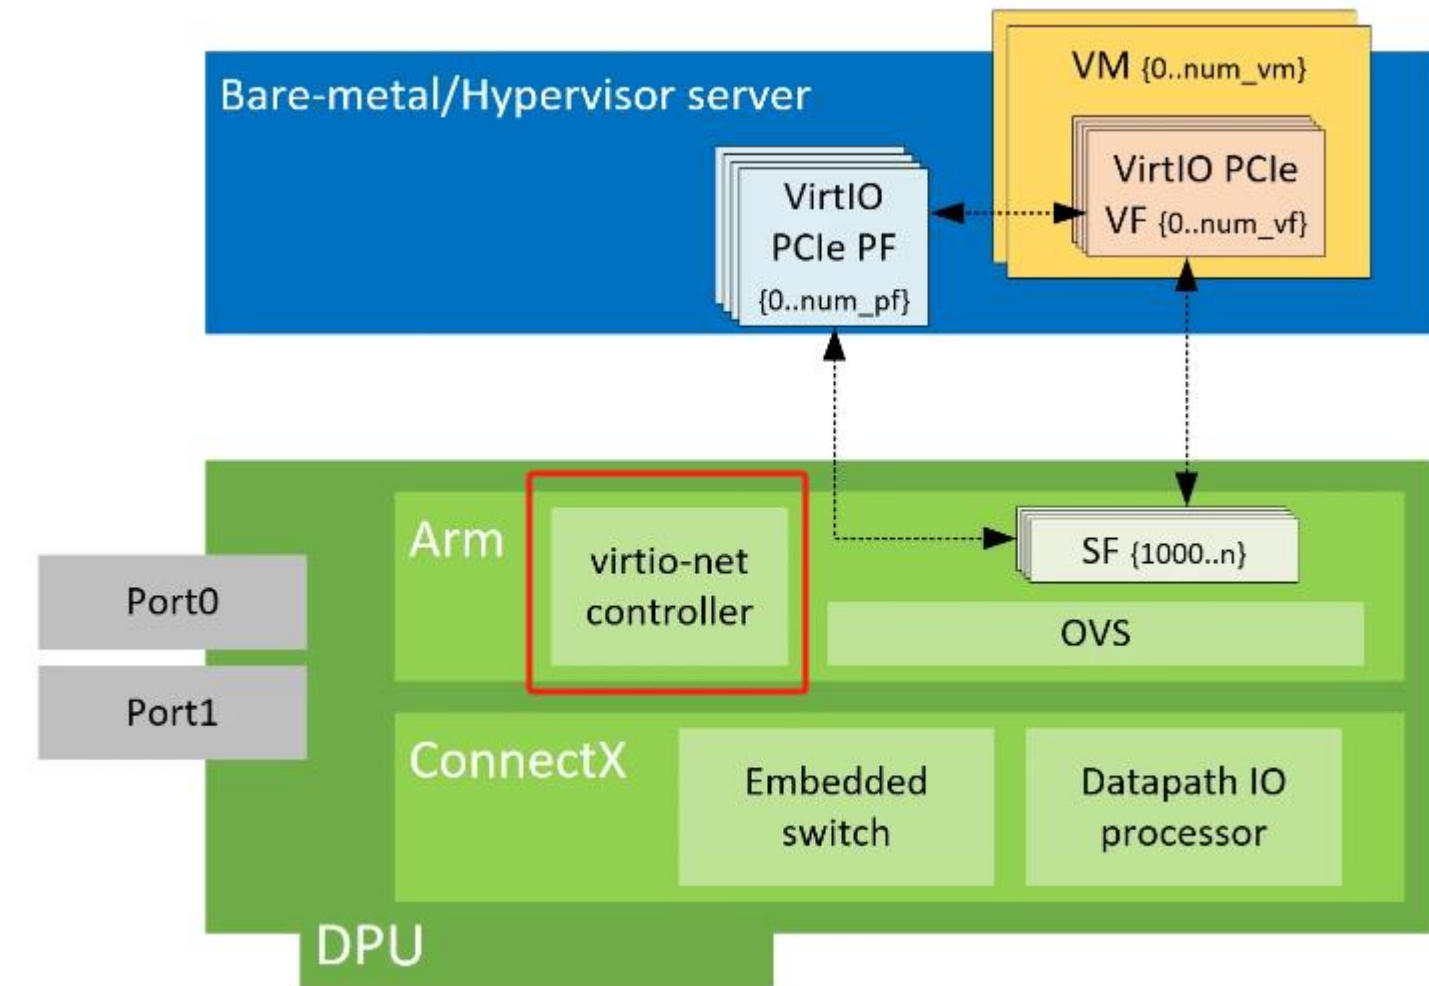
Virtio-net Hotplug for Bare-metal

Virtio-net overview

1. Service **virtio-net-controller** manager virtio_net devices, must have initialization conf
2. Communicated through rpc, commands through **virtnet**
3. OVS or DPDK through **rte_flow** or **TC** to control all virtio-device flow rules for offload or data path
4. Each virtio_net device will map to one **SF for rep** for flow rule inserting interface or other such as QoS interface.

Note:

1. Hot-plug 动态virtio必须是基于PF，硬件限制最多31.
2. Static 可以是PF，也可以是VF，VF需要从host上创建
3. Static virtio_net 必须是先配置好backen，再加载host上的 virtio_pci driver，否者会导致host kernel被挂住
4. OVS 或者DPDK或者TC必须配置好流规则后virtio才能通信



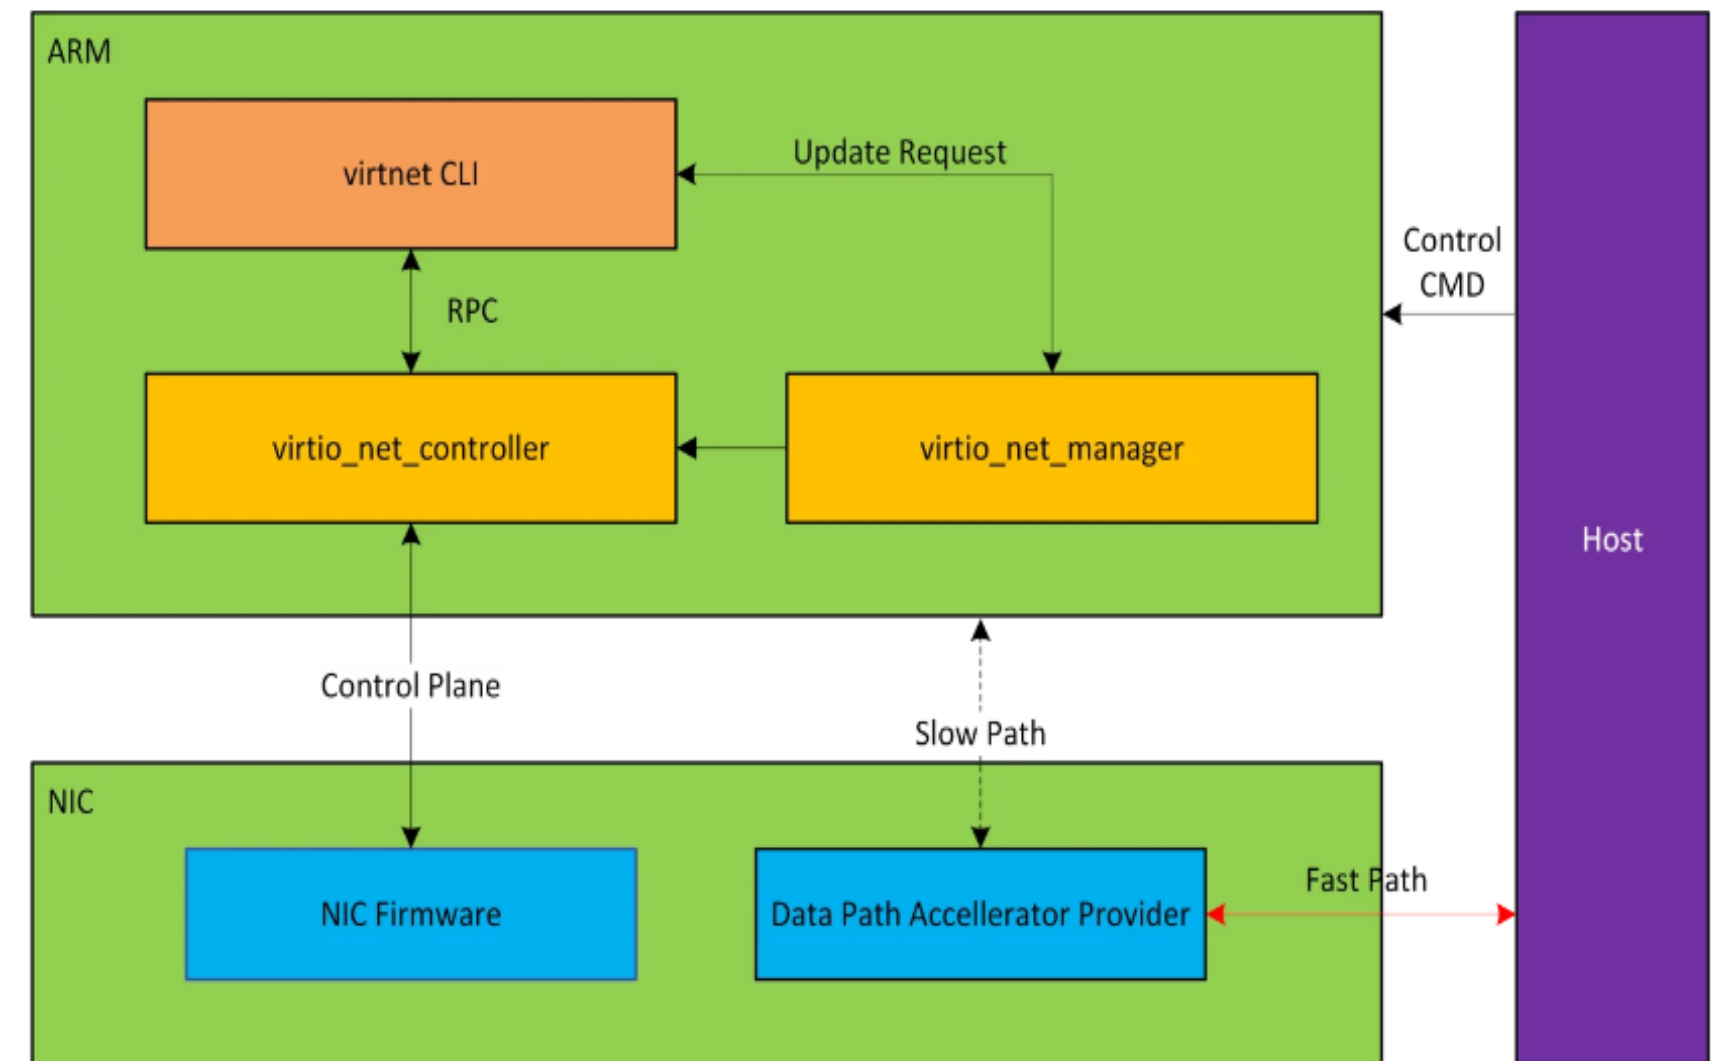
NVIDIA BlueField Virtio-net v1.9.0

<https://docs.nvidia.com/networking/display/bluefieldvirtionetv190>

Virtio-net how it works

Running Mechanism in DPU

1. virtio-net-controller service works as control center
2. virtnet cli is communicated with cmd “**virtnet**” through RPC
3. Virtio devices control messages exchange through FW path
4. RPC API can be integrated with upper applications
5. virt_net_manager is working for live update
6. Support fast recovery in case of crash
7. Static virtio-net is created from DPU power up
8. Hot-plug virtio-net is created dynamically
9. virtio-net capabilities : max number, max queues, features are limited by HW or FW setting.
10. Support virtio spec 1.x version (limited for 0.95)



Static PF	Hot-plug PF	VF
31	31	1008

Virtio-net hotplug – 1

1. FW config (on DPU)

```
mlxconfig -d 03:00.0 -y reset
mlxconfig -d 03:00.0 -y set \
    INTERNAL_CPU_MODEL=1 SRIOV_EN=0 \
    NUM_OF_VFS=0 NUM_OF_PF=0 \
    PCI_SWITCH_EMULATION_ENABLE=1 \
    PCI_SWITCH_EMULATION_NUM_PORT=32 \
    VIRTIO_NET_EMULATION_ENABLE=1 \
    VIRTIO_NET_EMULATION_NUM_VF=0 \
    VIRTIO_NET_EMULATION_NUM_PF=0 \
    VIRTIO_NET_EMULATION_NUM_MSIX=64 \
    LAG_RESOURCE_ALLOCATION=1 \
    PER_PF_NUM_SF=1 PF_TOTAL_SF=64 PF_SF_BAR_SIZE=10

mlxconfig -d 03:00.1 -y set \
    PER_PF_NUM_SF=1 PF_TOTAL_SF=2 PF_SF_BAR_SIZE=10
```

2. Enable virtio-net-controller service

```
systemctl enable virtio-net-controller
journalctl -u virtio-net-controller -n 100 -f
```

3. Edit /opt/mellanox/mlnx_virtnet/virtnet.conf

```
{
    "ib_dev_lag": "mlx5_bond_0",
    "ib_dev_for_static_pf": "mlx5_bond_0",
    "is_lag": 1,
    "recovery": 1,
    "sf_pool_percent": 0,
    "sf_pool_force_destroy": 0
}
```

4. Power cycle

```
[host]
Add "pci=realloc" into kernel boot grub options
```

```
[host]
ipmitool power cycle
```

Virtio-net usage :

/opt/mellanox/mlnx_virtnet/README.md

Virtio-net hotplug – 2

5. Check virtio-net-controller service

```
root@localhost:~# systemctl status virtio-net-controller
● virtio-net-controller.service - Nvidia VirtIO Net Controller Daemon
   Loaded: loaded (/etc/systemd/system/virtio-net-controller.service; enabled; vendor preset: enabled)
   Active: active (running) since Fri 2023-09-15 10:04:21 UTC; 1min 35s ago
     Docs: file:/opt/mellanox/mlnx_virtnet/README.md
   Process: 3036 ExecStartPost=/bin/sleep 3 (code=exited, status=0/SUCCESS)
    Main PID: 3035 (virtio_net_manager)
       Tasks: 29 (limit: 18994)
      Memory: 3.5M
         CPU: 6.968s
    CGroup: /system.slice/virtio-net-controller.service
            └─3035 /usr/sbin/virtio_net_manager
              └─3289 virtio_net_controller

Sep 15 10:04:21 localhost.localdomain systemd[1]: Started Nvidia VirtIO Net Controller Daemon.
Sep 15 10:04:22 localhost.localdomain virtio_net_manager[3289]: snap_channel.c:114 INFO registered migration
Sep 15 10:04:22 localhost.localdomain virtio-net-controller[3289]: [INFO] virtnet_controller.c:250:main: Sta
Sep 15 10:04:22 localhost.localdomain virtio-net-controller[3289]: [INFO] virtnet_controller.c:257:main: Cor
Sep 15 10:04:22 localhost.localdomain virtio-net-controller[3289]: [INFO] virtnet_controller.c:263:main: Cor
Sep 15 10:04:22 localhost.localdomain virtio-net-controller[3289]: [INFO] virtnet_controller.c:272:main: Cor
Sep 15 10:04:22 localhost.localdomain virtio-net-controller[3289]: [INFO] virtnet_providers.c:532:provider_l
Sep 15 10:04:22 localhost.localdomain virtio-net-controller[3289]: [DEBUG] virtnet_util.c:368:virtnet_ibdev_t
Sep 15 10:04:22 localhost.localdomain virtio-net-controller[3289]: [DEBUG] virtnet_util.c:395:virtnet_get_pc
Sep 15 10:04:22 localhost.localdomain virtio-net-controller[3289]: [INFO] virtnet_sf.c:2339:virtnet_sf_load:
root@localhost:~#
root@localhost:~# virtnet -v
v1.6.14
root@localhost:~# virtnet list
```

```
root@localhost:~# virtnet list
{
  "controller": {
    "emulation_manager": "mlx5_bond_0",
    "max_hotplug_devices": "15",
    "max_virt_net_devices": "15",
    "max_virt_queues": "64",
    "max_tunnel_descriptors": "6",
    "supported_features": {
      "value": "0x80000333004f982b",
      "0": "VIRTIO_NET_F_CSUM",
      "1": "VIRTIO_NET_F_GUEST_CSUM",
      "3": "VIRTIO_NET_F_MTU",
      "5": "VIRTIO_NET_F_MAC",
      "11": "VIRTIO_NET_F_HOST_TS04",
      "12": "VIRTIO_NET_F_HOST_TS06",
      "15": "VIRTIO_F_MRG_RX_BUFFER",
      "16": "VIRTIO_NET_F_STATUS",
      "17": "VIRTIO_NET_F_CTRL_VQ",
      "18": "VIRTIO_NET_F_CTRL_RX",
      "19": "VIRTIO_NET_F_CTRL_VLAN",
      "22": "VIRTIO_NET_F_MQ",
      "32": "VIRTIO_F_VERSION_1",
      "33": "VIRTIO_F_IOMMU_PLATFORM",
      "36": "VIRTIO_F_ORDER_PLATFORM",
      "37": "VIRTIO_F_SR_IOV",
      "40": "VIRTIO_F_RING_RESET",
      "41": "VIRTIO_F_ADMIN_VQ",
      "63": "VIRTIO_NET_F_SPEED_DUPLEX"
    },
    "supported_virt_queue_types": {
      "value": "0x1",
      "0": "SPLIT"
    },
    "supported_event_modes": {
      "value": "0x5",
      "0": "NO_MSIX_MODE",
      "2": "MSIX_MODE"
    }
  },
  "devices": []
}
```

virtio-net hotplug – 3

6. Create virtio_net hotplug devices

```
virtnet hotplug -i mlx5_bond_0 -f 0x80000 -m 00:11:22:33:44:11 -t 1500 -n 33 -s 1024
virtnet hotplug -i mlx5_bond_0 -f 0x80000 -m 00:11:22:33:44:22 -t 1500 -n 33 -s 1024
virtnet hotplug -i mlx5_bond_0 -f 0x80000 -m 00:11:22:33:44:33 -t 1500 -n 33 -s 1024
virtnet hotplug -i mlx5_bond_0 -f 0x80000 -m 00:11:22:33:44:44 -t 1500 -n 33 -s 1024
.....
virtnet list
```

```
virtnet hotplug -i mlx5_bond_0 -f 0x80000 -m 00:11:22:33:44:88 -t 1500 -n 33 -s 1024
{'ib_device': 'mlx5_bond_0', 'mac': '00:11:22:33:44:11', 'mtu': 1500, 'max_queues': 33}
{
  "bdf": "b2:00.0",
  "vuid": "VNETS1D0F0",
  "id": 0,
  "transitional": 0,
  "sf_rep_net_device": "en3f0pf0sf2000",
  "mac": "00:11:22:33:44:11",
  "errno": 0,
  "errstr": "Success"
}
{'ib_device': 'mlx5_bond_0', 'mac': '00:11:22:33:44:22', 'mtu': 1500, 'max_queues': 33}
{
  "bdf": "b3:00.0",
  "vuid": "VNETS2D0F0",
  "id": 1,
  "transitional": 0,
  "sf_rep_net_device": "en3f0pf0sf2001",
  "mac": "00:11:22:33:44:22",
  "errno": 0,
  "errstr": "Success"
}
{'ib_device': 'mlx5_bond_0', 'mac': '00:11:22:33:44:33', 'mtu': 1500, 'max_queues': 33}
{
  "bdf": "b4:00.0",
  "vuid": "VNETS3D0F0",
  "id": 2,
  "transitional": 0,
  "sf_rep_net_device": "en3f0pf0sf2002",
  "mac": "00:11:22:33:44:33",
  "errno": 0,
  "errstr": "Success"
}
{'ib_device': 'mlx5_bond_0', 'mac': '00:11:22:33:44:44', 'mtu': 1500, 'max_queues': 33}
```

```
"devices": [
  {
    "pf_id": 0,
    "function_type": "hotplug PF",
    "transitional": 0,
    "vuid": "VNETS1D0F0",
    "bdf": "b2:00.0",
    "sf_num": 2000,
    "sf_parent_device": "mlx5_bond_0",
    "sf_parent_device_pci_addr": "0000:03:00.0",
    "sf_rep_net_device": "en3f0pf0sf2000",
    "sf_rep_net_ifindex": 13,
    "sf_rdma_device": "mlx5_0",
    "sf_vhca_id": "0x13",
    "msix_config_vector": "0x0",
    "num_msix": 34,
    "max_queues": 33,
    "max_queues_size": 1024,
    "net_mac": "00:11:22:33:44:11",
    "net_mtu": 1500
  },
  {
    "pf_id": 1,
    "function_type": "hotplug PF",
    "transitional": 0,
    "vuid": "VNETS2D0F0",
    "bdf": "b3:00.0",
    "sf_num": 2001,
    "sf_parent_device": "mlx5_bond_0",
    "sf_parent_device_pci_addr": "0000:03:00.0",
    "sf_rep_net_device": "en3f0pf0sf2001",
    "sf_rep_net_ifindex": 14,
    "sf_rdma_device": "mlx5_1",
    "sf_vhca_id": "0x14",
    "msix_config_vector": "0x0",
    "num_msix": 34,
    "max_queues": 33,
    "max_queues_size": 1024,
    "net_mac": "00:11:22:33:44:22",
    "net_mtu": 1500
  },
  {
    "pf_id": 2,
    "function_type": "hotplug PF",
    "transitional": 0,
    "vuid": "VNETS3D0F0",
    "bdf": "b4:00.0",
    "sf_num": 2002,
    "sf_parent_device": "mlx5_bond_0",
    "sf_parent_device_pci_addr": "0000:03:00.0",
    "sf_rep_net_device": "en3f0pf0sf2002",
    "sf_rep_net_ifindex": 15,
    "sf_rdma_device": "mlx5_2",
    "sf_vhca_id": "0x15",
    "msix_config_vector": "0x0",
    "num_msix": 34,
    "max_queues": 33,
    "max_queues_size": 1024,
    "net_mac": "00:11:22:33:44:33",
    "net_mtu": 1500
  },
  {
    "pf_id": 3,
    "function_type": "hotplug PF",
    "transitional": 0,
    "vuid": "VNETS4D0F0",
    "bdf": "b5:00.0",
    "sf_num": 2003,
    "sf_parent_device": "mlx5_bond_0",
    "sf_parent_device_pci_addr": "0000:03:00.0",
    "sf_rep_net_device": "en3f0pf0sf2003",
    "sf_rep_net_ifindex": 16,
    "sf_rdma_device": "mlx5_3",
    "sf_vhca_id": "0x16",
    "msix_config_vector": "0x0",
    "num_msix": 34,
    "max_queues": 33,
    "max_queues_size": 1024,
    "net_mac": "00:11:22:33:44:44",
    "net_mtu": 1500
  }
]
```

```
root@localhost:~# virtnet query -p 0
{'all': 0, 'pf': 0, 'dbg_stats': False, 'brief': False, 'latency_stats': False}
{
  "devices": [
    {
      "pf_id": 0,
      "transitional": 0,
      "vuid": "VNETS1D0F0",
      "pci_bdf": "b2:00.0",
      "pci_dev_id": "0x1041",
      "pci_vendor_id": "0x1af4",
      "pci_class_code": "0x20000",
      "pci_subsys_id": "0x1",
      "pci_subsys_vendor_id": "0x1af4",
      "pci_revision_id": "1",
      "pci_max_vfs": "0",
      "enabled_vfs": "0",
      "device_feature": {
        "value": "0x103004f182b",
        "0": "VIRTIO_NET_F_CSUM",
        "1": "VIRTIO_NET_F_GUEST_CSUM",
        "3": "VIRTIO_NET_F_MTU",
        "5": "VIRTIO_NET_F_MAC",
        "11": "VIRTIO_NET_F_HOST_TS04",
        "12": "VIRTIO_NET_F_HOST_TS06",
        "16": "VIRTIO_NET_F_STATUS",
        "17": "VIRTIO_NET_F_CTRL_VQ",
        "18": "VIRTIO_NET_F_CTRL_RX",
        "19": "VIRTIO_NET_F_CTRL_VLAN",
        "22": "VIRTIO_NET_F_MQ",
        "32": "VIRTIO_F_VERSION_1",
        "33": "VIRTIO_F_IOMMU_PLATFORM",
        "40": "VIRTIO_F_RING_RESET"
      },
      "driver_feature": {
        "value": "0x3004f182b",
        "0": "VIRTIO_NET_F_CSUM",
        "1": "VIRTIO_NET_F_GUEST_CSUM",
        "3": "VIRTIO_NET_F_MTU",
        "5": "VIRTIO_NET_F_MAC",
        "11": "VIRTIO_NET_F_HOST_TS04",
        "12": "VIRTIO_NET_F_HOST_TS06",
        "16": "VIRTIO_NET_F_STATUS",
        "17": "VIRTIO_NET_F_CTRL_VQ",
        "18": "VIRTIO_NET_F_CTRL_RX",
        "19": "VIRTIO_NET_F_CTRL_VLAN",
        "22": "VIRTIO_NET_F_MQ",
        "32": "VIRTIO_F_VERSION_1",
        "33": "VIRTIO_F_IOMMU_PLATFORM"
      },
      "status": {
        "value": "0xf",
        "0": "STATUS_ACKNOWLEDGE",
        "1": "STATUS_DRIVER"
      }
    },
    {
      "pf_id": 1,
      "transitional": 0,
      "vuid": "VNETS2D0F0",
      "pci_bdf": "b3:00.0",
      "pci_dev_id": "0x1042",
      "pci_vendor_id": "0x1af4",
      "pci_class_code": "0x20000",
      "pci_subsys_id": "0x1",
      "pci_subsys_vendor_id": "0x1af4",
      "pci_revision_id": "1",
      "pci_max_vfs": "0",
      "enabled_vfs": "0",
      "device_feature": {
        "value": "0x103004f182b",
        "0": "VIRTIO_NET_F_CSUM",
        "1": "VIRTIO_NET_F_GUEST_CSUM",
        "3": "VIRTIO_NET_F_MTU",
        "5": "VIRTIO_NET_F_MAC",
        "11": "VIRTIO_NET_F_HOST_TS04",
        "12": "VIRTIO_NET_F_HOST_TS06",
        "16": "VIRTIO_NET_F_STATUS",
        "17": "VIRTIO_NET_F_CTRL_VQ",
        "18": "VIRTIO_NET_F_CTRL_RX",
        "19": "VIRTIO_NET_F_CTRL_VLAN",
        "22": "VIRTIO_NET_F_MQ",
        "32": "VIRTIO_F_VERSION_1",
        "33": "VIRTIO_F_IOMMU_PLATFORM",
        "40": "VIRTIO_F_RING_RESET"
      },
      "driver_feature": {
        "value": "0x3004f182b",
        "0": "VIRTIO_NET_F_CSUM",
        "1": "VIRTIO_NET_F_GUEST_CSUM",
        "3": "VIRTIO_NET_F_MTU",
        "5": "VIRTIO_NET_F_MAC",
        "11": "VIRTIO_NET_F_HOST_TS04",
        "12": "VIRTIO_NET_F_HOST_TS06",
        "16": "VIRTIO_NET_F_STATUS",
        "17": "VIRTIO_NET_F_CTRL_VQ",
        "18": "VIRTIO_NET_F_CTRL_RX",
        "19": "VIRTIO_NET_F_CTRL_VLAN",
        "22": "VIRTIO_NET_F_MQ",
        "32": "VIRTIO_F_VERSION_1",
        "33": "VIRTIO_F_IOMMU_PLATFORM"
      },
      "status": {
        "value": "0xf",
        "0": "STATUS_ACKNOWLEDGE",
        "1": "STATUS_DRIVER"
      }
    },
    {
      "pf_id": 2,
      "transitional": 0,
      "vuid": "VNETS3D0F0",
      "pci_bdf": "b4:00.0",
      "pci_dev_id": "0x1043",
      "pci_vendor_id": "0x1af4",
      "pci_class_code": "0x20000",
      "pci_subsys_id": "0x1",
      "pci_subsys_vendor_id": "0x1af4",
      "pci_revision_id": "1",
      "pci_max_vfs": "0",
      "enabled_vfs": "0",
      "device_feature": {
        "value": "0x103004f182b",
        "0": "VIRTIO_NET_F_CSUM",
        "1": "VIRTIO_NET_F_GUEST_CSUM",
        "3": "VIRTIO_NET_F_MTU",
        "5": "VIRTIO_NET_F_MAC",
        "11": "VIRTIO_NET_F_HOST_TS04",
        "12": "VIRTIO_NET_F_HOST_TS06",
        "16": "VIRTIO_NET_F_STATUS",
        "17": "VIRTIO_NET_F_CTRL_VQ",
        "18": "VIRTIO_NET_F_CTRL_RX",
        "19": "VIRTIO_NET_F_CTRL_VLAN",
        "22": "VIRTIO_NET_F_MQ",
        "32": "VIRTIO_F_VERSION_1",
        "33": "VIRTIO_F_IOMMU_PLATFORM",
        "40": "VIRTIO_F_RING_RESET"
      },
      "driver_feature": {
        "value": "0x3004f182b",
        "0": "VIRTIO_NET_F_CSUM",
        "1": "VIRTIO_NET_F_GUEST_CSUM",
        "3": "VIRTIO_NET_F_MTU",
        "5": "VIRTIO_NET_F_MAC",
        "11": "VIRTIO_NET_F_HOST_TS04",
        "12": "VIRTIO_NET_F_HOST_TS06",
        "16": "VIRTIO_NET_F_STATUS",
        "17": "VIRTIO_NET_F_CTRL_VQ",
        "18": "VIRTIO_NET_F_CTRL_RX",
        "19": "VIRTIO_NET_F_CTRL_VLAN",
        "22": "VIRTIO_NET_F_MQ",
        "32": "VIRTIO_F_VERSION_1",
        "33": "VIRTIO_F_IOMMU_PLATFORM"
      },
      "status": {
        "value": "0xf",
        "0": "STATUS_ACKNOWLEDGE",
        "1": "STATUS_DRIVER"
      }
    },
    {
      "pf_id": 3,
      "transitional": 0,
      "vuid": "VNETS4D0F0",
      "pci_bdf": "b5:00.0",
      "pci_dev_id": "0x1044",
      "pci_vendor_id": "0x1af4",
      "pci_class_code": "0x20000",
      "pci_subsys_id": "0x1",
      "pci_subsys_vendor_id": "0x1af4",
      "pci_revision_id": "1",
      "pci_max_vfs": "0",
      "enabled_vfs": "0",
      "device_feature": {
        "value": "0x103004f182b",
        "0": "VIRTIO_NET_F_CSUM",
        "1": "VIRTIO_NET_F_GUEST_CSUM",
        "3": "VIRTIO_NET_F_MTU",
        "5": "VIRTIO_NET_F_MAC",
        "11": "VIRTIO_NET_F_HOST_TS04",
        "12": "VIRTIO_NET_F_HOST_TS06",
        "16": "VIRTIO_NET_F_STATUS",
        "17": "VIRTIO_NET_F_CTRL_VQ",
        "18": "VIRTIO_NET_F_CTRL_RX",
        "19": "VIRTIO_NET_F_CTRL_VLAN",
        "22": "VIRTIO_NET_F_MQ",
        "32": "VIRTIO_F_VERSION_1",
        "33": "VIRTIO_F_IOMMU_PLATFORM",
        "40": "VIRTIO_F_RING_RESET"
      },
      "driver_feature": {
        "value": "0x3004f182b",
        "0": "VIRTIO_NET_F_CSUM",
        "1": "VIRTIO_NET_F_GUEST_CSUM",
        "3": "VIRTIO_NET_F_MTU",
        "5": "VIRTIO_NET_F_MAC",
        "11": "VIRTIO_NET_F_HOST_TS04",
        "12": "VIRTIO_NET_F_HOST_TS06",
        "16": "VIRTIO_NET_F_STATUS",
        "17": "VIRTIO_NET_F_CTRL_VQ",
        "18": "VIRTIO_NET_F_CTRL_RX",
        "19": "VIRTIO_NET_F_CTRL_VLAN",
        "22": "VIRTIO_NET_F_MQ",
        "32": "VIRTIO_F_VERSION_1",
        "33": "VIRTIO_F_IOMMU_PLATFORM"
      },
      "status": {
        "value": "0xf",
        "0": "STATUS_ACKNOWLEDGE",
        "1": "STATUS_DRIVER"
      }
    }
  ]
}
```

virtio-net hotplug – 4

7. virtio_net device sf rep into ovs for datapath offload

```
for i in {0..7};do ovs-vsctl add-port ovsbr en3f0pf0sf$((2000+i));done  
ovs-vsctl show
```

```
root@localhost:~#  
root@localhost:~# virtnet list | grep sf_rep_net_device  
    "sf_rep_net_device": "en3f0pf0sf2000",  
    "sf_rep_net_device": "en3f0pf0sf2001",  
    "sf_rep_net_device": "en3f0pf0sf2002",  
    "sf_rep_net_device": "en3f0pf0sf2003",  
    "sf_rep_net_device": "en3f0pf0sf2004",  
    "sf_rep_net_device": "en3f0pf0sf2005",  
    "sf_rep_net_device": "en3f0pf0sf2006",  
    "sf_rep_net_device": "en3f0pf0sf2007",  
root@localhost:~#
```

```
root@localhost:~# ovs-vsctl show  
0f4f06d1-347e-4757-873c-b9cd6bb82f03  
    Bridge ovsbr  
        Port en3f1pf1sf0  
            Interface en3f1pf1sf0  
        Port en3f0pf0sf2002  
            Interface en3f0pf0sf2002  
        Port en3f0pf0sf2006  
            Interface en3f0pf0sf2006  
        Port en3f0pf0sf2001  
            Interface en3f0pf0sf2001  
        Port en3f0pf0sf2000  
            Interface en3f0pf0sf2000  
        Port en3f0pf0sf0  
            Interface en3f0pf0sf0  
        Port en3f0pf0sf2005  
            Interface en3f0pf0sf2005  
        Port en3f0pf0sf2007  
            Interface en3f0pf0sf2007  
        Port en3f0pf0sf2003  
            Interface en3f0pf0sf2003  
        Port ovsbr  
            Interface ovsbr  
                type: internal  
        Port en3f0pf0sf2004  
            Interface en3f0pf0sf2004  
        Port bond0  
            Interface bond0  
    ovs_version: "2.17.8-3feee121f"
```


virtio-net hotplug – 5

8. verify host virtio_net devices and configure ip and ping to peer

```
[host]# modprobe -v virtio-pci && modprobe -v virtio-net
```

Check on host virtio-net
pcie devices

```
[root@localhost ~]#  
[root@localhost ~]# lspci | grep Virt  
03:00.0 Signal processing controller: Huawei Technologies Co., Ltd. iBMA Virtual Network Adapter (rev 01)  
b2:00.0 Ethernet controller: Red Hat, Inc. Virtio network device (rev 01)  
b3:00.0 Ethernet controller: Red Hat, Inc. Virtio network device (rev 01)  
b4:00.0 Ethernet controller: Red Hat, Inc. Virtio network device (rev 01)  
b5:00.0 Ethernet controller: Red Hat, Inc. Virtio network device (rev 01)  
b6:00.0 Ethernet controller: Red Hat, Inc. Virtio network device (rev 01)  
b7:00.0 Ethernet controller: Red Hat, Inc. Virtio network device (rev 01)  
b8:00.0 Ethernet controller: Red Hat, Inc. Virtio network device (rev 01)  
b9:00.0 Ethernet controller: Red Hat, Inc. Virtio network device (rev 01)  
[root@localhost ~]# modprobe virtio_net  
[root@localhost ~]# ll /sys/class/net/ | grep virtio  
lrwxrwxrwx. 1 root root 0 Sep 15 06:11 ens2 -> ../../devices/pci0000:ae/0000:ae:02.0/0000:af:00.0/0000:b0:02.0/0000:b3:00.0/virtio1/net/ens2  
lrwxrwxrwx. 1 root root 0 Sep 15 06:11 ens4 -> ../../devices/pci0000:ae/0000:ae:02.0/0000:af:00.0/0000:b0:04.0/0000:b5:00.0/virtio2/net/ens4  
lrwxrwxrwx. 1 root root 0 Sep 15 06:11 ens5 -> ../../devices/pci0000:ae/0000:ae:02.0/0000:af:00.0/0000:b0:05.0/0000:b6:00.0/virtio3/net/ens5  
lrwxrwxrwx. 1 root root 0 Sep 15 06:11 ens6 -> ../../devices/pci0000:ae/0000:ae:02.0/0000:af:00.0/0000:b0:06.0/0000:b7:00.0/virtio4/net/ens6  
lrwxrwxrwx. 1 root root 0 Sep 15 06:11 ens7 -> ../../devices/pci0000:ae/0000:ae:02.0/0000:af:00.0/0000:b0:07.0/0000:b8:00.0/virtio5/net/ens7  
lrwxrwxrwx. 1 root root 0 Sep 15 06:11 ens8 -> ../../devices/pci0000:ae/0000:ae:02.0/0000:af:00.0/0000:b0:08.0/0000:b9:00.0/virtio6/net/ens8  
lrwxrwxrwx. 1 root root 0 Sep 15 06:11 eth0 -> ../../devices/pci0000:ae/0000:ae:02.0/0000:af:00.0/0000:b0:03.0/0000:b4:00.0/virtio2/net/eth0  
lrwxrwxrwx. 1 root root 0 Sep 15 06:11 eth1 -> ../../devices/pci0000:ae/0000:ae:02.0/0000:af:00.0/0000:b0:06.0/0000:b7:00.0/virtio5/net/eth1  
[root@localhost ~]# ethtool -i ens2  
driver: virtio_net  
version: 1.0.0  
firmware-version:  
expansion-rom-version:  
bus-info: 0000:b3:00.0  
supports-statistics: yes  
supports-test: no  
supports-eprom-access: no  
supports-register-dump: no  
supports-priv-flags: no  
[root@localhost ~]#
```

Check on host virtio-net
driver is loaded and
ethernet devices is
working

Run ping and traffic from
host virtio_net

```
[root@localhost ~]#  
[root@localhost ~]# ifconfig ens2 192.168.200.200  
[root@localhost ~]# ping 192.168.200.100  
PING 192.168.200.100 (192.168.200.100) 56(84) bytes of data.  
64 bytes from 192.168.200.100: icmp_seq=1 ttl=64 time=138 ms  
64 bytes from 192.168.200.100: icmp_seq=2 ttl=64 time=0.131 ms  
64 bytes from 192.168.200.100: icmp_seq=3 ttl=64 time=0.135 ms
```

Check flows hw offloaded
in DPU OVS

```
root@localhost:~#  
root@localhost:~# ovs-appctl dpctl/dump-flows type=offloaded  
recirc_id(0),in_port(1),eth(src=0c:42:a1:6d:7c:76,dst=00:11:22:33:44:22),eth_type(0x0800),ipv4(frag=no), packets:21, bytes:2142, used:0.180s, actions:6  
recirc_id(0),in_port(6),eth(src=00:11:22:33:44:22,dst=0c:42:a1:6d:7c:76),eth_type(0x0800),ipv4(frag=no), packets:21, bytes:2058, used:0.180s, actions:1  
root@localhost:~#
```

Virtio-net Static for Virtualization

virtio-net static pf/vf – 1

1. FW config

```
mlxconfig -d 03:00.0 -y reset
mlxconfig -d 03:00.0 -y set \
    INTERNAL_CPU_MODEL=1 SRIOV_EN=1 \
    PCI_SWITCH_EMULATION_ENABLE=0 \
    PCI_SWITCH_EMULATION_NUM_PORT=0 \
VIRTIO_NET_EMULATION_ENABLE=1 \
VIRTIO_NET_EMULATION_NUM_VF=126 \
VIRTIO_NET_EMULATION_NUM_PF=4 \
VIRTIO_NET_EMULATION_NUM_MSIX=32 \
    LAG_RESOURCE_ALLOCATION=1 \
    PER_PF_NUM_SF=1 PF_TOTAL_SF=512 \
    PF_SF_BAR_SIZE=8

mlxconfig -d 03:00.1 -y set \
    PER_PF_NUM_SF=1 PF_TOTAL_SF=2 \
    PF_SF_BAR_SIZE=8
```

2. Enable virtio-net-controller service

```
systemctl enable virtio-net-controller
```

3. Edit /opt/mellanox/mlnx_virtnet/virtnet.conf

```
{
    "ib_dev_lag": "mlx5_bond_0",
    "ib_dev_for_static_pf": "mlx5_bond_0",
    "is_lag": 1,
    "recovery": 1,
    "sf_pool_percent": 0,
    "sf_pool_force_destroy": 0
}
```

4. Powe cycle

[host]

Add "**pci=realloc intel_iommu=on iommu=pt**" into kernel boot grub options

Add "**pci=assign-busses**" in case > 127 VFs

[host]

ipmitool power cycle

virtio-net static pf/vf – 2

5. Check after power up

```
systemctl status  
virtio-net-controller  
virtnet list  
virtnet query -p 0
```

- pci_max_vfs 是host virtio pci 设定
- msix 由
- Virtio Feature 由device和driver协商确定

```
{ 'all': '0x0', 'pf': '0x0', 'dbg_stats': '0x0', 'brief':  
{  
  "devices": [  
    {  
      "pf_id": 0,  
      "transitional": 0,  
      "vuid": "MT2333X02514VNETS0D0F1",  
      "pci_bdf": "02:00.1",  
      "pci_dev_id": "0x1041",  
      "pci_vendor_id": "0x1af4",  
      "pci_class_code": "0x20000",  
      "pci_subsys_id": "0x1041",  
      "pci_subsys_vendor_id": "0x1af4",  
      "pci_revision_id": "1",  
      "pci_max_vfs": "8",  
      "enabled_vfs": "8",  
      "device_feature": {  
        "value": "0x8800032300c7182b",  
        "0": "VIRTIO_NET_F_CSUM",  
        "1": "VIRTIO_NET_F_GUEST_CSUM",  
        "3": "VIRTIO_NET_F_MTU",  
        "5": "VIRTIO_NET_F_MAC",  
        "11": "VIRTIO_NET_F_HOST_TS04",  
        "12": "VIRTIO_NET_F_HOST_TS06",  
        "16": "VIRTIO_NET_F_STATUS",  
        "17": "VIRTIO_NET_F_CTRL_VQ",  
        "18": "VIRTIO_NET_F_CTRL_RX",  
        "22": "VIRTIO_NET_F_MQ",  
        "23": "VIRTIO_NET_F_CTRL_MAC_ADDR",  
        "32": "VIRTIO_F_VERSION_1",  
        "33": "VIRTIO_F_IOMMU_PLATFORM",  
        "37": "VIRTIO_F_SR_IOV",  
        "40": "VIRTIO_F_RING_RESET",  
        "41": "VIRTIO_F_ADMIN_VQ",  
        "59": "VIRTIO_NET_F_GUEST_HDRLEN",  
        "63": "VIRTIO_NET_F_SPEED_DUPLEX"  
      },  
      "driver_feature": {  
        "value": "0x8000002300c7182b",  
        "0": "VIRTIO_NET_F_CSUM",  
        "1": "VIRTIO_NET_F_GUEST_CSUM",  
        "3": "VIRTIO_NET_F_MTU",  
        "5": "VIRTIO_NET_F_MAC",  
        "11": "VIRTIO_NET_F_HOST_TS04",  
        "12": "VIRTIO_NET_F_HOST_TS06",  
        "16": "VIRTIO_NET_F_STATUS",  
        "17": "VIRTIO_NET_F_CTRL_VQ",  
        "18": "VIRTIO_NET_F_CTRL_RX",  
        "22": "VIRTIO_NET_F_MQ",  
        "23": "VIRTIO_NET_F_CTRL_MAC_ADDR",  
        "32": "VIRTIO_F_VERSION_1",  
        "33": "VIRTIO_F_IOMMU_PLATFORM",  
        "37": "VIRTIO_F_SR_IOV",  
        "63": "VIRTIO_NET_F_SPEED_DUPLEX"  
      },  
      "status": {  
        "value": "0xf",  
        "0": "ACK"  
      }  
    }  
  ],  
  "controller": {  
    "emulation_manager": "mlx5_bond_0",  
    "max_hotplug_devices": "0",  
    "max_virt_net_devices": "4",  
    "max_virt_queues": "64",  
    "max_tunnel_descriptors": "6",  
    "supported_features": {  
      "value": "0x80000333004f982b",  
      "0": "VIRTIO_NET_F_CSUM",  
      "1": "VIRTIO_NET_F_GUEST_CSUM",  
      "3": "VIRTIO_NET_F_MTU",  
      "5": "VIRTIO_NET_F_MAC",  
      "11": "VIRTIO_NET_F_HOST_TS04",  
      "12": "VIRTIO_NET_F_HOST_TS06",  
      "15": "VIRTIO_F_MRG_RX_BUFFER",  
      "16": "VIRTIO_NET_F_STATUS",  
      "17": "VIRTIO_NET_F_CTRL_VQ",  
      "18": "VIRTIO_NET_F_CTRL_RX",  
      "19": "VIRTIO_NET_F_CTRL_VLAN",  
      "22": "VIRTIO_NET_F_MQ",  
      "32": "VIRTIO_F_VERSION_1",  
      "33": "VIRTIO_F_IOMMU_PLATFORM",  
      "36": "VIRTIO_F_ORDER_PLATFORM",  
      "37": "VIRTIO_F_SR_IOV",  
      "40": "VIRTIO_F_RING_RESET",  
      "41": "VIRTIO_F_ADMIN_VQ",  
      "63": "VIRTIO_NET_F_SPEED_DUPLEX"  
    },  
    "supported_virt_queue_types": {  
      "value": "0x1",  
      "0": "SPLIT"  
    },  
    "supported_event_modes": {  
      "value": "0x5",  
      "0": "NO_MSIX_MODE",  
      "2": "MSIX_MODE"  
    }  
  },  
  "devices": [  
    {  
      "pf_id": 0,  
      "function_type": "static PF",  
      "transitional": 0,  
      "vuid": "VNETS0D0F0",  
      "msix_num_pool_size": 0,  
      "min_msix_num": 0,  
      "max_msix_num": 32,  
      "bdf": "af:00.0",  
      "sf_num": 1000,  
      "sf_parent_device": "mlx5_bond_0",  
      "sf_parent_device_pci_addr": "0000:03:00.0",  
      "sf_rep_net_device": "en3f0pf0sf1000",  
      "sf_rep_net_ifindex": 13,  
    }  
  ]  
}
```

virtio-net static pf/vf – 3

6. Create virtio_net vf and modify mac and features

[host] #static vf 是从host上 virtio-pci创建

```
#echo 0 > /sys/bus/pci/drivers/virtio-pci/0000:af:00.0/sriov_drivers_autoprobe
```

```
lspci | grep Virtio
```

```
cat /sys/bus/pci/drivers/virtio-pci/0000:af:00.0/sriov_totalvfs
```

```
cat /sys/bus/pci/drivers/virtio-pci/0000:af:00.1/sriov_totalvfs
```

```
modprobe -r virtio_net
```

```
echo 8 > /sys/bus/pci/drivers/virtio-pci/0000:af:00.0/sriov_numvfs
```

```
echo 8 > /sys/bus/pci/drivers/virtio-pci/0000:af:00.1/sriov_numvfs
```

[dpu]

```
for i in {0..7};do virtnet modify -p 0 -v $i device -m "00:11:22:33:44:0$((0+i))" -f 0x23004f182b;done
```

```
virtnet list
```

```
virtnet query -p 0 -v 0
```

[host]

```
modprobe -v virtio_net
```

```
[root@localhost ~]# lspci | grep Virtio
af:00.0 Ethernet controller: Red Hat, Inc. Virtio network device (rev 01)
af:00.1 Ethernet controller: Red Hat, Inc. Virtio network device (rev 01)
af:00.2 Ethernet controller: Red Hat, Inc. Virtio network device (rev 01)
af:00.3 Ethernet controller: Red Hat, Inc. Virtio network device (rev 01)
[root@localhost ~]# cat /sys/bus/pci/drivers/virtio-pci/0000:af:00.0/sriov_totalvfs
8
[root@localhost ~]# cat /sys/bus/pci/drivers/virtio-pci/0000:af:00.1/sriov_totalvfs
8
```

```
[root@localhost ~]# cat /sys/bus/pci/drivers/virtio-pci/0000:af:00.0/sriov_totalvfs
8
[root@localhost ~]# cat /sys/bus/pci/drivers/virtio-pci/0000:af:00.1/sriov_totalvfs
8
[root@localhost ~]# echo 8 > /sys/bus/pci/drivers/virtio-pci/0000:af:00.0/sriov_numvfs
[root@localhost ~]# lspci | grep Virtio
af:00.0 Ethernet controller: Red Hat, Inc. Virtio network device (rev 01)
af:00.1 Ethernet controller: Red Hat, Inc. Virtio network device (rev 01)
af:00.2 Ethernet controller: Red Hat, Inc. Virtio network device (rev 01)
af:00.3 Ethernet controller: Red Hat, Inc. Virtio network device (rev 01)
af:00.5 Ethernet controller: Red Hat, Inc. Virtio network device (rev 01)
af:00.6 Ethernet controller: Red Hat, Inc. Virtio network device (rev 01)
af:00.7 Ethernet controller: Red Hat, Inc. Virtio network device (rev 01)
af:01.0 Ethernet controller: Red Hat, Inc. Virtio network device (rev 01)
af:01.1 Ethernet controller: Red Hat, Inc. Virtio network device (rev 01)
af:01.2 Ethernet controller: Red Hat, Inc. Virtio network device (rev 01)
af:01.3 Ethernet controller: Red Hat, Inc. Virtio network device (rev 01)
af:01.4 Ethernet controller: Red Hat, Inc. Virtio network device (rev 01)
```

```
altname enp175s0f3
24: ens6f0v0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc mq state UP mode DEFAULT group default qlen 1000
    link/ether 00:11:22:33:44:00 brd ff:ff:ff:ff:ff:ff
    altname enp175s0f0v0
25: ens6f0v1: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc mq state UP mode DEFAULT group default qlen 1000
    link/ether 00:11:22:33:44:01 brd ff:ff:ff:ff:ff:ff
    altname enp175s0f0v1
26: ens6f0v2: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc mq state UP mode DEFAULT group default qlen 1000
    link/ether 00:11:22:33:44:02 brd ff:ff:ff:ff:ff:ff
    altname enp175s0f0v2
27: ens6f0v3: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc mq state UP mode DEFAULT group default qlen 1000
    link/ether 00:11:22:33:44:03 brd ff:ff:ff:ff:ff:ff
    altname enp175s0f0v3
28: ens6f0v4: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc mq state UP mode DEFAULT group default qlen 1000
    link/ether 00:11:22:33:44:04 brd ff:ff:ff:ff:ff:ff
    altname enp175s0f0v4
29: ens6f0v5: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc mq state UP mode DEFAULT group default qlen 1000
    link/ether 00:11:22:33:44:05 brd ff:ff:ff:ff:ff:ff
    altname enp175s0f0v5
30: ens6f0v6: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc mq state UP mode DEFAULT group default qlen 1000
    link/ether 00:11:22:33:44:06 brd ff:ff:ff:ff:ff:ff
    altname enp175s0f0v6
31: ens6f0v7: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc mq state UP mode DEFAULT group default qlen 1000
    link/ether 00:11:22:33:44:07 brd ff:ff:ff:ff:ff:ff
    altname enp175s0f0v7
```


virtio-net static pf/vf – 4

5. ovs and ping verify

```
[dpu]
virtnet list | grep sf_rep_net_device
for i in {0..7};do ovs-vsctl add-port ovsbr en3f0pf0sf$((3000+i));done
ovs-vsctl show

[host]
ifconfig ens6f0v0 192.168.200.2

[dpu]
ovs-appctl dpctl/dump-flows type=offloaded
```

```
root@localhost:~# virtnet list | grep sf_rep_net_device
"sf_rep_net_device": "en3f0pf0sf1000",
"sf_rep_net_device": "en3f0pf0sf3000",
"sf_rep_net_device": "en3f0pf0sf3001",
"sf_rep_net_device": "en3f0pf0sf3002",
"sf_rep_net_device": "en3f0pf0sf3003",
"sf_rep_net_device": "en3f0pf0sf3004",
"sf_rep_net_device": "en3f0pf0sf3005",
"sf_rep_net_device": "en3f0pf0sf3006",
"sf_rep_net_device": "en3f0pf0sf3007",
"sf_rep_net_device": "en3f0pf0sf1001",
"sf_rep_net_device": "en3f0pf0sf1002",
"sf_rep_net_device": "en3f0pf0sf1003",
```

```
[root@localhost ~]# ifconfig ens6f0v0 192.168.200.2
[root@localhost ~]# ping 192.168.200.100
PING 192.168.200.100 (192.168.200.100) 56(84) bytes of data.
64 bytes from 192.168.200.100: icmp_seq=1 ttl=64 time=133 ms
64 bytes from 192.168.200.100: icmp_seq=2 ttl=64 time=0.110 ms
```

```
root@localhost:~# ovs-appctl dpctl/dump-flows type=offloaded
recirc_id(0),in_port(1),eth(src=0c:42:a1:6d:7c:76,dst=00:11:22:33:44:00),eth_type(0x0806), packets:0, bytes:0, used:never, actions:5
recirc_id(0),in_port(1),eth(src=0c:42:a1:6d:7c:76,dst=00:11:22:33:44:00),eth_type(0x0800),ipv4(frag=no), packets:118, bytes:12036, used:0.110s, actions:5
recirc_id(0),in_port(5),eth(src=00:11:22:33:44:00,dst=0c:42:a1:6d:7c:76),eth_type(0x0806), packets:0, bytes:0, used:never, actions:1
recirc_id(0),in_port(5),eth(src=00:11:22:33:44:00,dst=0c:42:a1:6d:7c:76),eth_type(0x0800),ipv4(frag=no), packets:118, bytes:11564, used:0.110s, actions:1
root@localhost:~#
```

```
root@localhost:~# ovs-vsctl show
0f4f06d1-347e-4757-873c-b9cd6bb82f03
Bridge ovsbr
  Port en3f1pf1sf0
    Interface en3f1pf1sf0
  Port en3f0pf0sf3003
    Interface en3f0pf0sf3003
  Port en3f0pf0sf3005
    Interface en3f0pf0sf3005
  Port en3f0pf0sf3006
    Interface en3f0pf0sf3006
  Port en3f0pf0sf0
    Interface en3f0pf0sf0
  Port en3f0pf0sf3004
    Interface en3f0pf0sf3004
  Port en3f0pf0sf3007
    Interface en3f0pf0sf3007
  Port en3f0pf0sf3001
    Interface en3f0pf0sf3001
  Port en3f0pf0sf3002
    Interface en3f0pf0sf3002
  Port en3f0pf0sf3000
    Interface en3f0pf0sf3000
  Port ovsbr
    Interface ovsbr
      type: internal
  Port bond0
    Interface bond0
ovs_version: "2.17.8-3feee121f"
```

Virtio VF Dynamic MSIX

- 每个PF支持动态msi池, 缺省是0, VF按固定值分配

```
[DPU]# mlxconfig -y -d 03:00.0 s VIRTIO_NET_EMULATION_NUM_MSIX=32 NUM_VF_MSIX=32
```

- 查询当前PF msix池的大小, 初始化是0, 因为msix被固定分配出去了, 需要从已分配的VF中回收

```
[DPU]# virtnet list | grep -i '"pf_id": 0' -A 8 | grep -i msix_num_pool_size
```

- 查询VF可以分配的msix的范围

```
[DPU]# virtnet list | grep -i '"pf_id": 0' -A 8 | grep -i max_msix_num
```

```
[DPU]# virtnet list | grep -i '"pf_id": 0' -A 8 | grep -i min_msix_num
```

- 释放所有的VF的msix到PF的msix池, 然后重新按需分配

```
[host]# echo <vf0_bdf> > /sys/bus/pci/drivers/virtio-pci/unbind
```

```
[host]# echo <vf1_bdf> > /sys/bus/pci/drivers/virtio-pci/unbind
```

```
[DPU]# virtnet modify -p 0 -v 0 device -n 0
```

```
[DPU]# virtnet modify -p 0 -v 1 device -n 0
```

```
[DPU]# virtnet list | grep -i '"pf_id": 0' -A 8 | grep -i msix_num_pool_size
```

- 重新分配VF的msix输出

```
[DPU]# virtnet modify -p 0 -v 0 device -n 48
```

```
[DPU]# virtnet modify -p 0 -v 1 device -n 16
```

- 查询分配给VF的msix数


```
[DPU]# virtnet query -p 0 -v 0 | grep -i num_msix
```

```
[DPU]# virtnet query -p 0 -v 1 | grep -i num_msix
```

- 重新在host上加载VF的virtio驱动

```
[host]# echo <vf0_bdf> > /sys/bus/pci/drivers/virtio-pci/bind
```

```
[host]# echo <vf1_bdf> > /sys/bus/pci/drivers/virtio-pci/bind
```



```
root@localhost:~# virtnet query -p 0 -v 0 | grep queue
"max_queues": "32",
"enabled_queues": "31",
"net_max_queue_pairs": "15",
"enabled-queues-info": [
```

```
root@amdgen5sz:~# ethtool -l ens2f1v0
Channel parameters for ens2f1v0:
Pre-set maximums:
RX:                n/a
TX:                n/a
Other:             n/a
Combined:          15
Current hardware settings:
RX:                n/a
TX:                n/a
Other:             n/a
Combined:          15
```

```
root@localhost:~# virtnet query -p 0 -v 0 | grep msi
"msix_config_vector": "0x0",
"num_msix": "32",
"msix_vector": "0x1",
"msix_vector": "0x2",
"msix_vector": "0x3",
"msix_vector": "0x4",
"msix_vector": "0x5",
"msix_vector": "0x6",
"msix_vector": "0x7",
"msix_vector": "0x8",
"msix_vector": "0x9",
"msix_vector": "0xa",
"msix_vector": "0xb",
"msix_vector": "0xc",
"msix_vector": "0xd",
"msix_vector": "0xe",
"msix_vector": "0xf",
"msix_vector": "0x10",
"msix_vector": "0x11",
"msix_vector": "0x12",
"msix_vector": "0x13",
"msix_vector": "0x14",
"msix_vector": "0x15",
"msix_vector": "0x16",
"msix_vector": "0x17",
"msix_vector": "0x18",
"msix_vector": "0x19",
"msix_vector": "0x1a",
"msix_vector": "0x1b",
"msix_vector": "0x1c",
"msix_vector": "0x1d",
"msix_vector": "0x1e",
"msix_vector": "0xffff",
```

virtio-net run with ovs-dpdk example

■ LAG+Jumbo+vxlan+connection_tracking

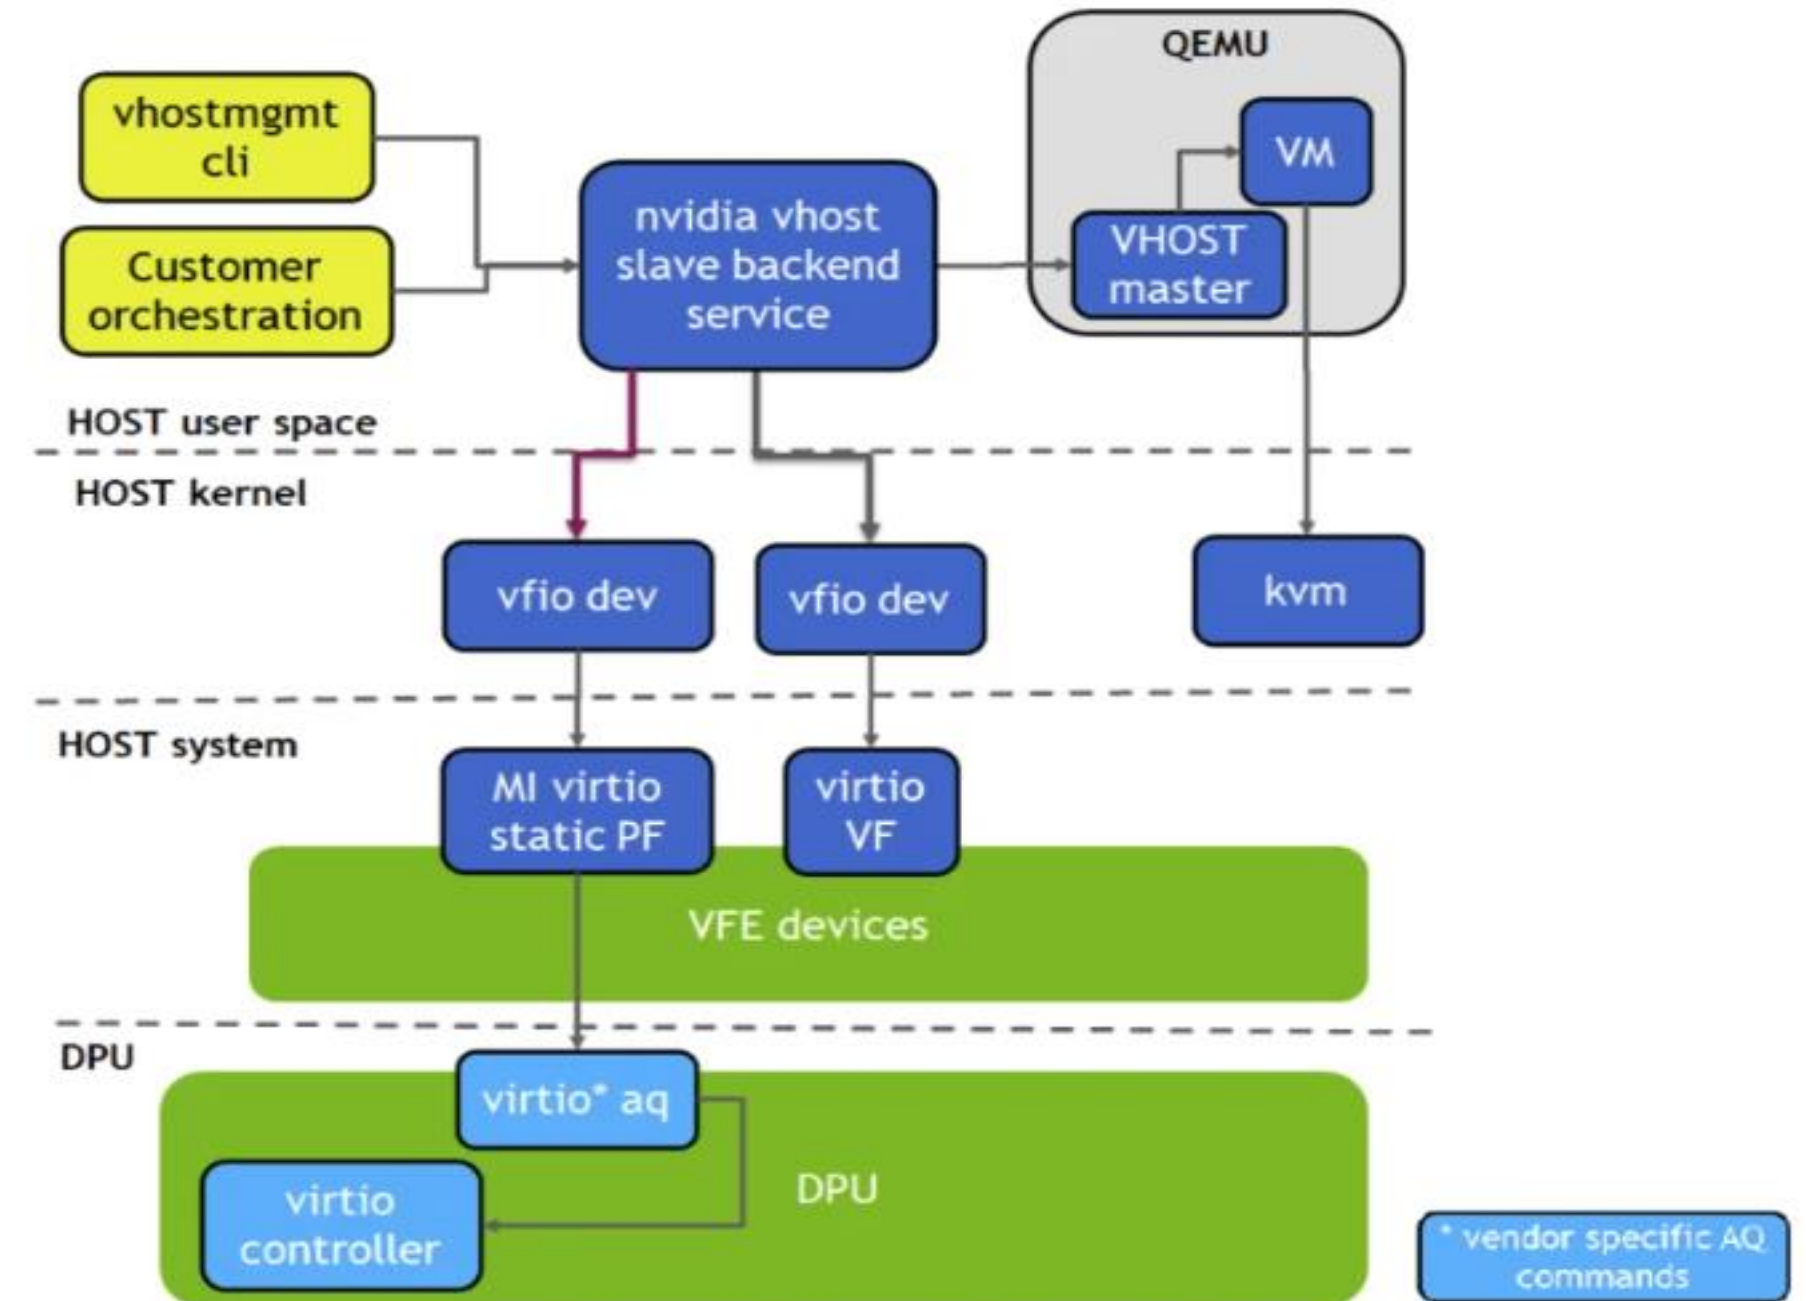
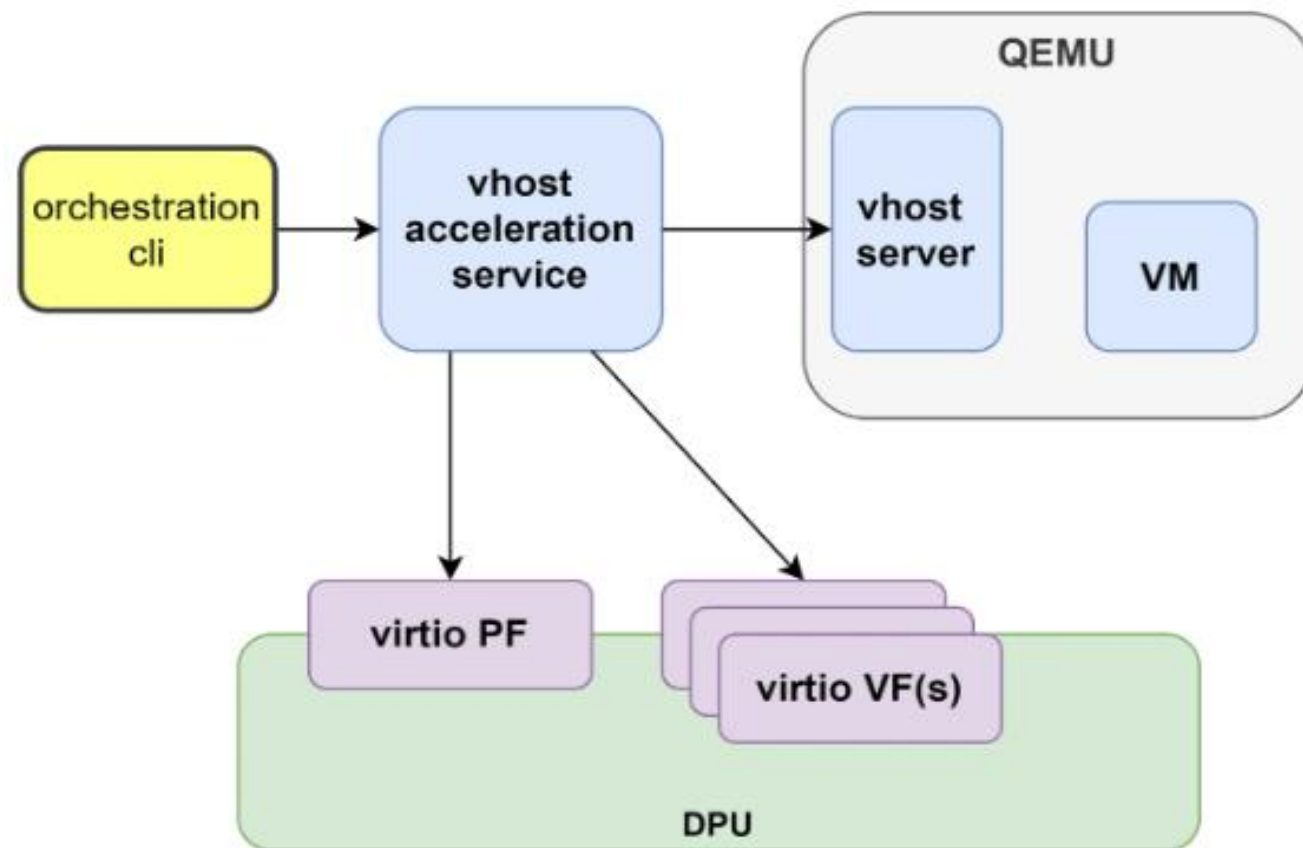
	Static PF	Hotplug PF	SR-IOV VF
SF Range	1000-999	2000-2999	3000 and above

```
echo 2048 > /sys/devices/system/node/node0/hugepages/hugepages-2048kB/nr_hugepages
ovs-vsctl add-br br-int -- set bridge br-int datapath_type=netdev
ovs-vsctl add-br br-ext -- set bridge br-ext datapath_type=netdev
ovs-vsctl --no-wait set o . other_config:dpdk-extra="-a 0000:03:00.0,representor=pf0sf[2000-2030,65535],dv_xmeta_en=1,sys_mem_en=1 mtu_request=9000",
ovs-vsctl --no-wait set o . other_config:dpdk-init=true
ovs-vsctl --no-wait set o . other_config:hw-offload=true other_config:max-idle=60000
ovs-vsctl add-port br-ext bond0 -- set Interface bond0 type=dpdk options:dpdk-
devargs="0000:03:00.0,dv_xmeta_en=1,sys_mem_en=1" mtu_request=9000
ovs-vsctl add-port br-int hostpf -- set Interface hostpf type=dpdk options:dpdk-
devargs="0000:03:00.0,representor=sf65535,dv_xmeta_en=1,sys_mem_en=1" mtu_request=9000
ovs-vsctl add-port br-int vxlan0 -- set interface vxlan0 type=vxlan options:remote_ip=11.0.0.2 options:local_ip=11.0.0.1
options:key=5 mtu_request=9000

ovs-ofctl del-flows br-int
rep=hostpf
vxlan_rep=vxlan0
ovs-ofctl add-flow br-int "table=0, priority=10, arp,in_port=$rep, actions=$vxlan_rep "
ovs-ofctl add-flow br-int "table=0, priority=10, arp,in_port=$vxlan_rep, actions=$rep"
ovs-ofctl add-flow br-int "table=0, priority=10, icmp,in_port=$rep, actions=$vxlan_rep"
ovs-ofctl add-flow br-int "table=0, priority=10, icmp,in_port=$vxlan_rep, actions=$rep"
ovs-ofctl add-flow br-int "table=0, priority=50, ct_state=-trk, tcp, in_port=$rep, actions=ct(table=0)"
ovs-ofctl add-flow br-int "table=0, priority=50, ct_state=-trk, tcp, in_port=$vxlan_rep, actions=ct(table=0)"
ovs-ofctl add-flow br-int "table=0, priority=50, ct_state=+trk+new, tcp, in_port=$vxlan_rep actions=ct(commit),$rep"
ovs-ofctl add-flow br-int "table=0, priority=50, ct_state=+trk+new, tcp, in_port=$rep, actions=ct(commit),$vxlan_rep"
ovs-ofctl add-flow br-int "table=0, priority=50, ct_state=+trk+est, tcp, in_port=$vxlan_rep, actions=$rep"
ovs-ofctl add-flow br-int "table=0, priority=50, ct_state=+trk+est, tcp, in_port=$rep, actions=$vxlan_rep"
ovs-ofctl dump-flows br-int
```


Virtio live migration

- Virtio热迁移满足虚拟化场景
- Host需要运行backend service, 利用vhost的机制和qemu接口
- Backend service利用静态PF的virtio 设备同DPU进行通信
- 有接口支持定制化适配客户的热迁移机制





Q&A

Q&A

- 更改FW配置后，server不能启动，BIOS hang住
 - 可能原因是FW更改配置导致BIOS申请的资源不可用，BIOS会hang住，解决方法是方法进入DPU OS或者BMC，复位FW配置或者更改FW配置，掉电重启尝试
 - DPU的NIC mode如果配置virtio emulation会导致BIOS hang，这个可以通过DPU的BMC进入UEFI menu更改到DPU mode
- 在配置静态virtio-net时，OS的virtio 驱动crash，kernel hang
 - 可能原因是OS主动加载了virtio pci驱动，但是因为DPU上的virtio net没有配置完全，或者services没有启动，导致virito net的controll message没有响应，被hang住。
 - 解决方式是，DPU的virtnet先在host os启动之前配置好，或者host先禁止加载virtio pci驱动
- Virtnet 性能
 - virtio-net性能和队列数相关，单条队列的性能有上限，尽量使用多个队列多流提高性能，开启checksum，LRO、TSO 硬件卸载
 - 单virtio-net的最多队列数是31, ((64-1)/2)
- Virtio host UEFI
 - Virtio(blk and net) uefi driver 包含在DPU的 exp rom里，不建议用BIOS自带的
 - Enable或者disable只能从host上设置，和DPU上设置是分开的

```
[Host]
root@amdgen5sz:~# mlxconfig -d 02:00.0 q | grep UEFI
    UEFI_HII_EN                True(1)
    UEFI_LOGS                   DISABLED(0)
    EXP_ROM_VIRTIO_NET_UEFI_ARM_ENABLE    False(0)
    EXP_ROM_VIRTIO_NET_UEFI_x86_ENABLE    False(0)
    EXP_ROM_VIRTIO_BLK_UEFI_ARM_ENABLE    False(0)
    EXP_ROM_VIRTIO_BLK_UEFI_x86_ENABLE    False(0)
    EXP_ROM_NVME_UEFI_x86_ENABLE          True(1)
    EXP_ROM_UEFI_ARM_ENABLE               True(1)
    EXP_ROM_UEFI_x86_ENABLE               True(1)
```