

宁夏培训内容的总结

1. 链接的总结

a. 系统下载

正常系统下载使用官方网站

Ubuntu: <https://ubuntu.com/download/server>

CentOS: <https://www.centos.org/download/>

RockyLinux: <https://rockylinux.org/download>

特殊的系统下载使用google搜索页可以找到对应的下载页面。

例如，会议上提到的ubuntu22.04.4，可以使用下面这个链接

<https://old-releases.ubuntu.com/releases/22.04.4/>

这个链接也是通过google搜索得来。

b. IB相关驱动

- 各种工具及固件匹配关系的最新版本查询的链接

<https://www.nvidia.com/en-us/networking/products/lts-releases/>

<https://developer.nvidia.com/networking/infiniband-software>

- IB网卡驱动及对应的工具包

IB网卡的驱动可以通过关键词搜索 **OFED download**，链接如下

https://network.nvidia.com/products/infiniband-drivers/linux/mlnx_ofed/

选择LTS（长期支持）版本和对应的系统和架构即可。下载时选择ISO或者tar

未来IB网卡会从OFED过渡到DOCA，链接如下

<https://developer.nvidia.com/doca-downloads>

- MFT (Mellanox Firmware Toolkit)

<https://network.nvidia.com/products/adapter-software/firmware-tools/>

- 交换机固件及网卡固件

<https://network.nvidia.com/support/firmware/firmware-downloads/>

- CPLD的说明文档

<https://docs.nvidia.com/networking/display/mftv4301508/cpldupdate+-+tool+for+programming+on-board+cplds+on+nvidia+devices>

c. GPU相关驱动以及CUDA包

- CUDA工具

<https://developer.nvidia.com/cuda-toolkit-archive>

选择对应版本，选择runfile。如何安装下面会有对应的介绍。

- GPU驱动

<https://www.nvidia.cn/drivers/lookup/>

- nvlink驱动

https://developer.download.nvidia.cn/compute/cuda/repos/ubuntu2204/x86_64/

找到cuda对应的fabricmanager的驱动版本。配套版本查找方法如下

<https://docs.nvidia.com/cuda/cuda-toolkit-release-notes/index.html#id9>

- container (docker) 的支持包

https://developer.download.nvidia.cn/compute/cuda/repos/ubuntu2204/x86_64/

nvidia-container-toolkit_1.17.8-1_amd64.deb

nvidia-container-toolkit-base_1.17.8-1_amd64.deb

```
libnvidia-container1_1.17.8-1_amd64.deb  
libnvidia-container-tools_1.17.8-1_amd64.deb
```

安装方法：`dpkg -i *.deb`

- 关于apt源，也可以使用cuda的源来进行安装cuda

```
wget https://developer.download.nvidia.com/compute/cuda/repos/ubuntu2404/x86_64/cuda-keyring_1.1-1_all.deb  
sudo dpkg -i cuda-keyring_1.1-1_all.deb  
sudo apt-get update  
sudo apt-get -y install cuda-toolkit-13-0
```

d. apt源，一般使用清华或者阿里云

清华源举例

```
deb https://mirrors.tuna.tsinghua.edu.cn/ubuntu/ jammy main restricted  
deb https://mirrors.tuna.tsinghua.edu.cn/ubuntu/ jammy-updates main restricted  
deb https://mirrors.tuna.tsinghua.edu.cn/ubuntu/ jammy universe  
deb https://mirrors.tuna.tsinghua.edu.cn/ubuntu/ jammy-updates universe  
deb https://mirrors.tuna.tsinghua.edu.cn/ubuntu/ jammy multiverse  
deb https://mirrors.tuna.tsinghua.edu.cn/ubuntu/ jammy-updates multiverse  
deb https://mirrors.tuna.tsinghua.edu.cn/ubuntu/ jammy-backports main restricted universe multiverse  
deb https://mirrors.tuna.tsinghua.edu.cn/ubuntu/ jammy-security main restricted  
deb https://mirrors.tuna.tsinghua.edu.cn/ubuntu/ jammy-security universe  
deb https://mirrors.tuna.tsinghua.edu.cn/ubuntu/ jammy-security multiverse
```

阿里源举例

```
deb https://mirrors.aliyun.com/ubuntu/ jammy main restricted universe multiverse  
deb-src https://mirrors.aliyun.com/ubuntu/ jammy main restricted universe multiverse  
deb https://mirrors.aliyun.com/ubuntu/ jammy-security main restricted universe multiverse  
deb-src https://mirrors.aliyun.com/ubuntu/ jammy-security main restricted universe multiverse  
deb https://mirrors.aliyun.com/ubuntu/ jammy-updates main restricted universe multiverse  
deb-src https://mirrors.aliyun.com/ubuntu/ jammy-updates main restricted universe multiverse  
deb-src https://mirrors.aliyun.com/ubuntu/ jammy-backports main restricted universe multiverse
```

2. 操作系统安装步骤的总结

- a. 安装操作系统，一些基础软件的推荐。

```
apt install -y openssh-server curl wget vim htop tree ansible ipmitool lldpd net-tools unzip
```

- openssh-server
- curl
- wget
- vim
- htop
- tree
- ansible
- ipmitool
- lldpd
- net-tools
- unzip

- b. 更新内核（视需求更新）

```

# 查询内核
root@UFM-StorNet-Node01:/# cat /proc/version
Linux version 5.15.0-151-generic (buildd@lcy02-amd64-092) (gcc (Ubuntu 11.4.0-1ubuntu1~22.04) 11.4.0, GNU
ld (GNU Binutils for Ubuntu) 2.38) #161-Ubuntu SMP Tue Jul 22 14:25:40 UTC 2025
root@UFM-StorNet-Node01:/# uname -r
5.15.0-151-generic
root@UFM-StorNet-Node01:/# uname -a
Linux UFM-StorNet-Node01 5.15.0-151-generic #161-Ubuntu SMP Tue Jul 22 14:25:40 UTC 2025 x86_64 x86_6
4 x86_64 GNU/Linux

# 查看可用内核版本
apt search linux-image

# 安装特定内核版本
sudo apt install linux-image-[版本号]
sudo apt install linux-headers-[版本号]

# 更新引导加载器，更改下次启动的加载的内核
sudo update-grub

```

c. 关闭系统自动更新

```

sed -i.bak 's/1/0/' /etc/apt/apt.conf.d/10periodic

vi /etc/apt/apt.conf.d/50unattended-upgrades
加入
Unattended-Upgrade::Package-Blacklist {
    "linux-image-*";
    "linux-headers-*";
}

```

d. 设置BIOS

需要设置的项目

1. iommu / vt-d
2. cpu改为性能模式
3. 关闭 ACS

查询的命令

```

dmesg | grep -i iommu
cpupower frequency-info
lspci -vvv | grep ACSCtl

```

查询结果

```

root@H800GPU-ComServer-Node09:~# dmesg | grep -i iommu
[ 11.287745] DMAR-IR: IOAPIC id 8 under DRHD base 0x9c3fc000 IOMMU 19
[ 16.839506] iommu: Default domain type: Translated
[ 16.839506] iommu: DMA domain TLB invalidation policy: lazy mode

```

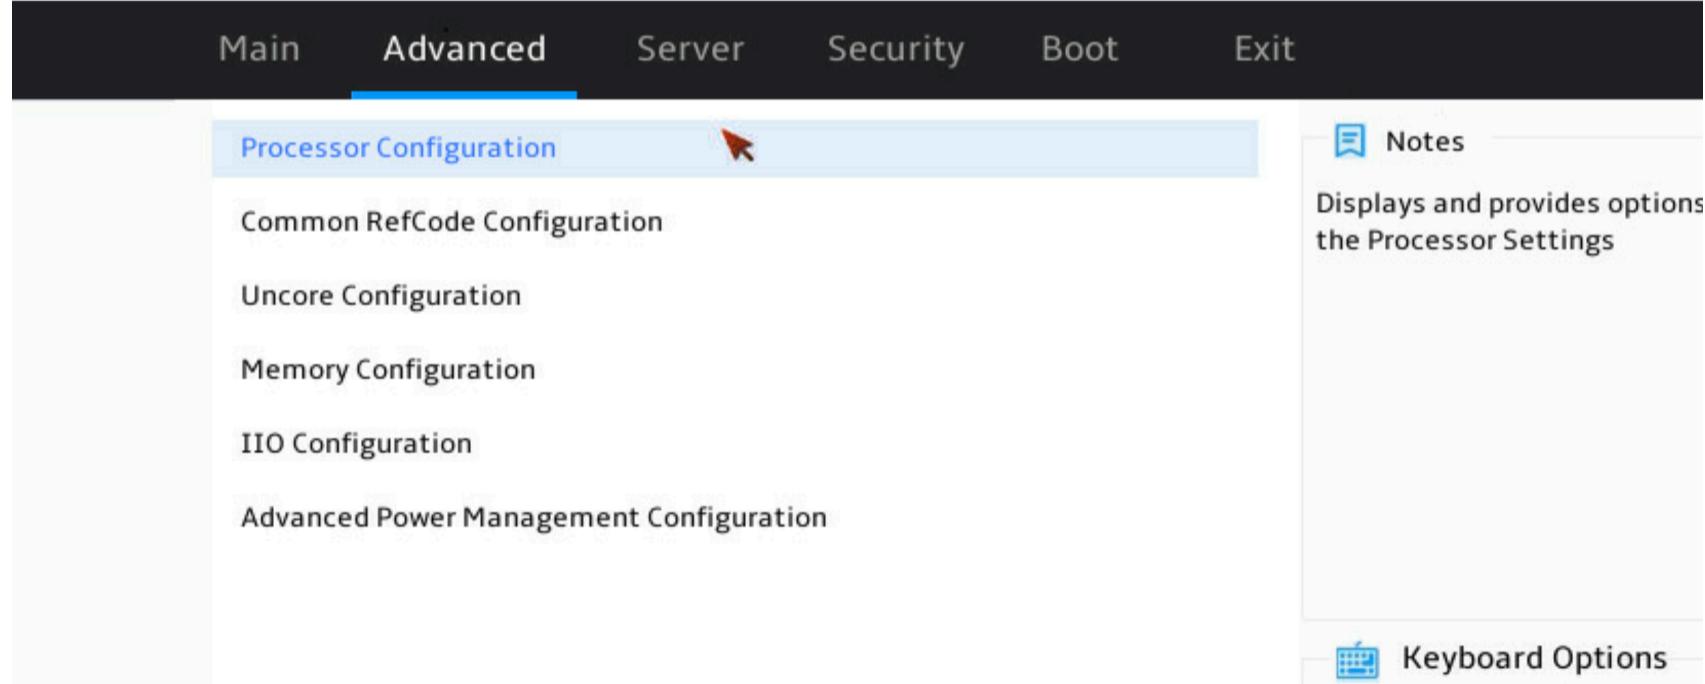
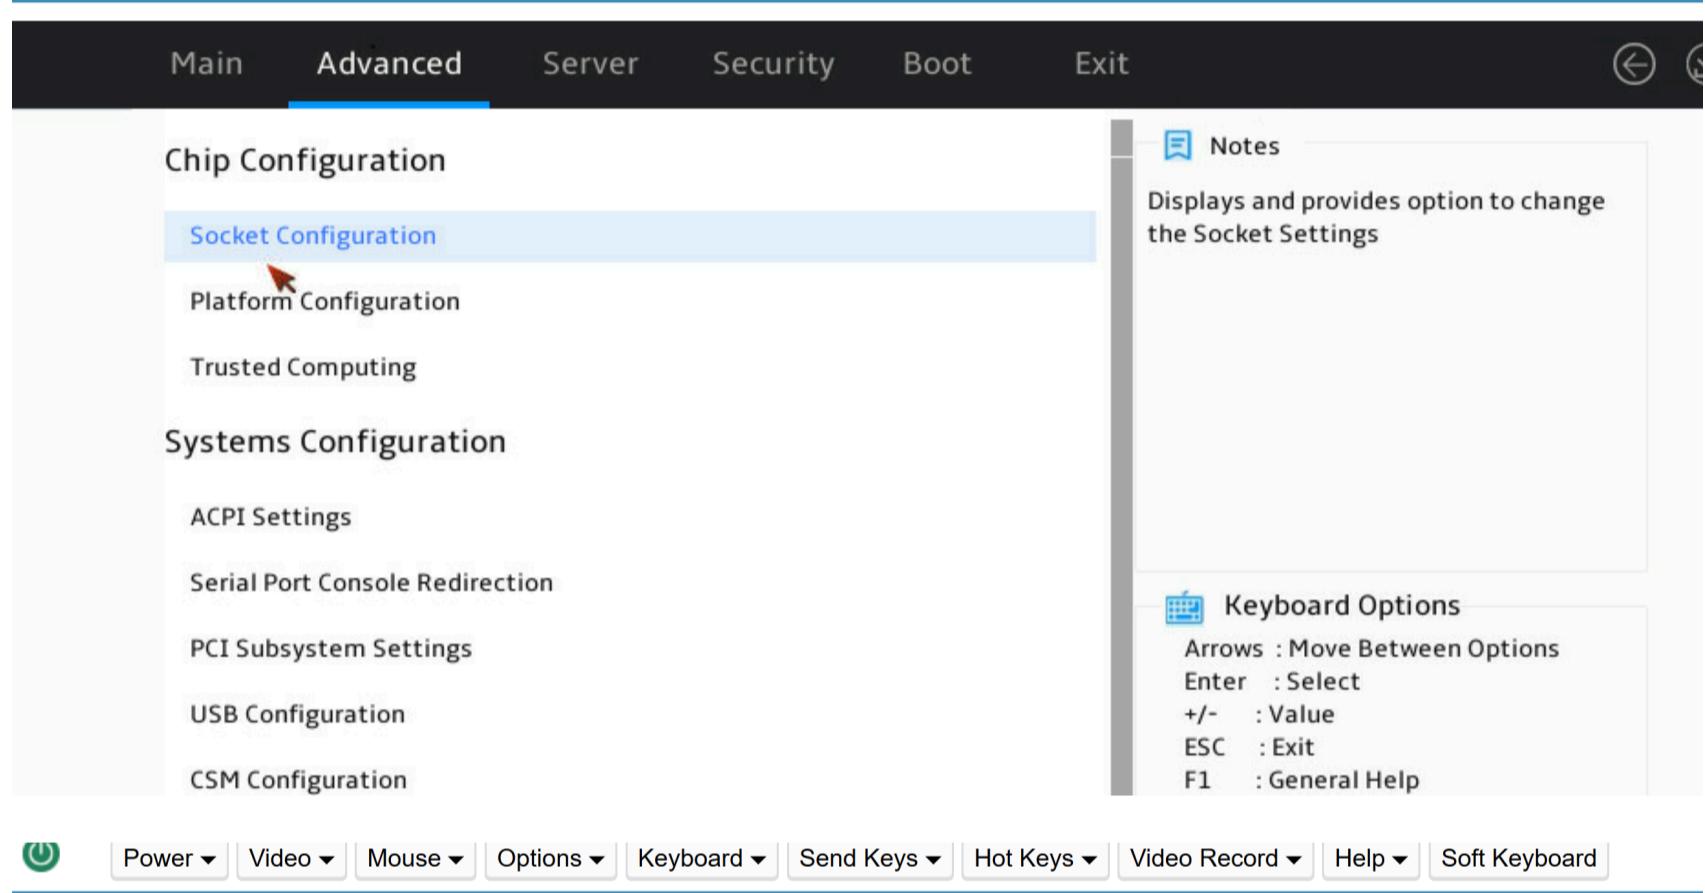
```

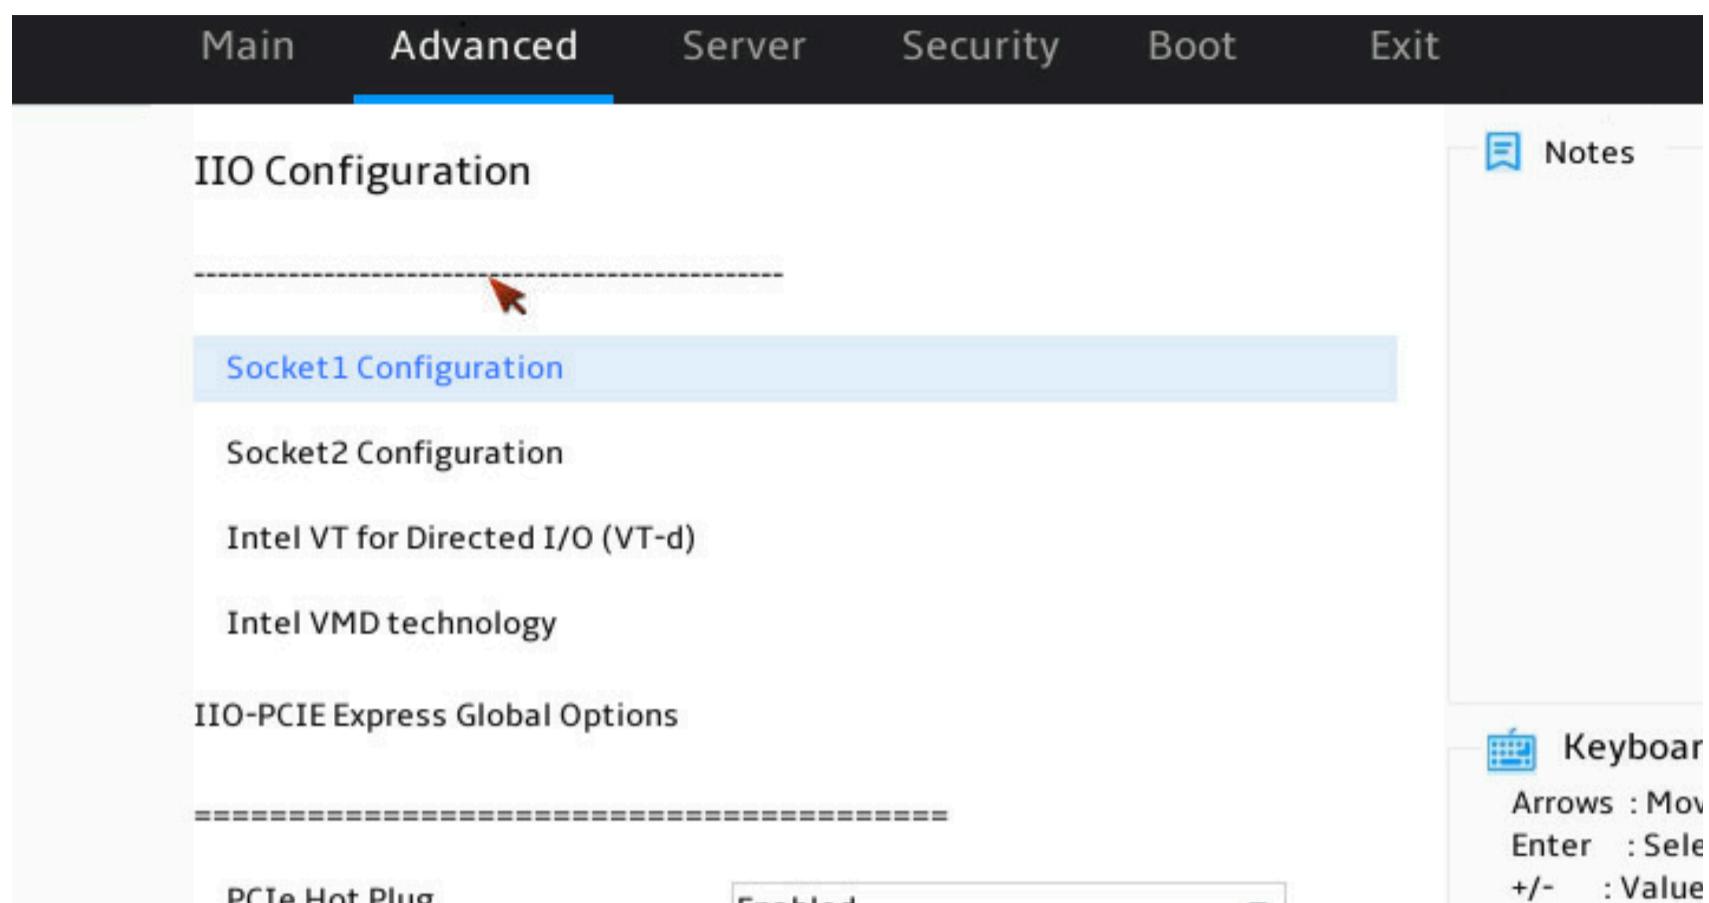
root@H800GPU-ComServer-Node09:~# cpupower frequency-info
analyzing CPU 0:
  driver: intel_pstate
  CPUs which run at the same hardware frequency: 0
  CPUs which need to have their frequency coordinated by software: 0

```

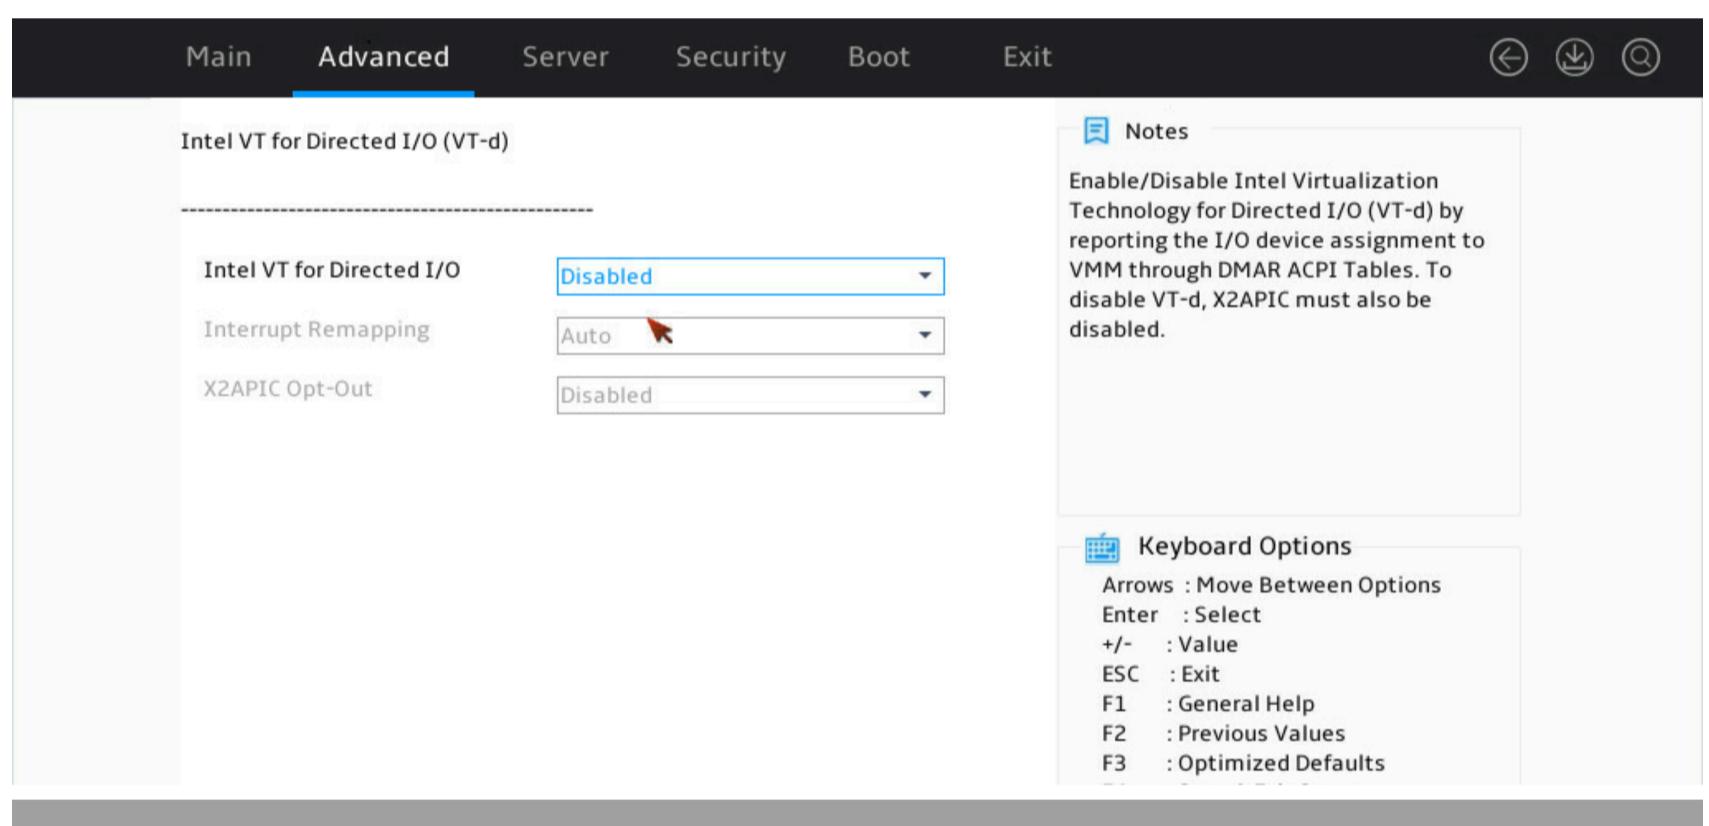
maximum transition latency: Cannot determine or is not supported.
hardware limits: 800 MHz - 3.40 GHz
available cpufreq governors: performance powersave
current policy: frequency should be within 800 MHz and 3.40 GHz.
The governor "powersave" may decide which speed to use
within this range.
current CPU frequency: Unable to call hardware
current CPU frequency: 768 MHz (asserted by call to kernel)
boost state support:
Supported: yes
Active: yes

▼ 关闭vt-d





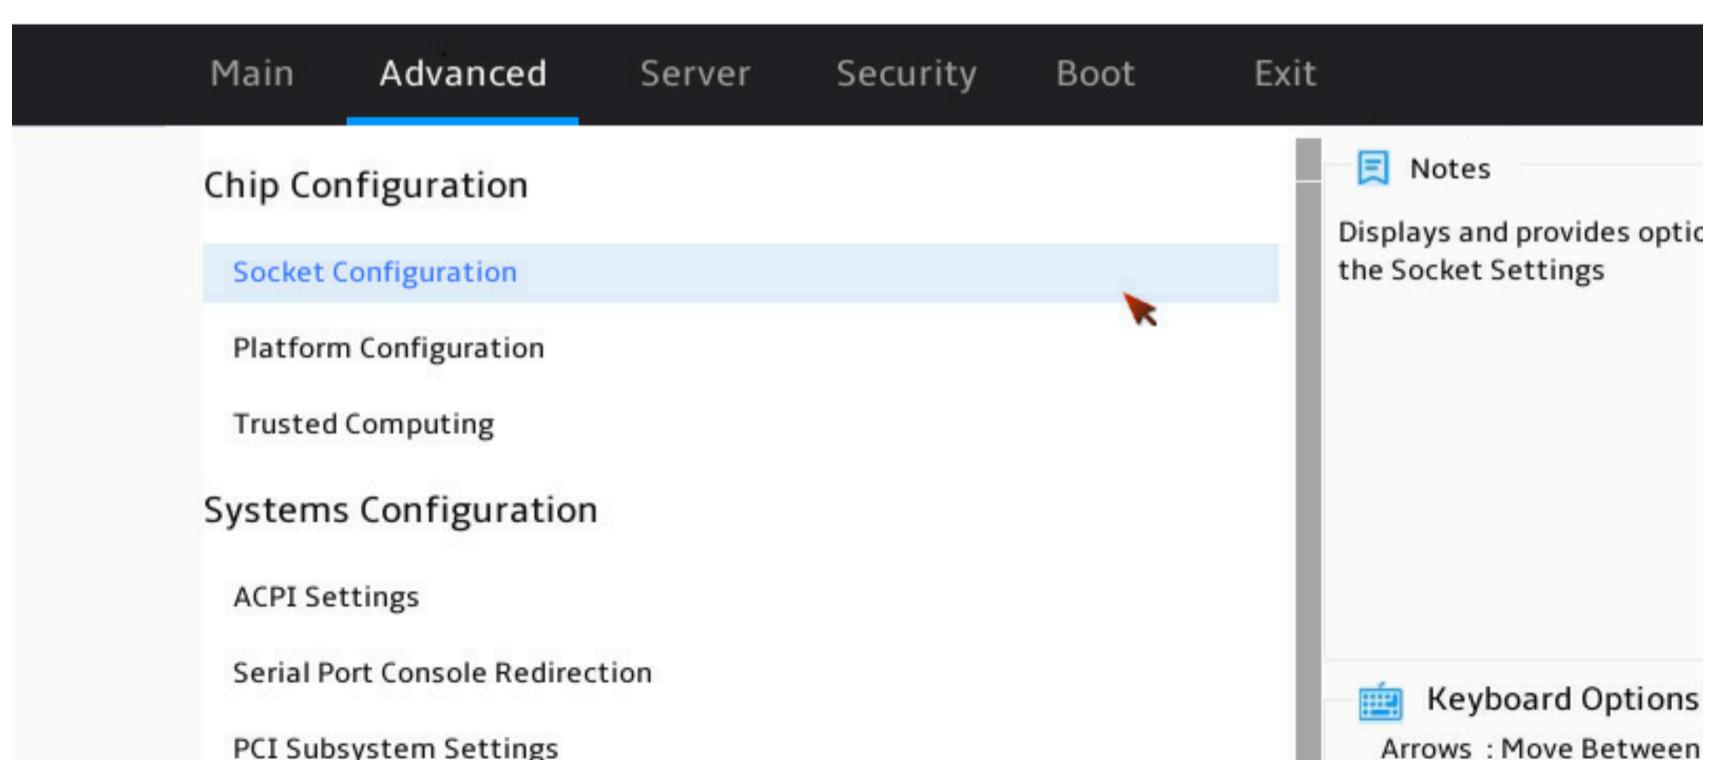
Keyboard
Arrows : Move
Enter : Select
+/- : Value



Notes
Enable/Disable Intel Virtualization Technology for Directed I/O (VT-d) by reporting the I/O device assignment to VMM through DMAR ACPI Tables. To disable VT-d, X2APIC must also be disabled.

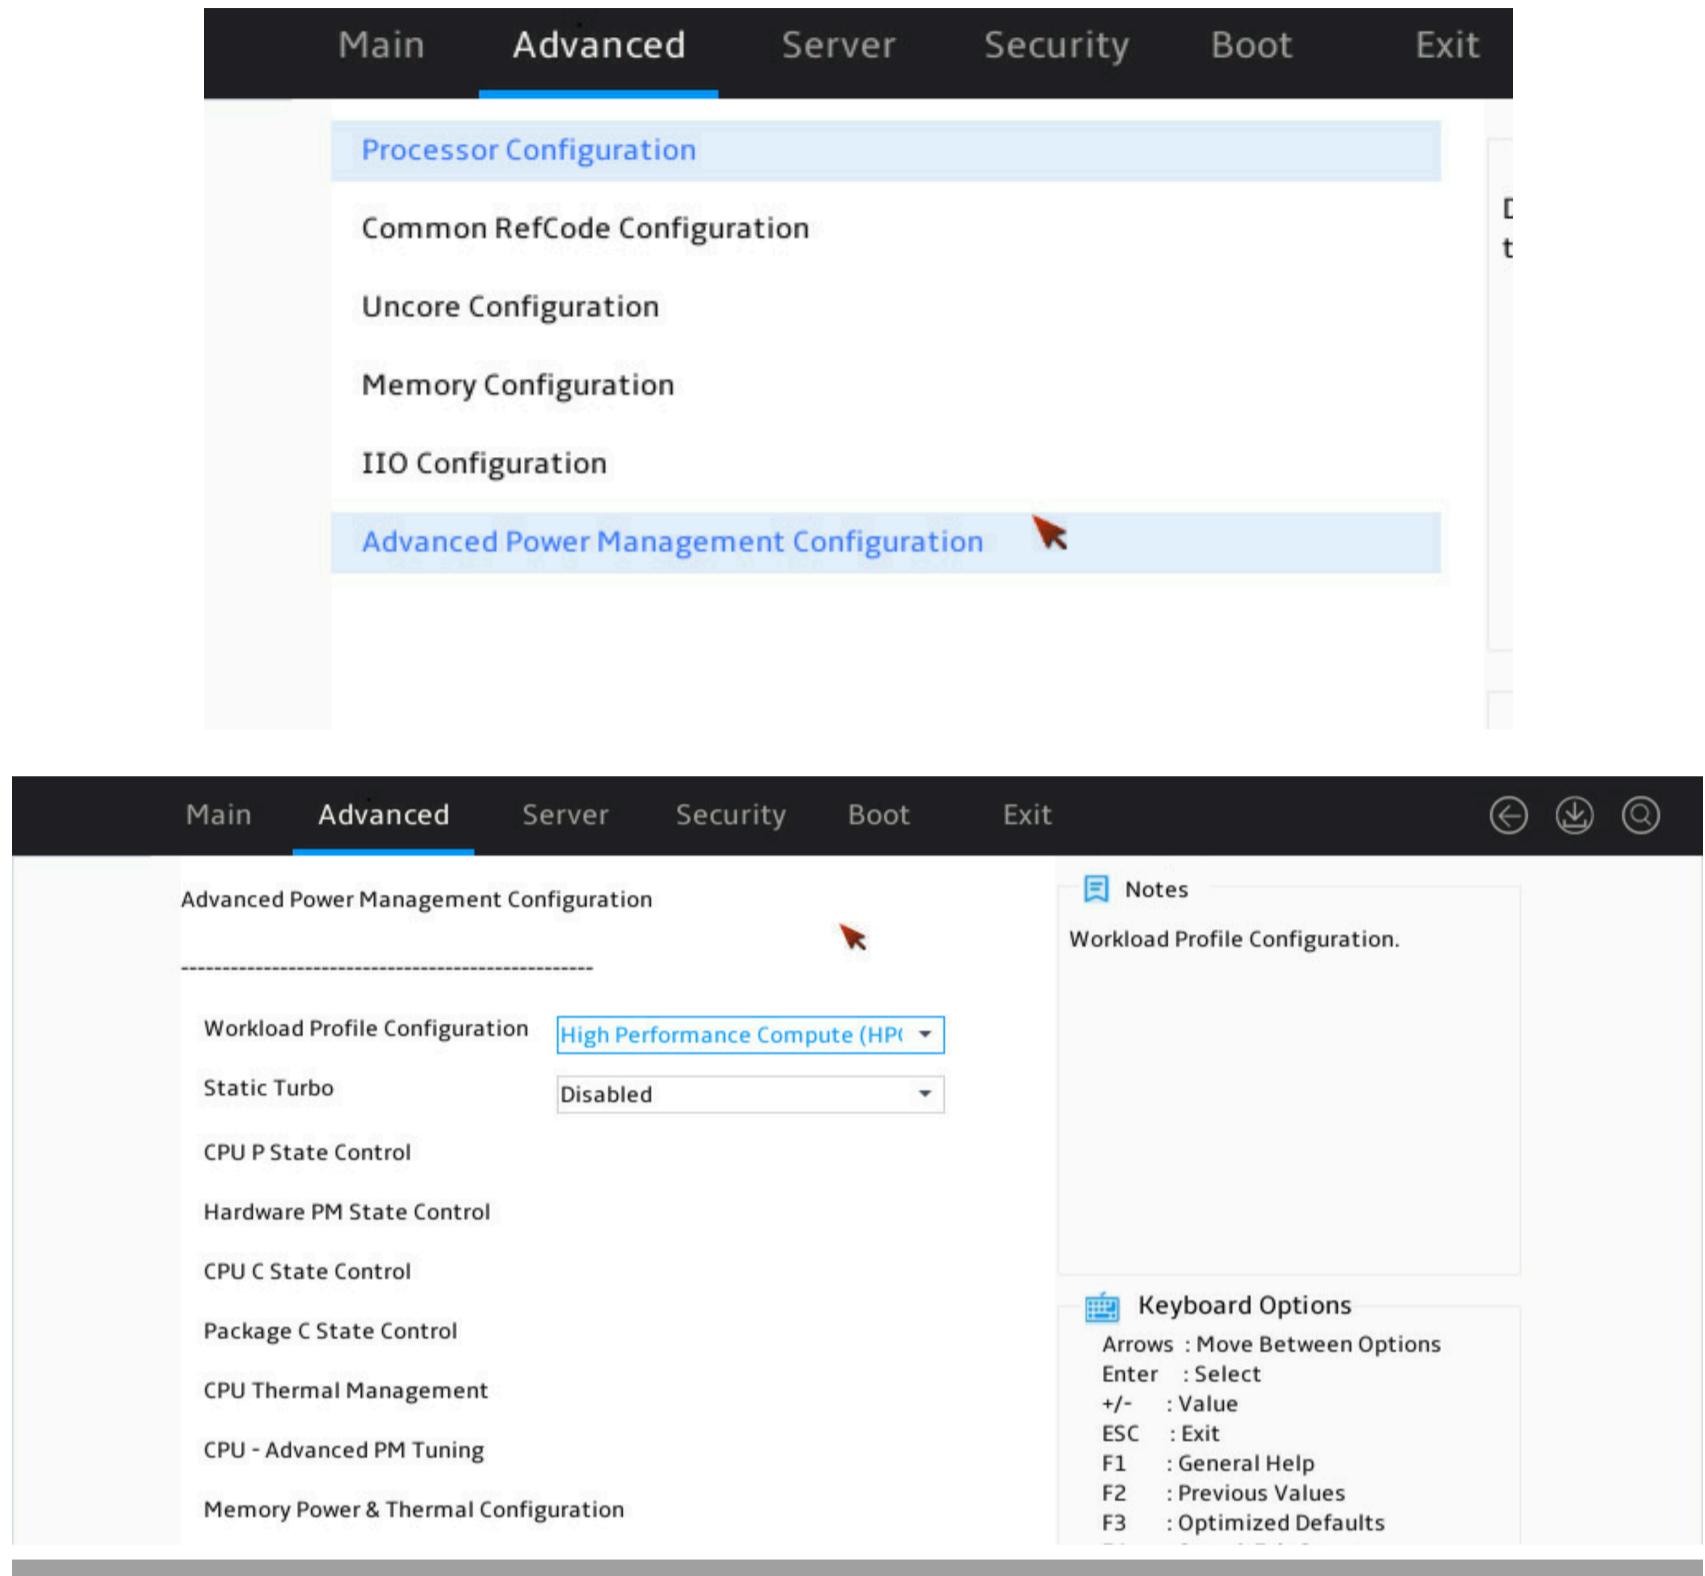
Keyboard Options
Arrows : Move Between Options
Enter : Select
+/- : Value
ESC : Exit
F1 : General Help
F2 : Previous Values
F3 : Optimized Defaults

▼ 设置CPU的性能模式



Notes
Displays and provides options for the Socket Settings

Keyboard Options
Arrows : Move Between



3. IB网卡的驱动及固件

- 安装网卡驱动OFED（开局安装驱动，包含了固件）

下载链接见上

```
tar xf MLNX_OFED_LINUX-<version>-<os>-x86_64.tar.gz
cd MLNX_OFED_LINUX-<version>-<os>
./mlnxofedinstall --all --force
/etc/init.d/openibd restart
```

- 更换网卡后，单独刷IB网卡固件

```
flint -d mlx5_0 query
flint -d mlx5_0 -i fw-ConnectX7-rel-28_43_2566-MCX75310AAS-NEA_Ax-UEFI-14.37.13-FlexBoot-3.7.500.signed.bin burn
```

4. 升级IB交换机固件（包括FW和CPLD）

- 对交换机的操作建议在对应网络（计算、存储）的UFM上进行操作。因为在拥有两个网络的设备上操作的时候，无法指定对应的网卡，从而无法升级对应的网络的交换
- 升级固件（只能升级不带管理的，升级带管理的需联系代理商）

- 检查交换机固件

查询交换机lid列表 `ibswitches`

查询基础命令 `flint -d lid-496 query`

查询批量命令

```
ibswitches | awk '{print $11}' | grep -o '[0-9]\+' | xargs -I {} sh -c 'echo "LID: {}" && flint -d lid-{} query | grep "FW Version"'
```

ii. 升级的步骤

```
wget https://www.mellanox.com/downloads/firmware/fw-Quantum-2-rel-31_2016_1028-MQM9790-NS2X_Ax.signed.bin.zip  
unzip fw-Quantum-2-rel-31_2016_1028-MQM9790-NS2X_Ax.signed.bin.zip  
flint -d lid-496 -i fw-Quantum-2-rel-31_2016_1028-MQM9790-NS2X_Ax.signed.bin burn
```

iii. 重启交换机生效

```
flint -d lid-19 swreset
```

c. 升级CPLD固件

i. 安装MFT (找一台安装即可)

```
tar xzf mft-4.xx.x-xxx.tgz  
cd mft-4.xx.x-xxx  
./install.sh --oem  
mst start  
mst status
```

ii. 检查CPLD的固件 (一般查看完以后都需要升级到最新的)

CPLD工具需要向代理商获取

基础命令 `./updateswitchcpld --unmanaged -d lid-26 --check_cpld --verbose`

批量命令

```
ibswitches | awk '{print $11}' | grep -o '[0-9]\+' | xargs -I {} sh -c 'echo "LID: {}" && ./updateswitchcpld --unmanaged -d lid-{} --check_cpld --verbose'
```

iii. 升级CPLD的固件 (永远升级最新版即可, 没有配套关系)

```
#不带管理的交换机  
./updateswitchcpld --unmanaged -d lid-<lid of switch>  
#带管理的交换机  
./updateswitchcpld --managed -t <switch's IP> -u <username> -p <password> --os mlrx-os
```

iv. 重启交换机

```
flint -d lid-19 swreset
```

5. 安装显卡驱动 (CUDA)

a. 下载CUDA

```
wget https://developer.download.nvidia.com/compute/cuda/13.0.0/local_installers/cuda_13.0.0_580.65.06_linux.run
```

b. 安装CUDA(gpu的驱动)

i. 先检查是否有默认的 (系统自带的) 显卡驱动

```
lsmod | grep nouveau
```

如果没有输出证明没有加载这个模块。如果有输出, 需要关闭

```
echo "blacklist nouveau " >>/etc/modprobe.d/blacklist.conf  
echo "options nouveau modeset=0">>>/etc/modprobe.d/blacklist.conf
```

ii. 安装CUDA

```
sh cuda_12.9.1_575.57.08_linux.run --silent --driver --toolkit --override-driver-check
```

iii. 配置环境变量

```
# 添加环境 cuda 变量  
# echo "export PATH=/usr/local/cuda-12.x/bin${PATH:+:$PATH}" >> ~/.bashrc  
echo "export PATH=/usr/local/cuda-12.6/bin${PATH:+:$PATH}" >> ~/.bashrc  
# echo "export LD_LIBRARY_PATH=/usr/local/cuda-12.x/lib64${LD_LIBRARY_PATH:+:$LD_LIBRARY_PATH}" >> ~/.bashrc  
echo "export LD_LIBRARY_PATH=/usr/local/cuda-12.6/lib64${LD_LIBRARY_PATH:+:$LD_LIBRARY_PATH}" >> ~/.bashrc  
  
source ~/.bashrc  
# 选择修改对应的cuda版本  
# 注 : export CUDA_HOME=/usr/local/cuda
```

iv. 持久化

```
cd /usr/share/doc/NVIDIA_GLX-1.0/samples/  
# 解压nvidia-persistenced-init.tar.bz2  
tar -xvf nvidia-persistenced-init.tar.bz2  
# 进入nvidia-persistenced-init文件夹  
cd nvidia-persistenced-init/  
# 执行  
. ./install.sh #,安装nvidia-persistenced  
#启动持久化服务  
systemctl start nvidia-persistenced.service  
#设置nvidia-persistenced服务开机启动  
systemctl enable nvidia-persistenced.service
```

6. 安装nv switch驱动 (fabricmanager)

a. 安装驱动之前，查看各个GPU之间的连接状态

▼ 图片如下

```
root@H800GPU-ComServer-Node09:/home/ubuntu/startup# nvidia-smi topo -p2p a  
      GPU0    GPU1    GPU2    GPU3    GPU4    GPU5    GPU6    GPU7  
GPU0     X      NS      NS      NS      NS      NS      NS      NS  
GPU1     NS     X      NS      NS      NS      NS      NS      NS  
GPU2     NS      NS     X      NS      NS      NS      NS      NS  
GPU3     NS      NS      NS     X      NS      NS      NS      NS  
GPU4     NS      NS      NS      NS     X      NS      NS      NS  
GPU5     NS      NS      NS      NS      NS     X      NS      NS  
GPU6     NS      NS      NS      NS      NS      NS     X      NS  
GPU7     NS      NS      NS      NS      NS      NS      NS     X  
  
Legend:  
X = Self  
OK = Status Ok  
CNS = Chipset not supported  
GNS = GPU not supported  
TNS = Topology not supported  
NS = Not supported  
U = Unknown  
root@H800GPU-ComServer-Node09:/home/ubuntu/startup#
```

可以看到连接状态是NS

b. 安装

```
wget https://developer.download.nvidia.com/compute/cuda/repos/ubuntu2204/x86_64/nvidia-fabricmanager-560.28.03-1_amd64.deb  
sudo dpkg -i nvidia-fabricmanager-560.28.03-1_amd64.deb  
systemctl start nvidia-fabricmanager  
systemctl enable nvidia-fabricmanager
```

c. 配置对应的服务

```
systemctl start nvidia-fabricmanager # 启动 nvidia-fabricmanager 服务  
systemctl status nvidia-fabricmanager # 查看 nvidia-fabricmanager 服务  
systemctl enable nvidia-fabricmanager # 设置开机自启
```

d. 安装完以后再次观察

▼ 图片如下

```
NIC) . mlx5_ /  
root@H800GPU-ComServer-Node09:/home/ubuntu/startup# nvidia-smi topo -p2p a  
    GPU0   GPU1   GPU2   GPU3   GPU4   GPU5   GPU6   GPU7  
GPU0   X      OK     OK     OK     OK     OK     OK     OK  
GPU1   OK     X      OK     OK     OK     OK     OK     OK  
GPU2   OK     OK     X      OK     OK     OK     OK     OK  
GPU3   OK     OK     OK     X      OK     OK     OK     OK  
GPU4   OK     OK     OK     OK     X      OK     OK     OK  
GPU5   OK     OK     OK     OK     OK     X      OK     OK  
GPU6   OK     OK     OK     OK     OK     OK     X      OK  
GPU7   OK     OK     OK     OK     OK     OK     OK     X  
  
Legend:  
X = Self  
OK = Status OK  
CNS = Chipset not supported  
GNS = GPU not supported  
TNS = Topology not supported  
NS = Not supported  
U = Unknown  
root@H800GPU-ComServer-Node09:/home/ubuntu/startup#
```

7. 安装docker支持(视需要而定)

```
wget https://developer.download.nvidia.com/compute/cuda/repos/ubuntu2204/x86_64/cuda-keyring_1.1-1_all.deb  
sudo dpkg -i cuda-keyring_1.1-1_all.deb  
sudo apt-get update  
apt install libnvidia-container-tools nvidia-container-toolkit
```

8. 做系统基础配置例如network ntp dns log压缩 配置

9. 多数情况下，可以把相关的配置写成一个ansible的playbook

10. 打流测试包括IB的，以太的

a. IB打流测试

i. 带宽

- 前提条件两台设备要可以相互ping。这个ip不一定必须要配置在IB网卡上面，只需要相互ping通就可以了
- 一端设备开始被测试（这台设备的IP是 10.232.87.102）

```
ib_write_bw -d -q 10 mlx5_ 0 --report_gbits
```

- 一端设备开始测试

```
ib_write_bw -d mlx5_0 -q 10 --report_gbits --run_infinitely 10.232.87.102
```

ii. 延时

- 前提条件两台设备要可以相互ping。这个ip不一定必须要配置在IB网卡上面，只需要相互ping通就可以了
- 一端设备开始被测试（这台设备的IP是 10.232.87.102）

```
ib_write_lat -d mlx5_0
```

- 一端设备开始测试
`ib_write_lat -d mlx5_0 10.232.87.102`
- 关于结果，主要参考绿色部分的三个值

```
root@UFM-StorNet-Node01:~# ib_write_lat -d mlx5_0 10.232.87.102
-----
          RDMA_Write Latency Test
Dual-port      : OFF      Device      : mlx5_0
Number of qps  : 1       Transport type : IB
Connection type : RC     Using SRQ    : OFF
PCIe relax order: OFF    Lock-free   : OFF
ibv_wr* API    : ON
TX depth       : 1
Mtu           : 4096[B]
Link type      : IB
Max inline data : 220[B]
rdma_cm QPs   : OFF
Data ex. method : Ethernet
-----
local address: LID 0x01 QPN 0x002f PSN 0x5c5957 RKey 0x1ffccb VAddr 0x0056000d5ac000
remote address: LID 0x13d QPN 0x002f PSN 0x6ac2da RKey 0x1ff5b3 VAddr 0x00555bbb6cf000
-----
#bytes #iterations  t_min[usec]  t_max[usec]  t_typical[usec]  t_avg[usec]  t_stdev[usec]  99% per
centile[usec]  99.9% percentile[usec]
Conflicting CPU frequency values detected: 2100.000000 != 1957.948000. CPU Frequency is not max.
2      1000      1.60      3.88      1.84      1.84      0.03      1.89      3.88
-----
```

b. 以太打流测试

i. 带宽

- 前提条件两台设备要可以相互ping。这个ip不一定必须要配置在IB网卡上面，只需要相互ping通就可以了
- 一端设备开始被测试（这台设备的IP是 `10.232.87.102`）

`iperf3 -s`

- 一端设备开始测试
`iperf3 -c 10.232.87.102`
- 注意

lacp模式下，不加参数只能测得一根链路的带宽，如果想要测试bond链路的带宽，需要根据lacp的模式调整测试方案

11. gpu burn测试

a. 下载并编译gpu burn

```
git clone https://github.com/wilicc/gpu-burn.git
cd gpu-burn
make
```

b. 指定使用那张卡进行测试

`CUDA_VISIBLE_DEVICES=0,1,2,3,4,5,6,7 ./gpu_burn 300`

c. 查看GPU的状态

`watch -n 1 nvidia-smi`

d. 结果

```
Uninitted cublas
Freed memory for dev 0
Uninitted cublas
done

Tested 8 GPUs:
  GPU 0: OK
  GPU 1: OK
  GPU 2: OK
  GPU 3: OK
  GPU 4: OK
  GPU 5: OK
  GPU 6: OK
  GPU 7: OK
root@H800GPU-ComServer-Node09:/home/ubuntu/startup/gpu-burn#
```

12. nccl 安装及测试

a. 先安装mpirun

```
apt-get install openmpi-bin openmpi-common libopenmpi-dev
```

b. nccl

下载和编译

```
git clone https://github.com/NVIDIA/nccl.git
cd nccl
make -j src.build
```

安装

```
# Install tools to create debian packages
sudo apt install build-essential devscripts debhelper fakeroot
# Build NCCL deb package
make pkg.debian.build
# 还在nccl那个工作目录下
ls build/pkg/deb/
# 还在nccl那个工作目录下
dpkg -i build/pkg/deb/*.deb
```

c. nccl_test

下载并编译

```
git clone https://github.com/NVIDIA/nccl-tests.git
# 这里如果如果CUDA没有安装到/usr/local/cuda需要执行CUDA的安装目录
# 这里如果没有指定nccl的安装命令，需要指定nccl的安装目录
# 如果都是默认安装这不需要指定
# make CUDA_HOME=/path/to/cuda NCCL_HOME=/path/to/nccl
#进入nccl-tests目录
cd nccl-tests
make MPI=1 MPI_HOME=/usr/lib/x86_64-linux-gnu/openmpi
```

d. 使用nccl_test进行测试

i. 单点nccl-tests执行

```
./build/all_reduce_perf -b 128M -e 8G -f 2 -g 8
```

或者

```
mpirun --allow-run-as-root \
```

```
-np 8 ./build/all_reduce_perf -b 128M -e 8G -f 2
```

结果

```
root@H800GPU-ComServer-Node01:/home/ubuntu/startup/nccl-tests# 
root@H800GPU-ComServer-Node01:/home/ubuntu/startup/nccl-tests# mpirun --allow-run-as-root \
-np 8 ./build/all_reduce_perf -b 128M -e 8G -f 2
# Collective test starting: all_reduce_perf
# nThread 1 nGpus 1 minBytes 134217728 maxBytes 8589934592 step: 2(factor) warmup iters: 5 iters: 20 agg iters: 1 validation
: 1 graph: 0
#
# Using devices
# Rank 0 Group 0 Pid 3928604 on H800GPU-ComServer-Node01 device 0 [0000:19:00] NVIDIA H800
# Rank 1 Group 0 Pid 3928605 on H800GPU-ComServer-Node01 device 1 [0000:3a:00] NVIDIA H800
# Rank 2 Group 0 Pid 3928606 on H800GPU-ComServer-Node01 device 2 [0000:4c:00] NVIDIA H800
# Rank 3 Group 0 Pid 3928607 on H800GPU-ComServer-Node01 device 3 [0000:5c:00] NVIDIA H800
# Rank 4 Group 0 Pid 3928608 on H800GPU-ComServer-Node01 device 4 [0000:9a:00] NVIDIA H800
# Rank 5 Group 0 Pid 3928610 on H800GPU-ComServer-Node01 device 5 [0000:ba:00] NVIDIA H800
# Rank 6 Group 0 Pid 3928612 on H800GPU-ComServer-Node01 device 6 [0000:ca:00] NVIDIA H800
# Rank 7 Group 0 Pid 3928614 on H800GPU-ComServer-Node01 device 7 [0000:da:00] NVIDIA H800
#
#          out-of-place           in-place
# size      count     type   redop   root    time   algbw busbw #wrong   time   algbw busbw #wrong
# (B)      (elements)
134217728  33554432   float   sum    -1  1184.7 113.29 198.27   0  1185.6 113.20 198.11   0
268435456  67108864   float   sum    -1  2282.2 117.62 205.83   0  2281.8 117.64 205.87   0
536870912  134217728   float   sum    -1  4469.0 120.13 210.23   0  4464.5 120.25 210.44   0
1073741824 268435456   float   sum    -1  8825.5 121.66 212.91   0  8824.3 121.68 212.94   0
2147483648 536870912   float   sum    -1  17530 122.50 214.38   0  17525 122.54 214.45   0
4294967296 1073741824   float   sum    -1  34979 122.79 214.88   0  34972 122.81 214.92   0
8589934592 2147483648   float   sum    -1  69877 122.93 215.13   0  69846 122.98 215.22   0
#
# Out of bounds values : 0 OK
# Avg bus bandwidth : 210.255
#
# Collective test concluded: all_reduce_perf
```

ii. 多机nccl测试

官方的troubleshooting页面

<https://docs.nvidia.com/deeplearning/nccl/user-guide/docs/troubleshooting.html#pci-access-control-services-acss>

1. 提前配置

查询SrcValid+

```
lspci -vvv | grep ACSCtl | grep SrcValid+
```

关闭SrcValid+

```
lspci -vvv | grep 'PCI bridge' | cut -d' ' -f 1 | xargs -l{} setpci -s {} ECAP_ACS+06.w=0000
```

2. 文件准备 myhosts2

```
root@H800GPU-ComServer-Node09:/home/ubuntu/startup/nccl-tests# cat myhosts2
gpu2 slots=8
gpu4 slots=8
gpu5 slots=8
gpu6 slots=8
gpu1 slots=8
```

3. 通过mpirun运行

```
mpirun --allow-run-as-root \
-hostfile myhosts2 \
-np 32 \
-npernode 8 \
-x LD_LIBRARY_PATH \
-x NCCL_DEBUG=INFO \
-x NCCL_IB_HCA=mlx5_1,mlx5_2,mlx5_4,mlx5_7 \
-x NCCL_SOCKET_IFNAME=bond0 \
./build/all_reduce_perf -b 128M -e 8G -f 2 -g 1
```

```
root@H800GPU-ComServer-Node09:/home/ubuntu/startup/nccl-tests# mpirun --allow-run-as-root -h
ostfile myhosts2 -np 32 -npernode 8 -x LD_LIBRARY_PATH -x NCCL_DEBUG=INFO -x NCCL_IB_HCA=
```

```
mlx5_1,mlx5_2,mlx5_4,mlx5_7 -x NCCL_SOCKET_IFNAME=bond0 ./build/all_reduce_perf -b 769M -e 8  
G -f 2 -g 1 -i 0
```

```
root@H800GPU-ComServer-Node09:/home/ubuntu/startup/nccl-tests# mpirun --allow-run-as-root -hostfile myhosts2 -np 32 -npernode 8 -x LD_LIBRARY_PATH -x  
NCCL_DEBUG=INFO -x NCCL_IB_HCA=mlx5_1,mlx5_2,mlx5_4,mlx5_7 -x NCCL_SOCKET_IFNAME=bond0 ./build/all_reduce_perf -b 769M -e 8G -f 2 -g 1 -i 0
```

总结

1. 需要安装的驱动

ib卡驱动OFED, gpu卡驱动CUDA, nvswitch驱动fabric manager

2. 概念的总结

IB网卡需要的安装包OFED是一个包含了固件、库、驱动、工具集的合集

GPU网卡安装包 CUDA是一个包含了固件、库、驱动、工具集的合集

正常安装驱动，一般就是一个deb包，当然这个包也有对应的依赖

3. 关于硬件选型

400G QSFP112 OSFP112 (APC端面，端面是指的模块和线缆的接口)

200G 都是Q(sfp)56的。Q的是向后兼容的。Q112向下兼容Q56 Q28

Q112 400G APC

Q56 200G UPC/APC (其实是Q112降速) 较贵

Q28 100G UPC

VR SR DR FR LR这个选型要清晰

多模中OM3 OM4

模式中，单模虽然贵，但是贵在稳定。即使短距离使用单模也还是稳定的。

4. 应该如何升级关注什么

a. 交换机要和网卡的固件版本匹配

b. 在官方建议中（见上“查询最新版本的连接”）可以看到最新的版本中对应的交换机的和网卡的固件版本。建议升级到最新的LTS版本。因为如果遇到问题开case nvidia基本上会要求先升版本

c. 交换机需要关注的点：固件、CPLD (其实就是电源风扇模块的固件)。带管理的交换机，需要升级系统。升级系统的同时，就升级了固件，所以不能单独升级固件。

交换机的固件要统一

d. 网卡需要关注的点：OFED的版本应该和固件所匹配。尤其是在换网卡的时候需要注意。

网络中的网卡固件要统一。

5. 网卡的对应关系

a. PCI插槽和网卡的对应关系

```
root@UFM-StorNet-Node01:~# lspci | grep -i mellanox  
17:00.0 Infiniband controller: Mellanox Technologies MT28908 Family [ConnectX-6]  
31:00.0 Infiniband controller: Mellanox Technologies MT28908 Family [ConnectX-6]  
32:00.0 Ethernet controller: Mellanox Technologies MT27710 Family [ConnectX-4 Lx]  
32:00.1 Ethernet controller: Mellanox Technologies MT27710 Family [ConnectX-4 Lx]
```

网卡的状态

```
root@UFM-StorNet-Node01:~# ibstat  
CA 'mlx5_0'  
    CA type: MT4123  
    Number of ports: 1  
    Firmware version: 20.43.3608  
    Hardware version: 0  
    Node GUID: 0x946dae03000ed0a0
```

```

System image GUID: 0x946dae03000ed0a0
Port 1:
    State: Active
    Physical state: LinkUp
    Rate: 200
    Base lid: 1
    LMC: 0
    SM lid: 1
    Capability mask: 0xa651e84a
    Port GUID: 0x946dae03000ed0a0
    Link layer: InfiniBand

CA 'mlx5_1'
    CA type: MT4123
    Number of ports: 1
    Firmware version: 20.40.1000
    Hardware version: 0
    Node GUID: 0xb83fd203007b92c2
    System image GUID: 0xb83fd203007b92c2
    Port 1:
        State: Down
        Physical state: Polling
        Rate: 10
        Base lid: 65535
        LMC: 0
        SM lid: 0
        Capability mask: 0xa651e848
        Port GUID: 0xb83fd203007b92c2
        Link layer: InfiniBand

CA 'mlx5_bond_0'
    CA type: MT4117
    Number of ports: 1
    Firmware version: 14.32.1010
    Hardware version: 0
    Node GUID: 0x30b9300300354afd
    System image GUID: 0x30b9300300354afd
    Port 1:
        State: Active
        Physical state: LinkUp
        Rate: 10
        Base lid: 0
        LMC: 0
        SM lid: 0
        Capability mask: 0x00010000
        Port GUID: 0x28de01ffe4f2803
        Link layer: Ethernet

```

网卡和PCI-e的关系

ls -l /sys/class/infiniband/mlx5_1/device

```

root@UFM-StorNet-Node01:~# ls -l /sys/class/infiniband/mlx5_1/device
lrwxrwxrwx 1 root root 0 Aug  6 11:30 /sys/class/infiniband/mlx5_1/device → ../../0000:31:00.0
root@UFM-StorNet-Node01:~# ls -l /sys/class/infiniband/mlx5_0/device
lrwxrwxrwx 1 root root 0 Aug  6 11:30 /sys/class/infiniband/mlx5_0/device → ../../0000:17:00.0

```

b. 网卡和对应的以太的关系

```
root@UFM-StorNet-Node01:~# ibdev2netdev
mlx5_0 port 1 => ibs5 (Down)
mlx5_1 port 1 => ibs3 (Down)
mlx5_bond_0 port 1 => bond0 (Up)
```

6. IP和IB的对应关系

两种概念

- a. 我的网卡工作在IP模式下

这种情况下，可以把IB网卡就理解为正常的IP网卡就可以了。完全一模一样

- b. 我的网卡工作在IB模式下

例如上面我们看到IB和IP的对应关系

```
root@UFM-StorNet-Node01:~# ibdev2netdev
mlx5_0 port 1 => ibs5 (Down)
mlx5_1 port 1 => ibs3 (Down)
mlx5_bond_0 port 1 => bond0 (Up)
```

首先在这里看到Down与Up需要配置一个IP地址，且执行 `ip link set dev ibs5 up` 就可以up对应的以太网端口了。

IP地址是跑在IP over IB的模式下。正常的RDMA访问仍然是IB数据包。IP仅用于ping或者找到主机方便或者握手或者心跳。